


Data Plane and P4

Version 1, Presented at QuTech (Delft) on 12 September 2019



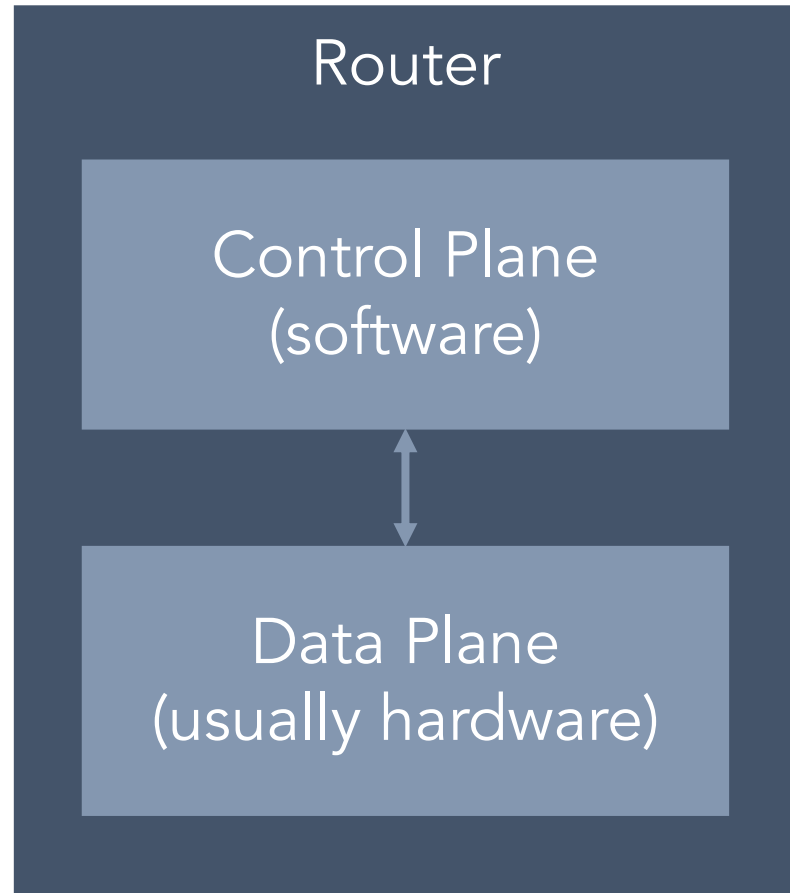
Agenda

1. Control plane versus data plane
2. Fixed-function classical data plane
3. Programmable classical data plane
4. Control plane – data plane interface
5. Programmable quantum data plane

Section 1

Control plane versus data plane

Control plane versus data plane



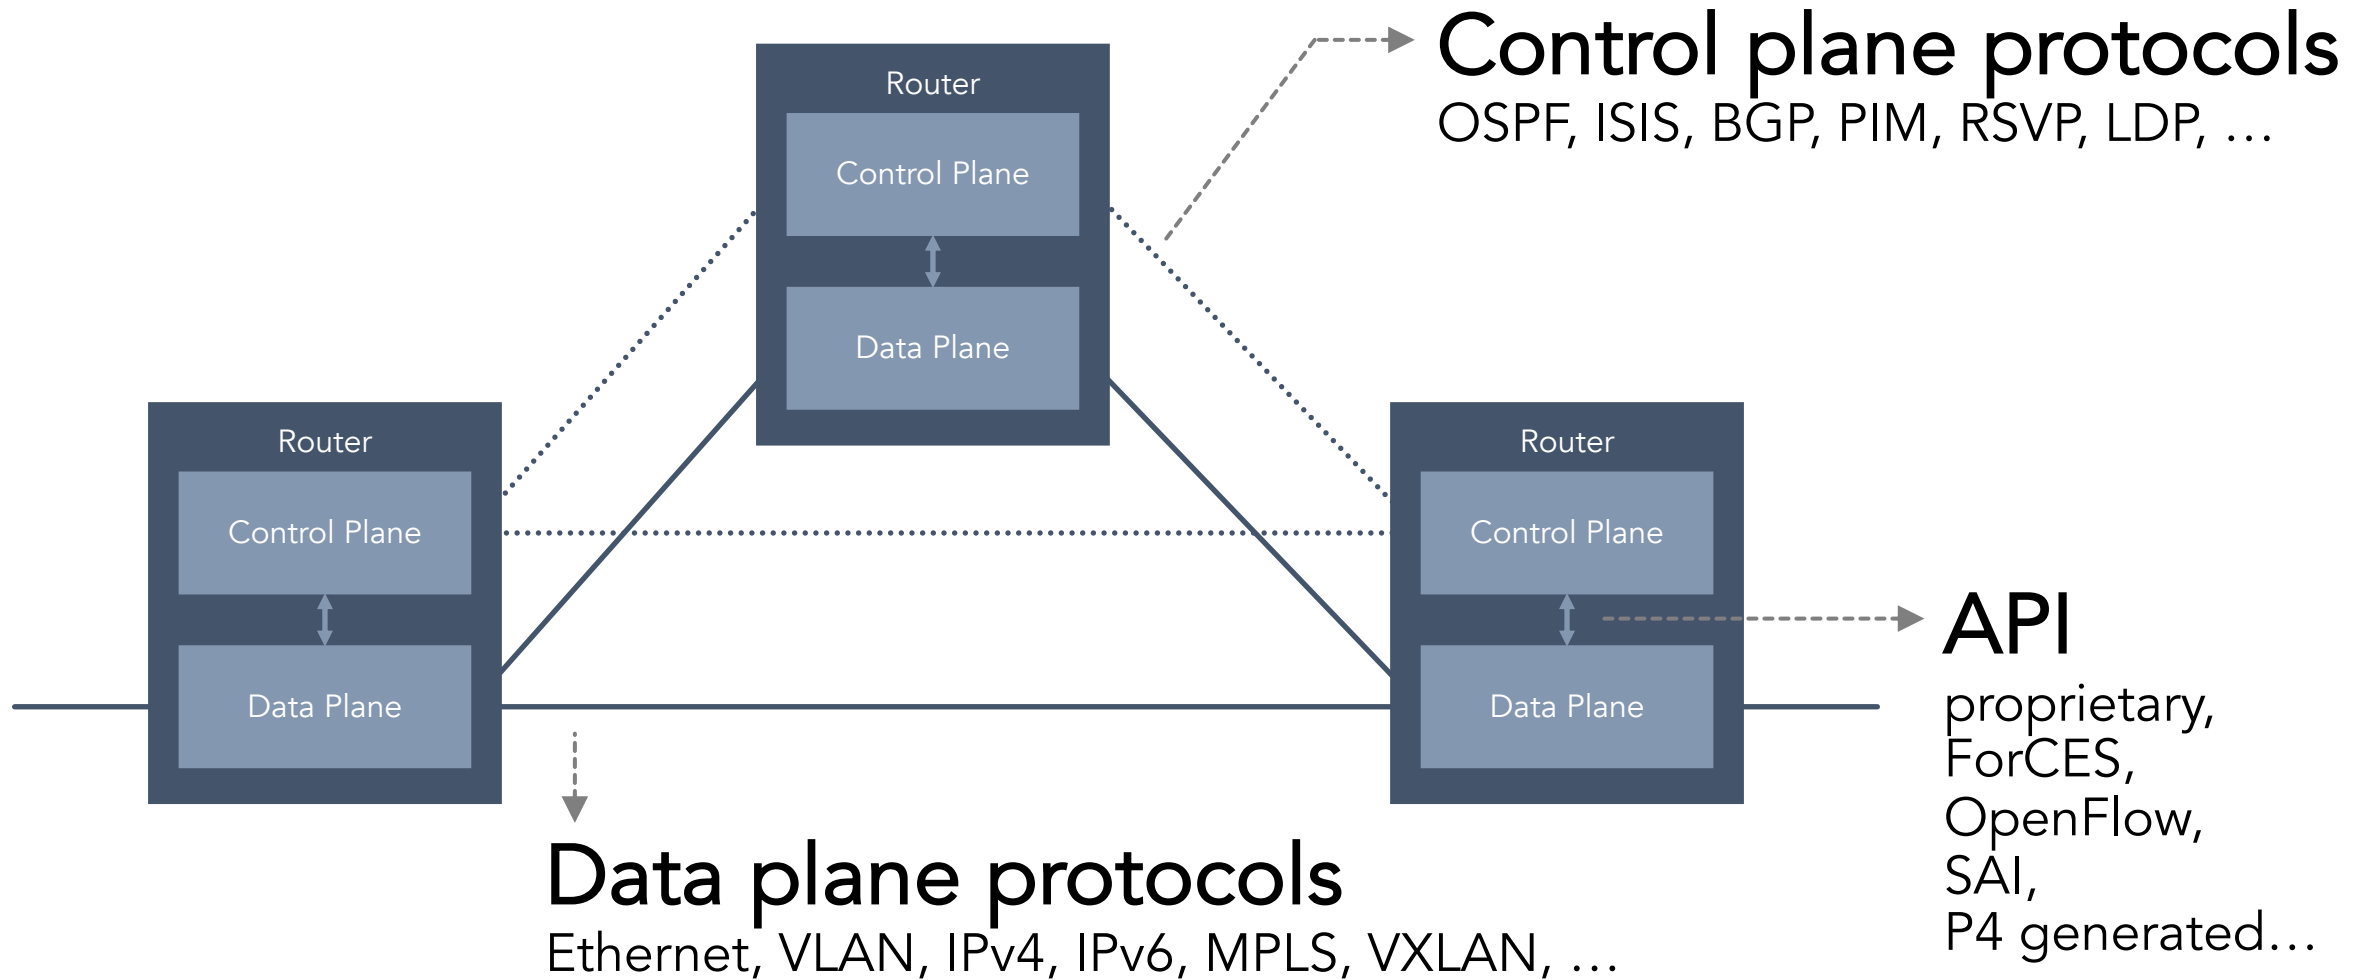
Control plane

- Discover network topology and resources
- Program the forwarding plane (e.g. forwarding table)
- Flexibility is main concern
- Software on general purpose OS and CPU

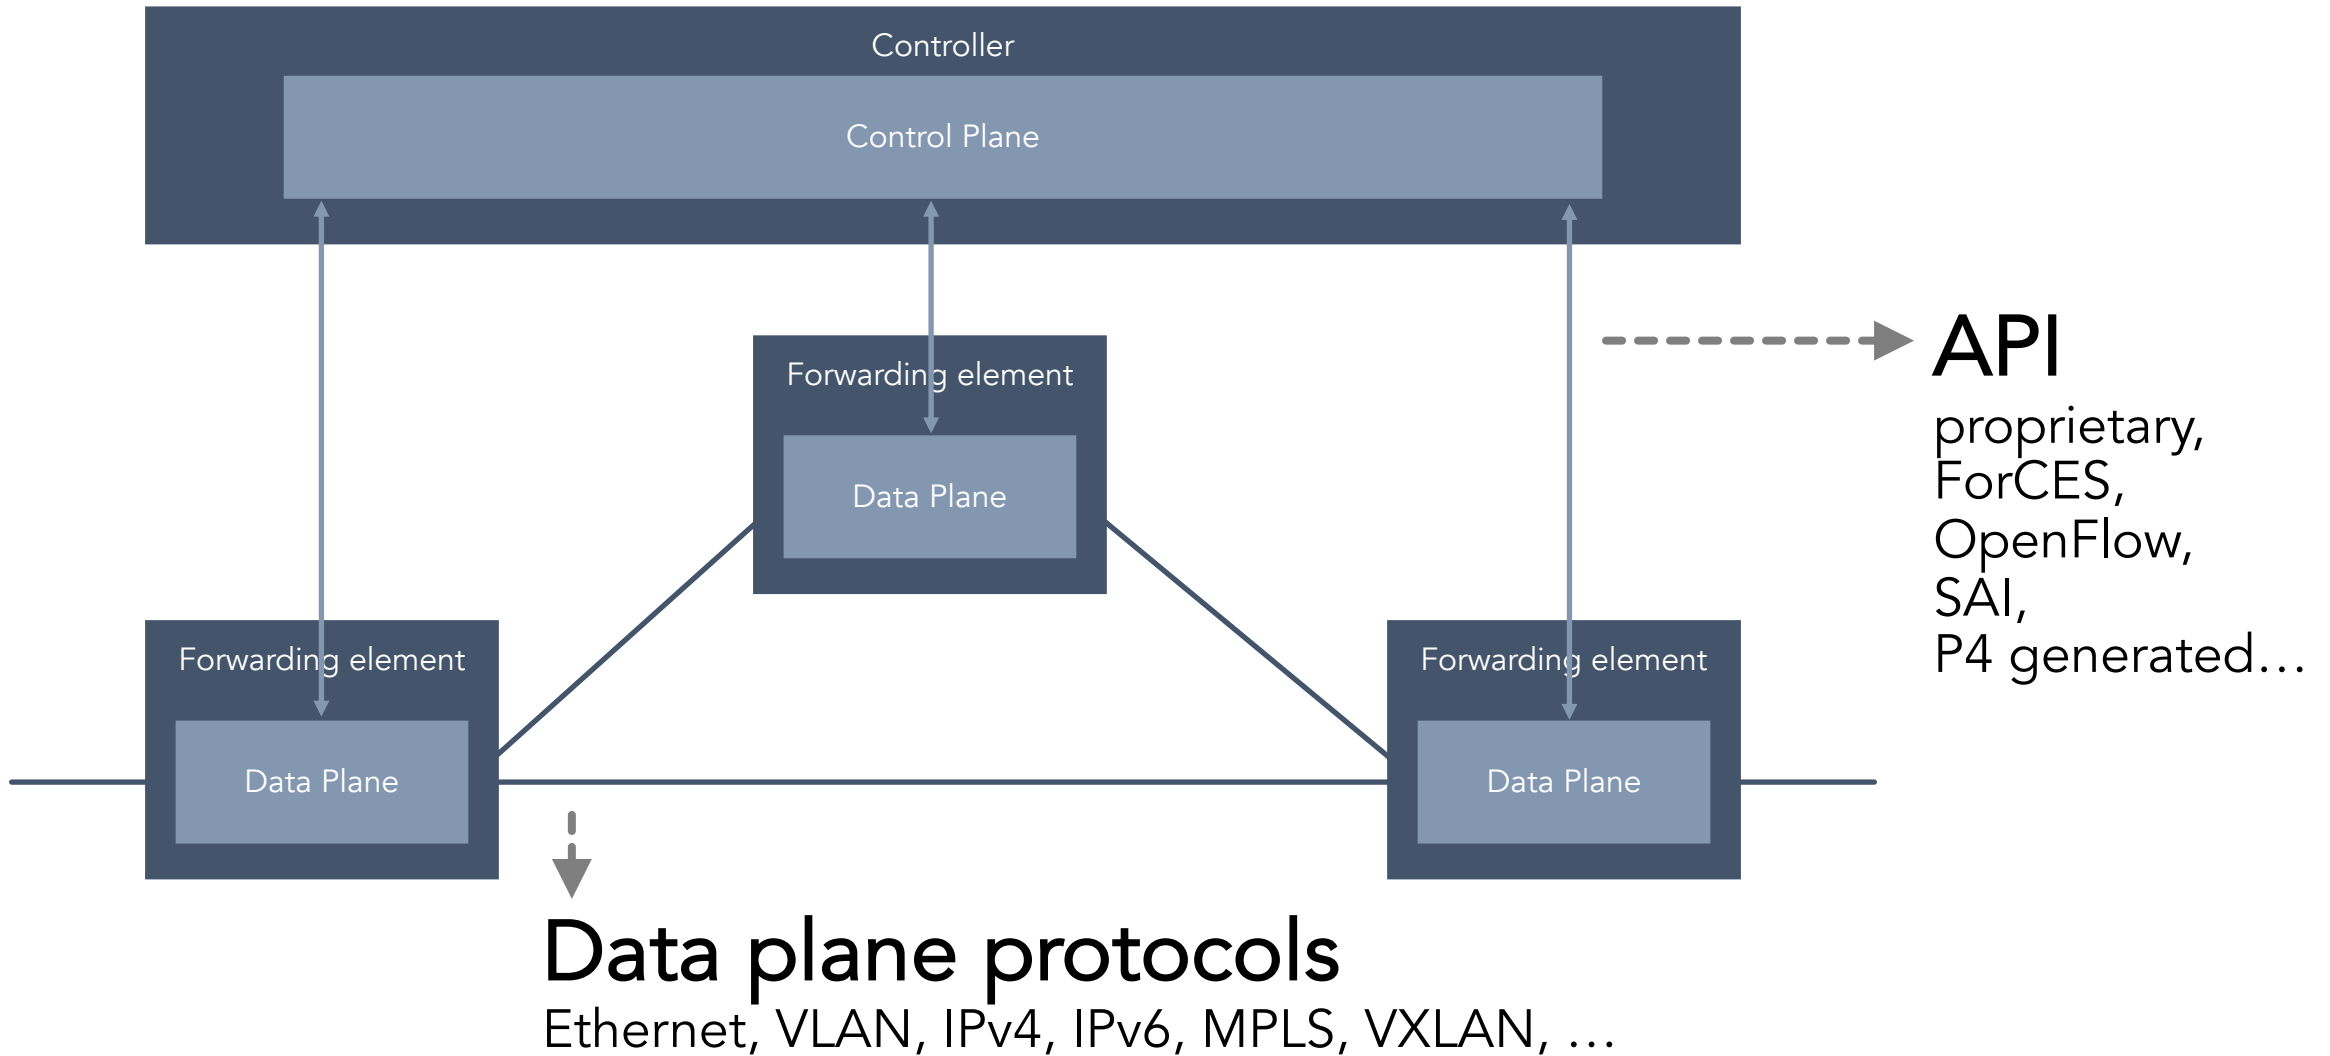
Data plane (aka forwarding plane)

- Forward packets
- Using the forwarding table populated by the control plane
- Speed is main concern
- Fixed-function hardware, programmable hardware, software

"Traditional" distributed control plane



SDN: centralized control plane (simplified)



More planes



Management plane

Configure and monitor device
General purpose CPU with general purpose OS
Examples: SSH, NETCONF, SNMP, ...

Control plane

Routing protocols, signaling protocols
General purpose CPU with general purpose OS
Examples: OSPF, ISIS, BGP, LDP, PIM, ...

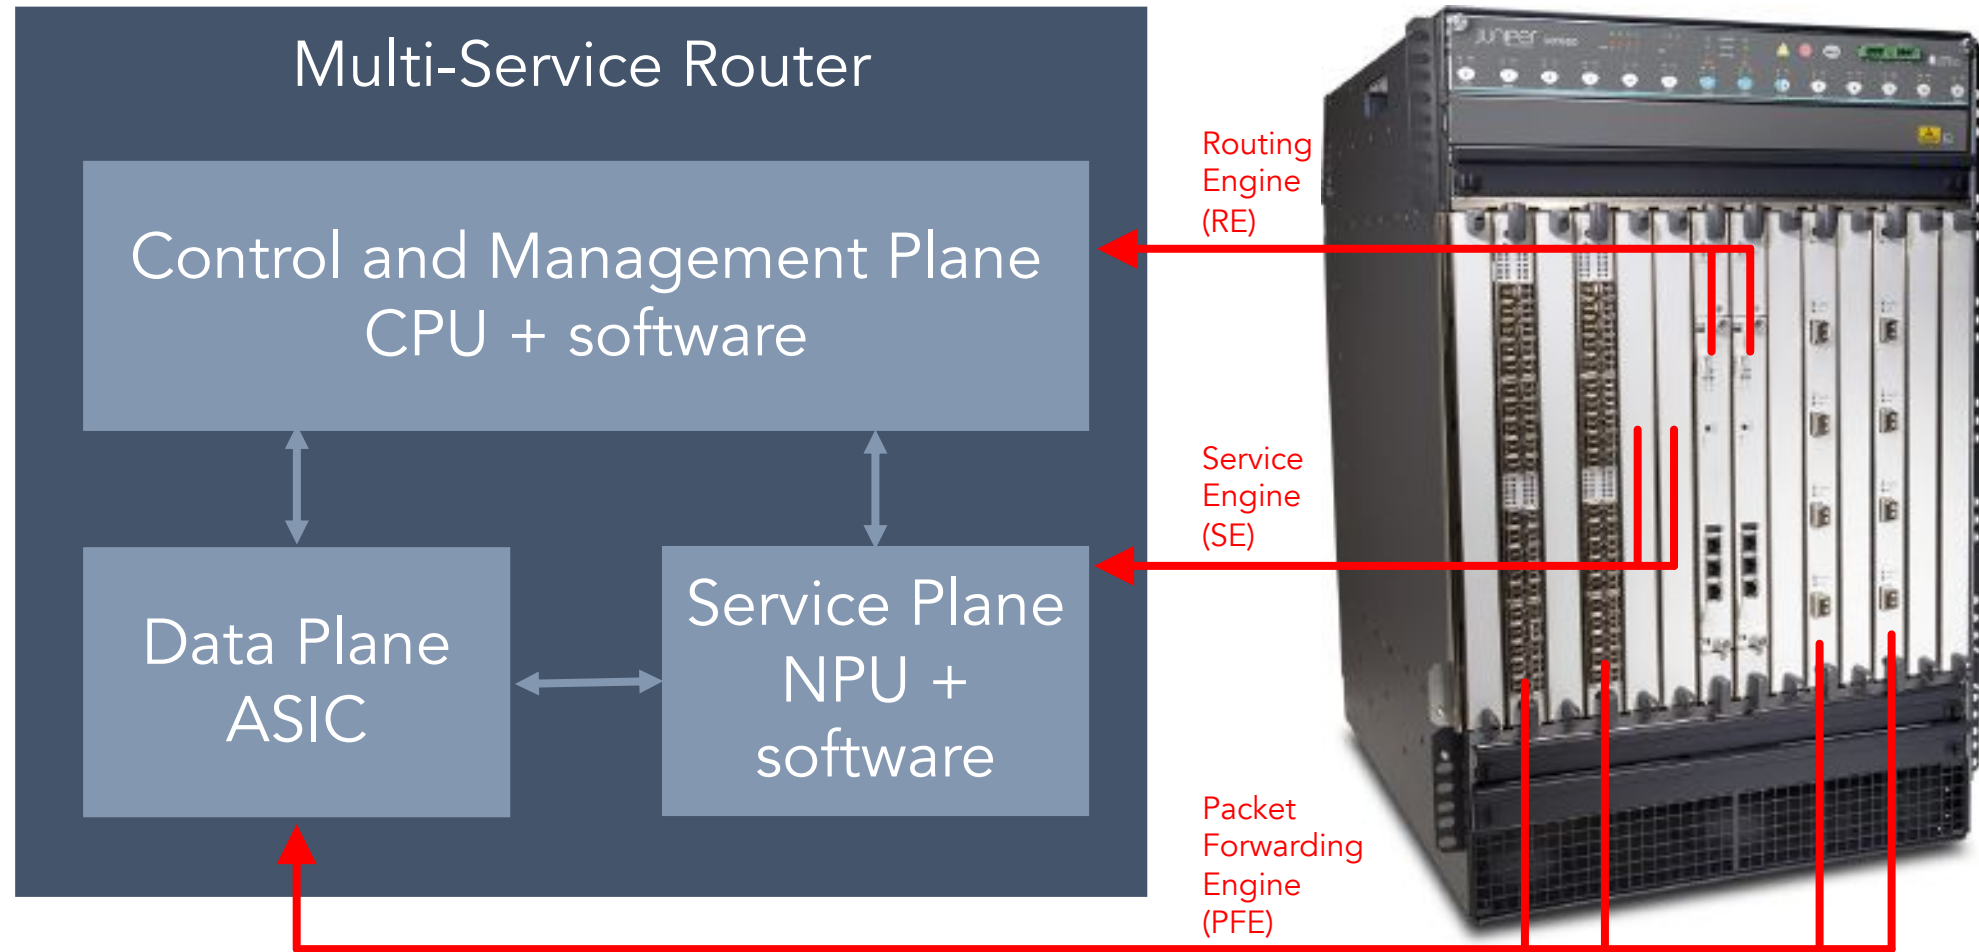
Service plane

Stateful flow processing
Networking NPU with real-time OS
Example: Application layer firewall (reconstruct TCP stream)

Data plane (aka forwarding plane)

Stateless packet-by-packet forwarding
Networking ASIC
Examples: Ethernet, IPv4, IPv6, MPLS, ..

Physical implementation of the four planes



Section 2

Fixed-function classical data plane

What does the data plane do? (A lot!)

- Parse received packets, compose ("de-parse") sent packets
- Layer 3 longest prefix match lookup of destination IP address
- Layer 2 exact match lookup of destination Ethernet address
- Exact match lookup of VLAN tag(s) and MPLS tag(s)
- Lookup of source address for Ethernet address learning and IP reverse path validation
- Handle congestion: buffer packets and implement drop policies
- Implement Quality of Service (QoS): scheduling and packet marking
- Implement (Non-) Equal Cost Multipath (N)ECMP: compute hashes, balance flows
- Apply policies, also known as Access Control Lists (ACLs)
- Track metrics (counters, thresholds, ...), generate in-line telemetry
- Make copies for multicast and mirroring
- Interface with the control plane (punt control packets, program tables, ...)
- ... etc. etc. etc. ... (*lots more*)

Fixed pipeline data plane

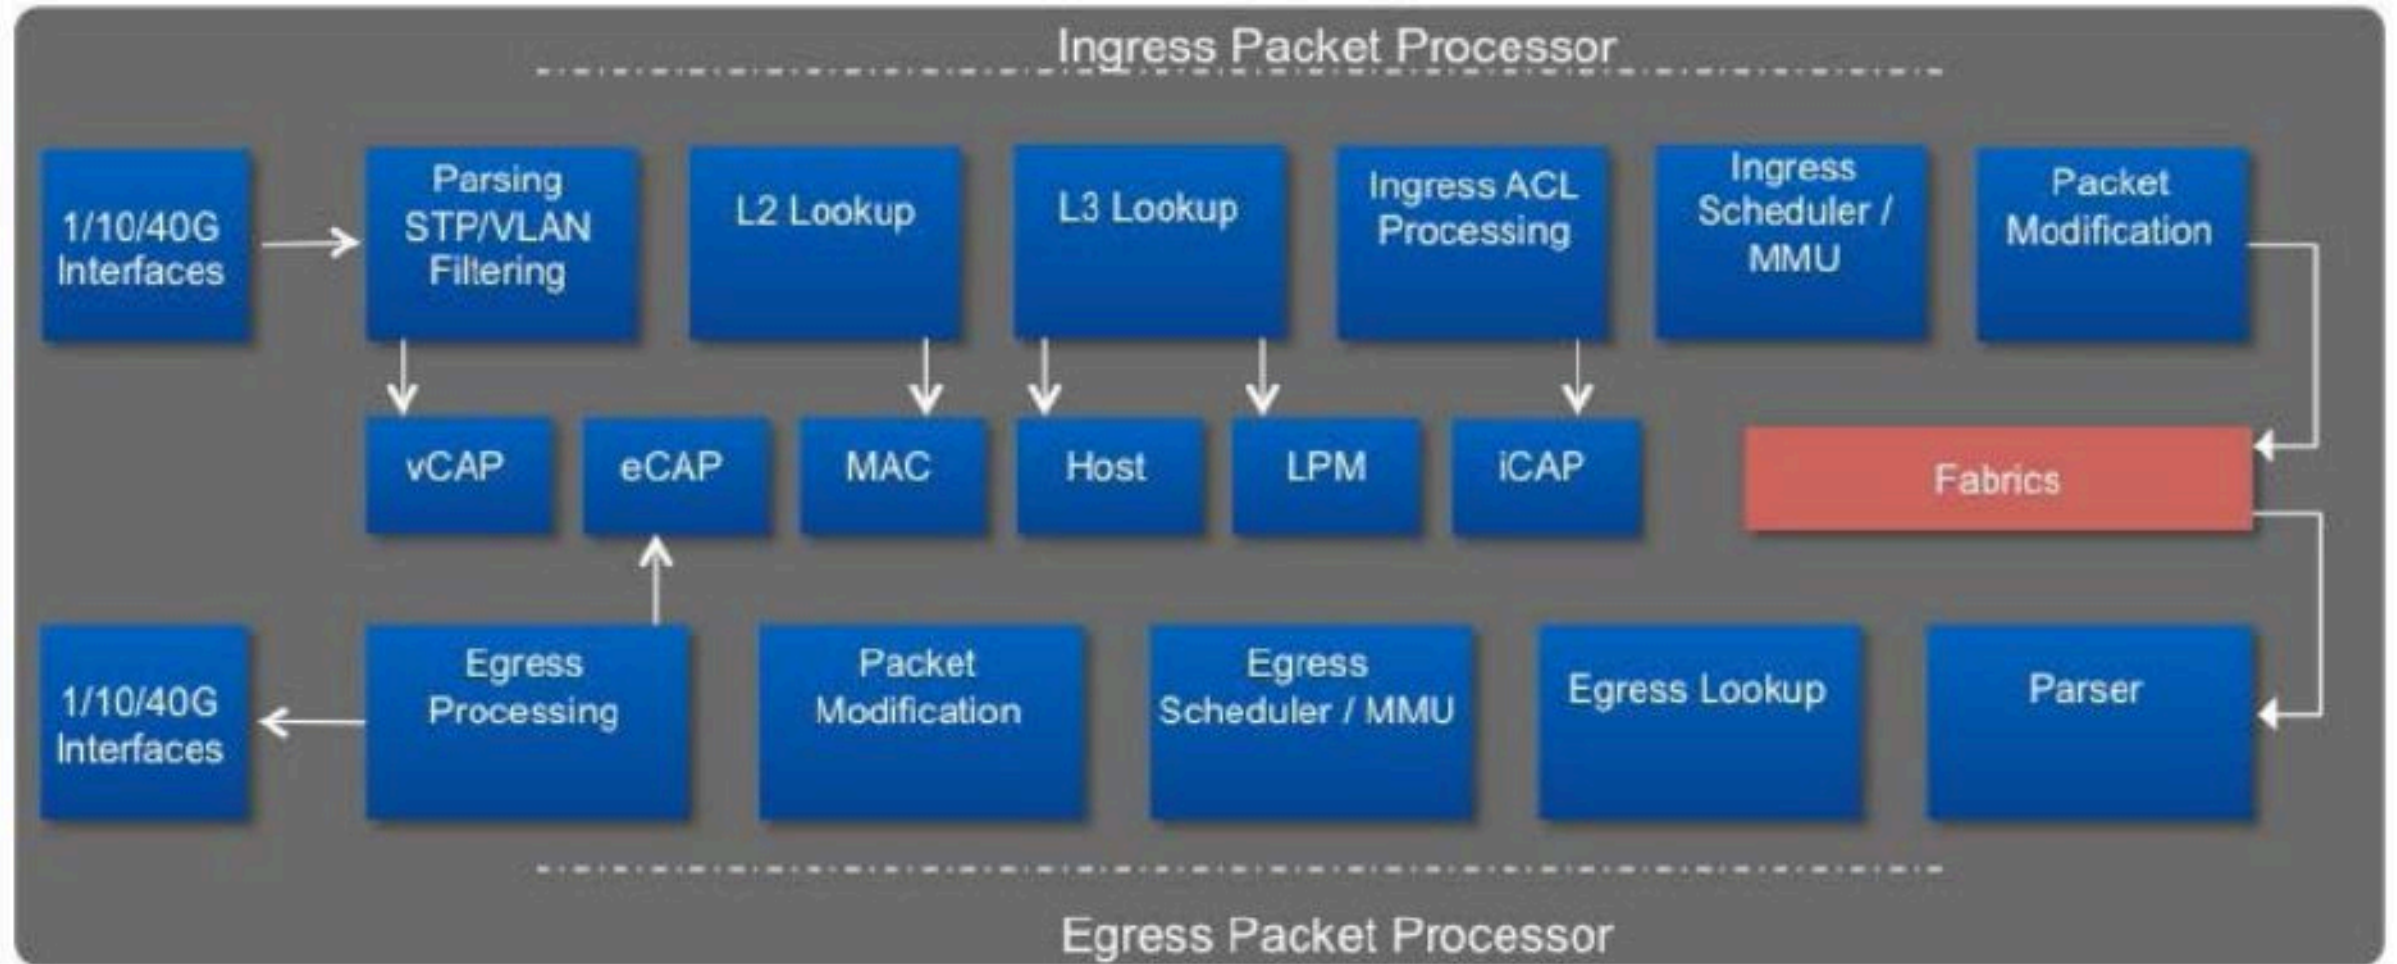
Most things are fixed in hardware and cannot be changed:

- The processing steps for each packet are fixed in hardware
- Fixed set of encapsulations (e.g. Ethernet, IPv4, IPv6, GRE, VXLAN)
- Fixed sequence of lookup tables (e.g. VLAN, Ethernet, IP, ACLs, ...)
- Fixed set of features (e.g. multicast, ECMP, telemetry, ...)

However, not everything is fixed:

- Dynamically populate contents of tables (e.g. routes, ACLs, ...)
- Dynamically configure certain features (e.g. speed of interface 10/40)
- Confusingly, this is sometimes called "the control plane programs the data plane"

Example fixed forwarding pipeline



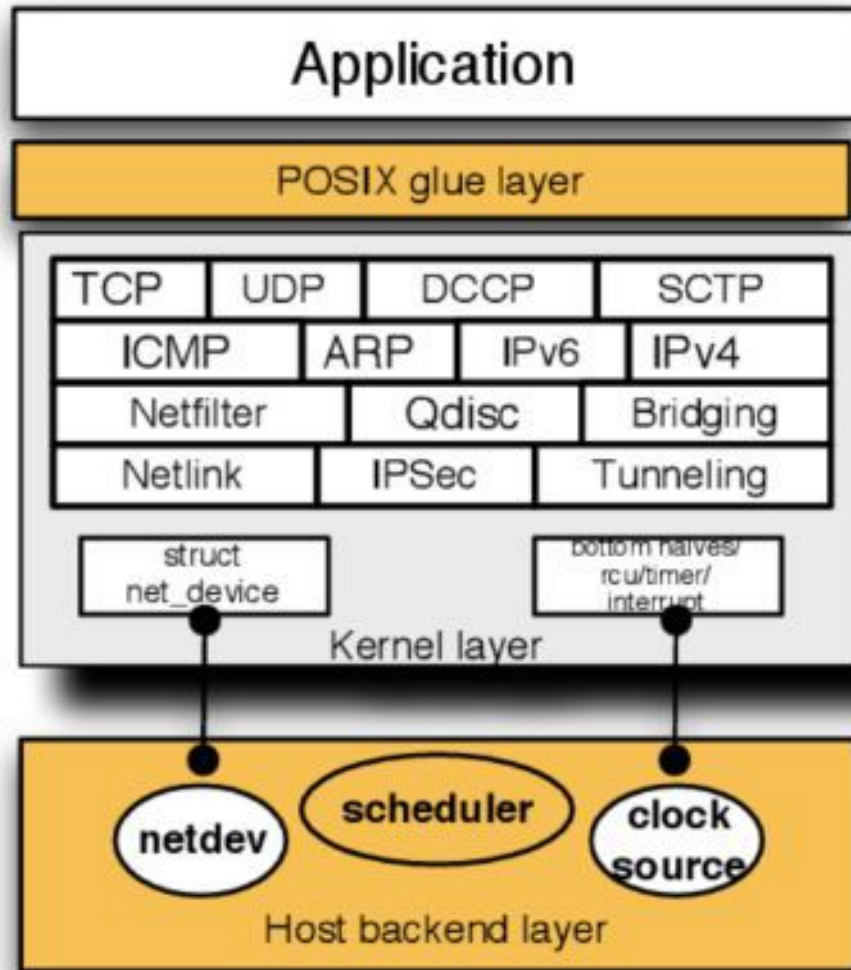
Fixed pipeline in hardware: Tomahawk 3



Broadcom Tomahawk 3

- ~ Dec 2017
- StrataXGS BCM56980
- 12.8 terabits per second (Tbps)
- 128x100Gbps / 64x200Gbps / 32x400Gbps
- ~ 8 billion packets per second (0.04 nsec/packet)
- 500ns or 300ns latency
- Relatively small forwarding tables
- Relatively shallow buffers
- Lean feature set
- **Data plane is not programmable**

Fixed pipeline in software: Linux kernel stack



Legacy Linux networking stack

- Linux can be configured to be a software switch
- Ethernet, IPv4, IPv6, VLANs, ACLs, ...
- Everything implemented in software (with option for hardware acceleration)
- Socket interface to send and receive packets (inet sockets, raw sockets)
- Socket interface to populate the forwarding tables and to configure options (netlink sockets)
- Note: recent versions of Linux have a programmable network stack (discussed later)

Data plane considerations

Optimize for (trade-off):

- Cost
- Throughput
- Latency
- Buffer depth
- Table sizes (routes, ACLs,...)
- Features (protocols, telemetry, ...)
- Flexibility (programmability)
- Power (cooling)
- Density

Constrained by:

- Die size limits (800 mm²)
- Process geometry
- Internal / external memory
- IO speed (SERDES)
- Power dissipation
- Development cost
- Device cost
- Time to market

Single-stage versus multi-stage data plane



	FB Backpack	Next Gen
Capacity	128x100GbE	32x400GbE / 128x100GbE
Front panel	128x QSFP28	32x QSFP-DD or OSFP
Size	8U Chassis	1U Fixed
# Switch Chips	12 x 3.2T	1 x 12.8T

Source: Facebook, OCP

75% Reduction in System Power, 85% reduction in System Cost *


Section 3

Programmable classical data plane


(Using P4 as the main example)

Programmable pipeline vs fixed pipeline

Traditional switch



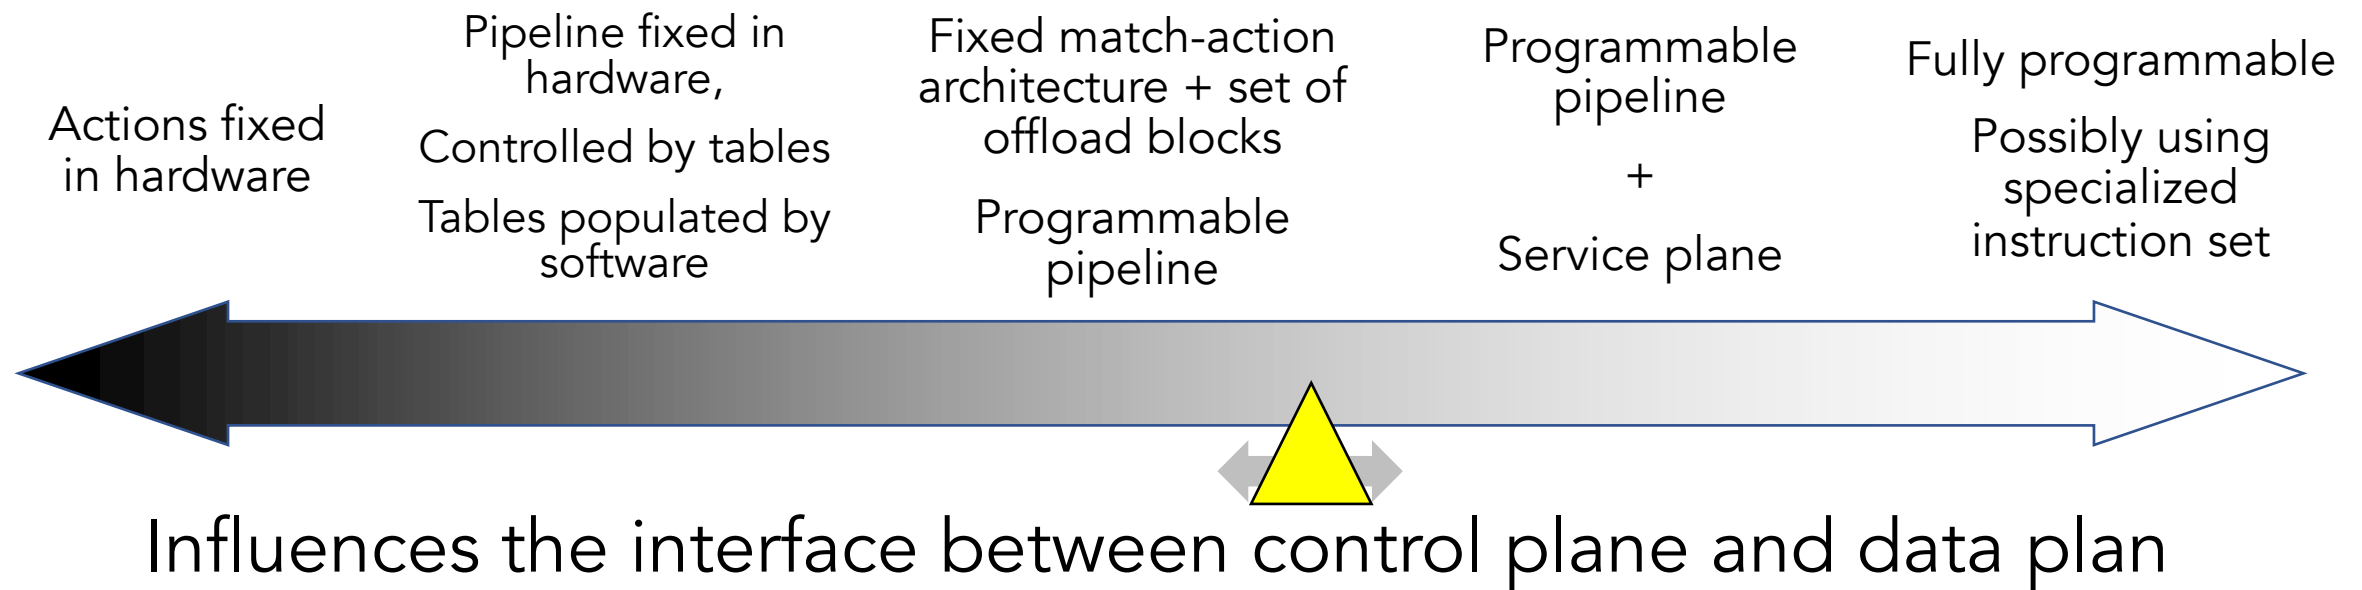
P4-defined switch




Controllable vs automated quantum nodes

Definitions from "*A Link Layer Protocol for Quantum Networks*":

- **Controllable node:** offers the possibility to perform controllable quantum operations as well as storing qubits. Specifically, these nodes enable decision making
- **Automated node:** are typically only timing controlled, i.e. they perform the same preprogrammed action in each time step.



Parser, match-action tables, de-parser







Advantages of a programmable pipeline

- Ability to make the data plane do what you want it to do
- Ability to define new data plane protocols
- Remove unused data plane protocols (reduced complexity)
- Flexible and efficient use of table resources
- Greater visibility (new telemetry protocols)
- Software-style of development
- Keep your own ideas

The data plane for quantum networks is not yet well-defined!
So we absolutely need a programmable data plane

Architectural diversity of P4 targets

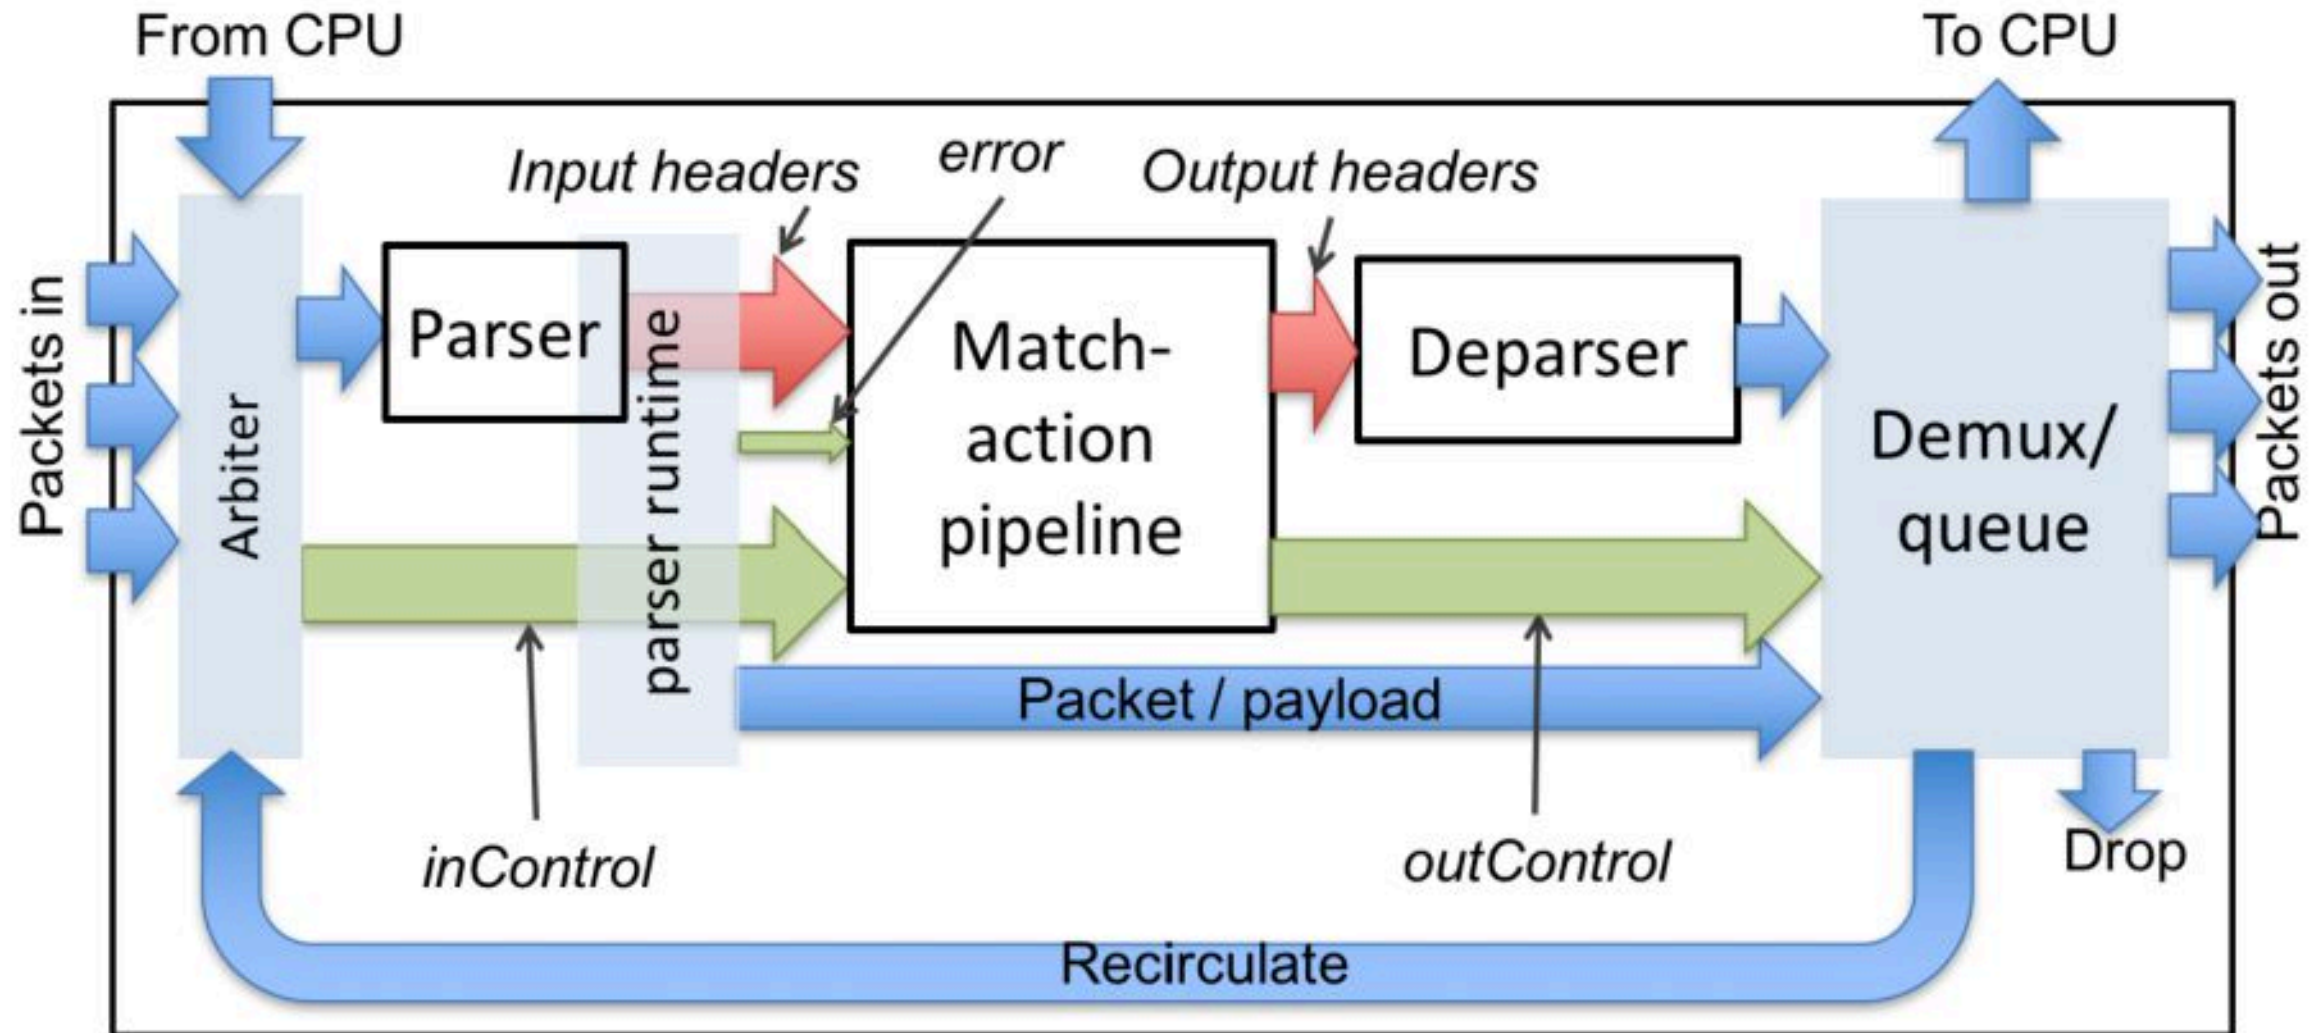


P4 architecture definition and externs library

Term	Explanation
P4 Target	An embodiment of a specific hardware implementation
P4 Architecture	Provides an interface to program a target via some set of P4-programmable components, externs, fixed components

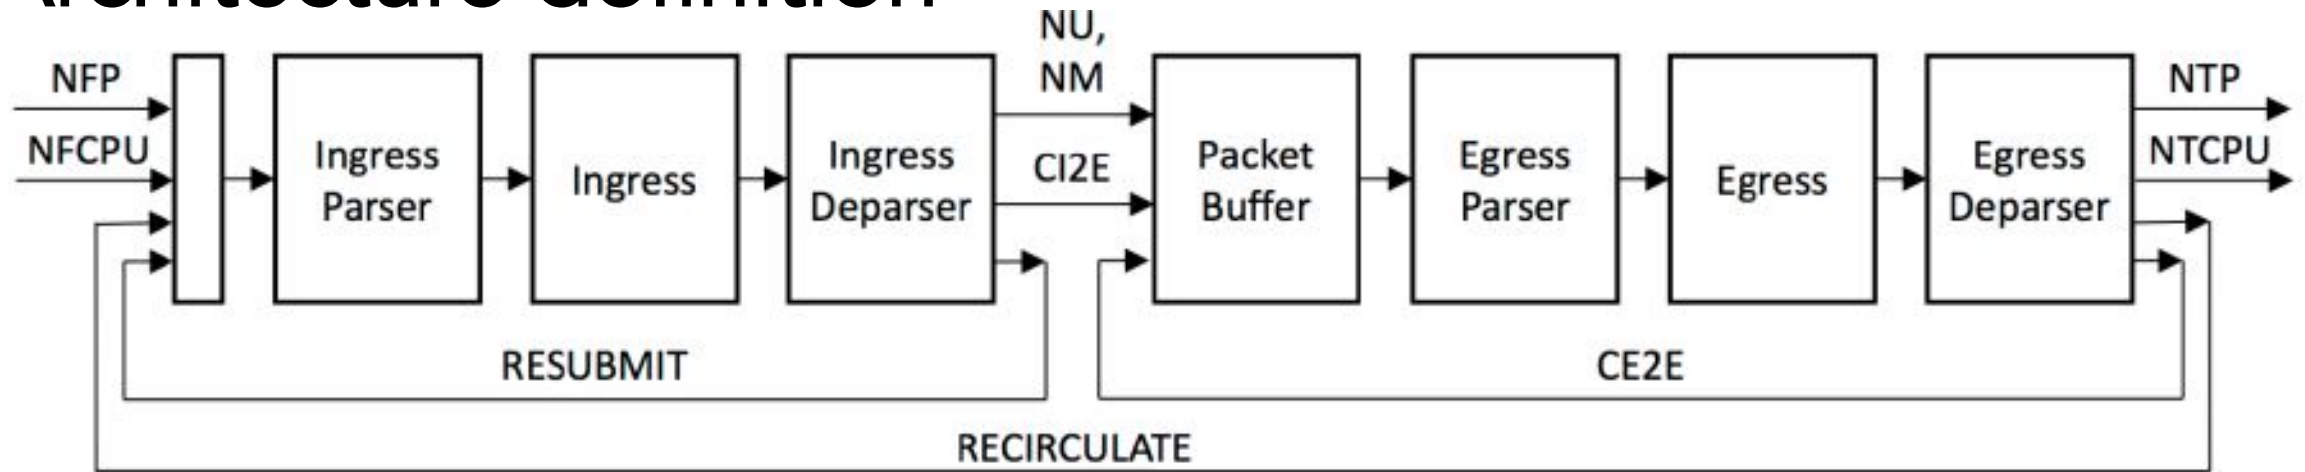


Very simple switch architecture



Portable Switch Architecture (PSA)

Architecture definition



Externs library

Packet Replication Engine	Buffering and Queueing Engine	Hashes	Checksums	Counters	Meters	Registers
		Random	Action Profile	Action Selector	Timestamp	Packet Digest

Portable Switch Architecture (PSA)

Architecture definition



Externs library

Packet Replication Engine	Buffering and Queueing Engine	Hashes	Checksums	Counters	Meters	Registers
		Random	Action profile	Action Selector	Timestamp	Packet Digest

?

Example P4 header definitions

```
header Ethernet_h {  
    EthernetAddress dstAddr;  
    EthernetAddress srcAddr;  
    bit<16> etherType;  
}
```

```
header IPv4_h {  
    bit<4> version;  
    bit<4> ihl;  
    bit<8> diffserv;  
    bit<16> totalLen;  
    bit<16> identification;  
    bit<3> flags;  
    bit<13> fragOffset;  
    bit<8> ttl;  
    bit<8> protocol;  
    bit<16> hdrChecksum;  
    IPv4Address srcAddr;  
    IPv4Address dstAddr;  
}
```

Example P4 meta data definitions (structs)

Architecture defined meta-data

```
struct psa_egress_input_metadata_t {
    ClassOfService_t    class_of_service;
    PortId_t            egress_port;
    PSA_PacketPath_t    packet_path;
    EgressInstance_t    instance;
    Timestamp_t         egress_timestamp;
    ParserError_t       parser_error;
}

struct psa_egress_output_metadata_t {
    bool                clone;
    CloneSessionId_t    clone_session_id;
    bool                drop;
}
```

Program defined meta-data

```
struct Parsed_packet {
    Ethernet_h ethernet;
    IPv4_h      ip;
}
```

Example P4 parser (state machine)

```
parser TopParser(packet_in b, out Parsed_packet p) {  
  
    state start {  
        b.extract(p.ethernet);  
        transition select(p.ethernet.etherType) {  
            0x0800: parse_ipv4;  
        }  
    }  
  
    state parse_ipv4 {  
        b.extract(p.ip);  
        verify(p.ip.version == 4w4, error.Ipv4IncorrectVersion);  
        verify(p.ip.ihl == 4w5, error.Ipv4OptionsNotSupported);  
        transition accept;  
    }  
}
```

Example P4 control block (match-action tables)

```
control TopPipe(inout Parsed_packet headers,  
                in error parseError,  
                in InControl inCtrl,  
                out OutControl outCtrl) {  
...  
    action Set_nhop(IPv4Address ipv4_dest, PortId port) {  
        nextHop = ipv4_dest;  
        headers.ip.ttl = headers.ip.ttl - 1;  
        outCtrl.outputPort = port;  
    }  
...  
    table ipv4_match {  
        key = { headers.ip.dstAddr: lpm; } // longest-prefix match  
        actions = {  
            Drop_action;  
            Set_nhop;  
        }  
        size = 1024;  
        default_action = Drop_action;  
    }
```

Example P4 control block (apply section)

```
...
    apply {
...
        // Lookup destination IPv4 address in IPv4 route table to determine next-hop
        ipv4_match.apply();
        if (outCtrl.outputPort == DROP_PORT) return;

        // Lookup remaining TTL in TTL table to decide whether to punt to CPU
        check_ttl.apply();
        if (outCtrl.outputPort == CPU_OUT_PORT) return;


        // Lookup next-hop IP address in dmac table to determine destination MAC address
        dmac.apply();
        if (outCtrl.outputPort == DROP_PORT) return;

        // Lookup outgoing port in smac table to determine source MAC address
        smac.apply();
    }
}
```

Programmable pipeline hardware (P4)

Barefoot Tofino 2 (Now Intel)

- ~ Jan 2019
- 12.8 Tbps
- Rich feature set
- Programmable using P4 high-level language
<https://p4.org>



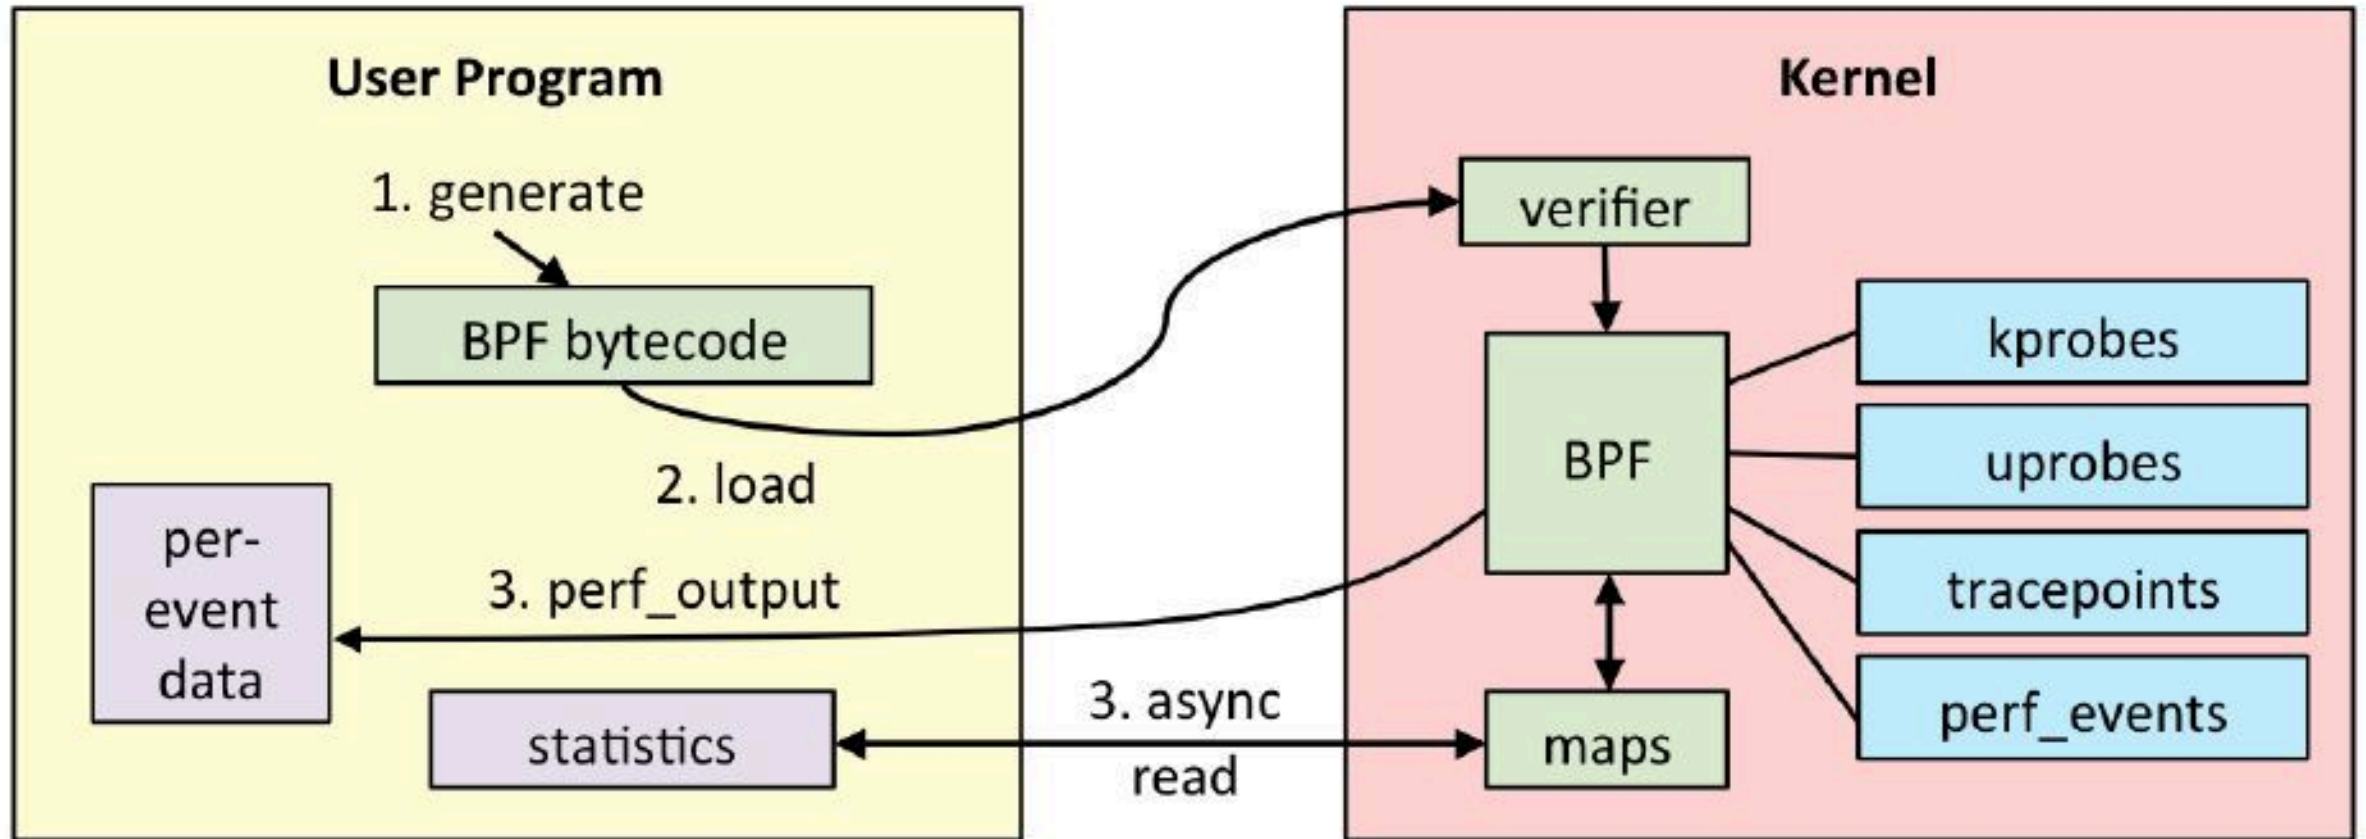
Programmable pipeline hardware (vendor-specific)

Broadcom Trident 4

- ~ June 2019
- StrataXGS BCM56880
- 12.8 terabits per second (Tbps)
- Rich feature set
- Programmable using NPL high-level language
<https://nplang.org/>



Programmable pipeline in software: Linux eBPF




Section 4


Control plane - data data interface

What interface are we talking about?

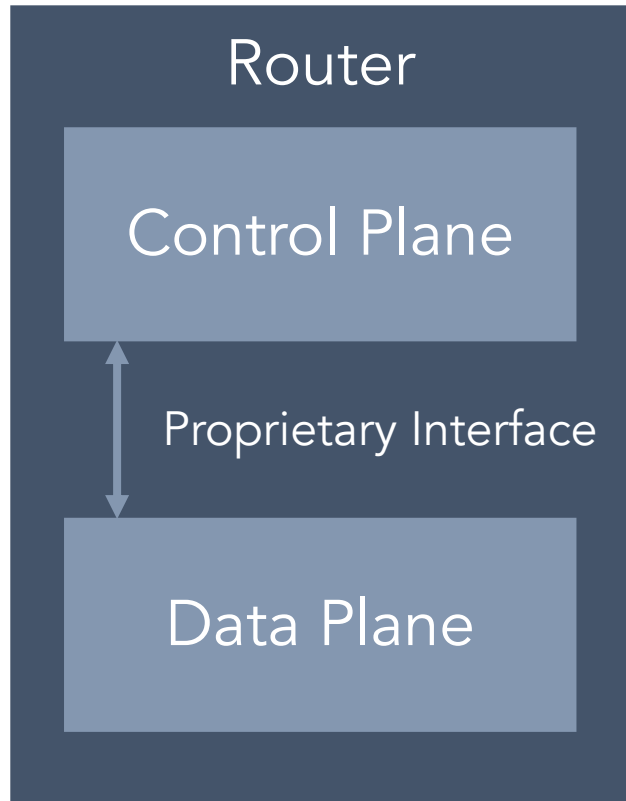
Traditional switch



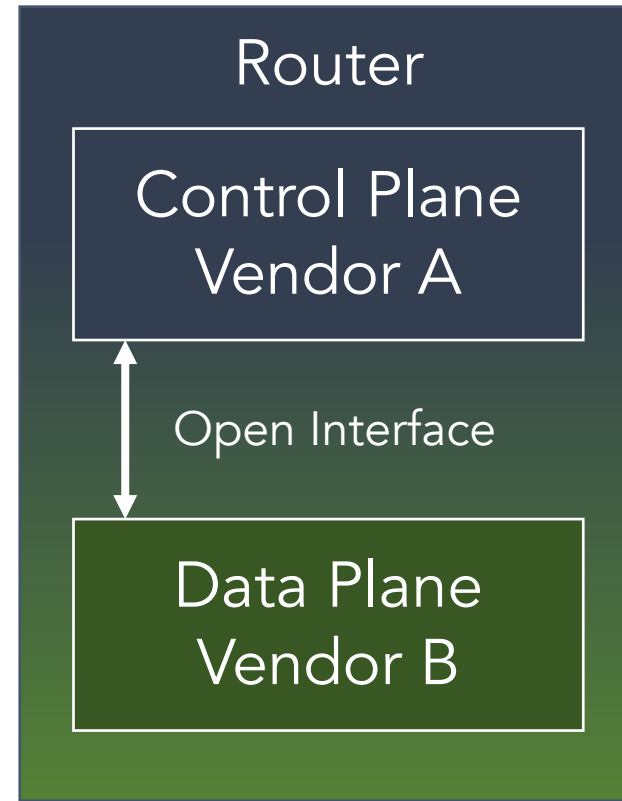

P4-defined switch



Disaggregation and open interfaces



Monolithic closed
single-vendor
router

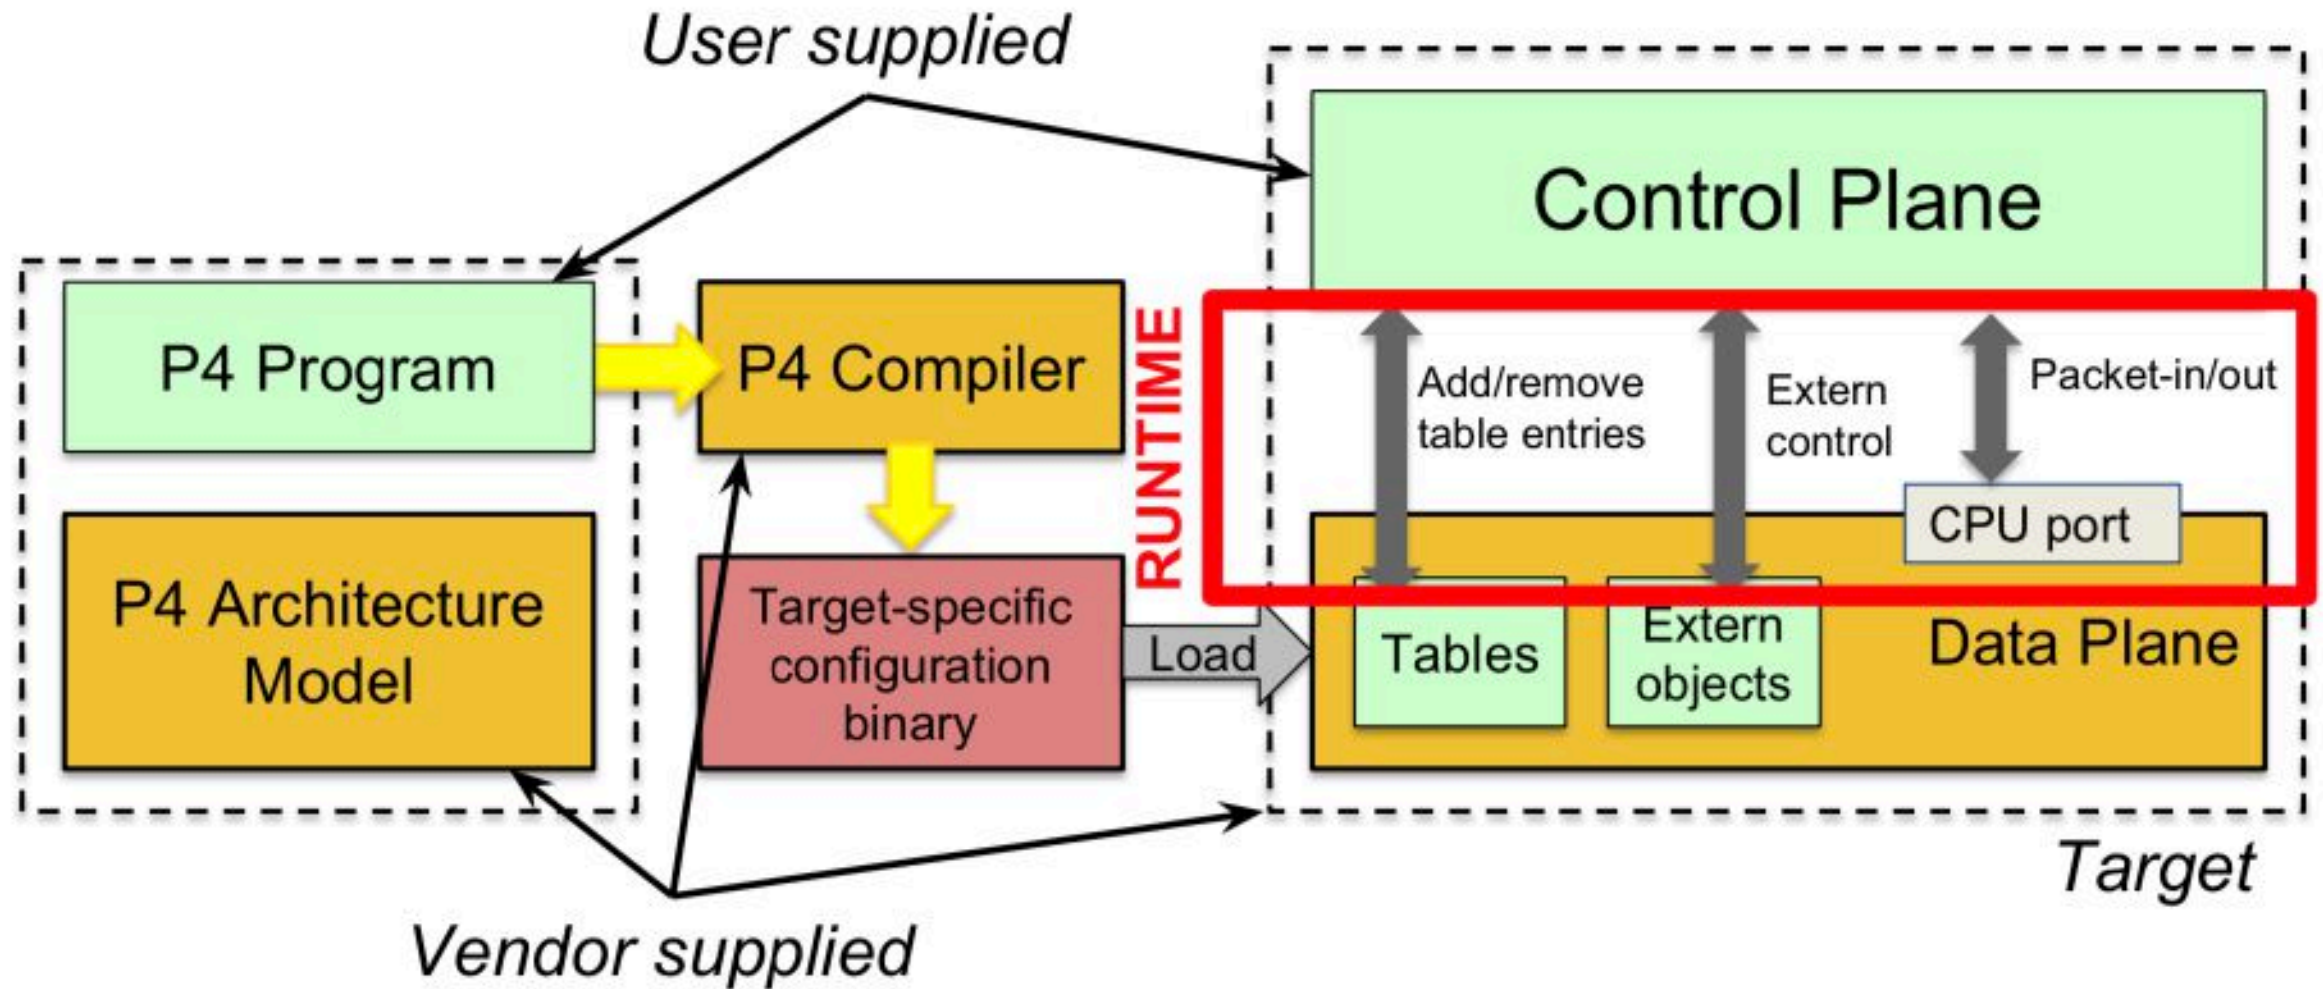


Disaggregated open
multi-vendor
router

"Standard" control plane - data plane interfaces



P4 runtime project

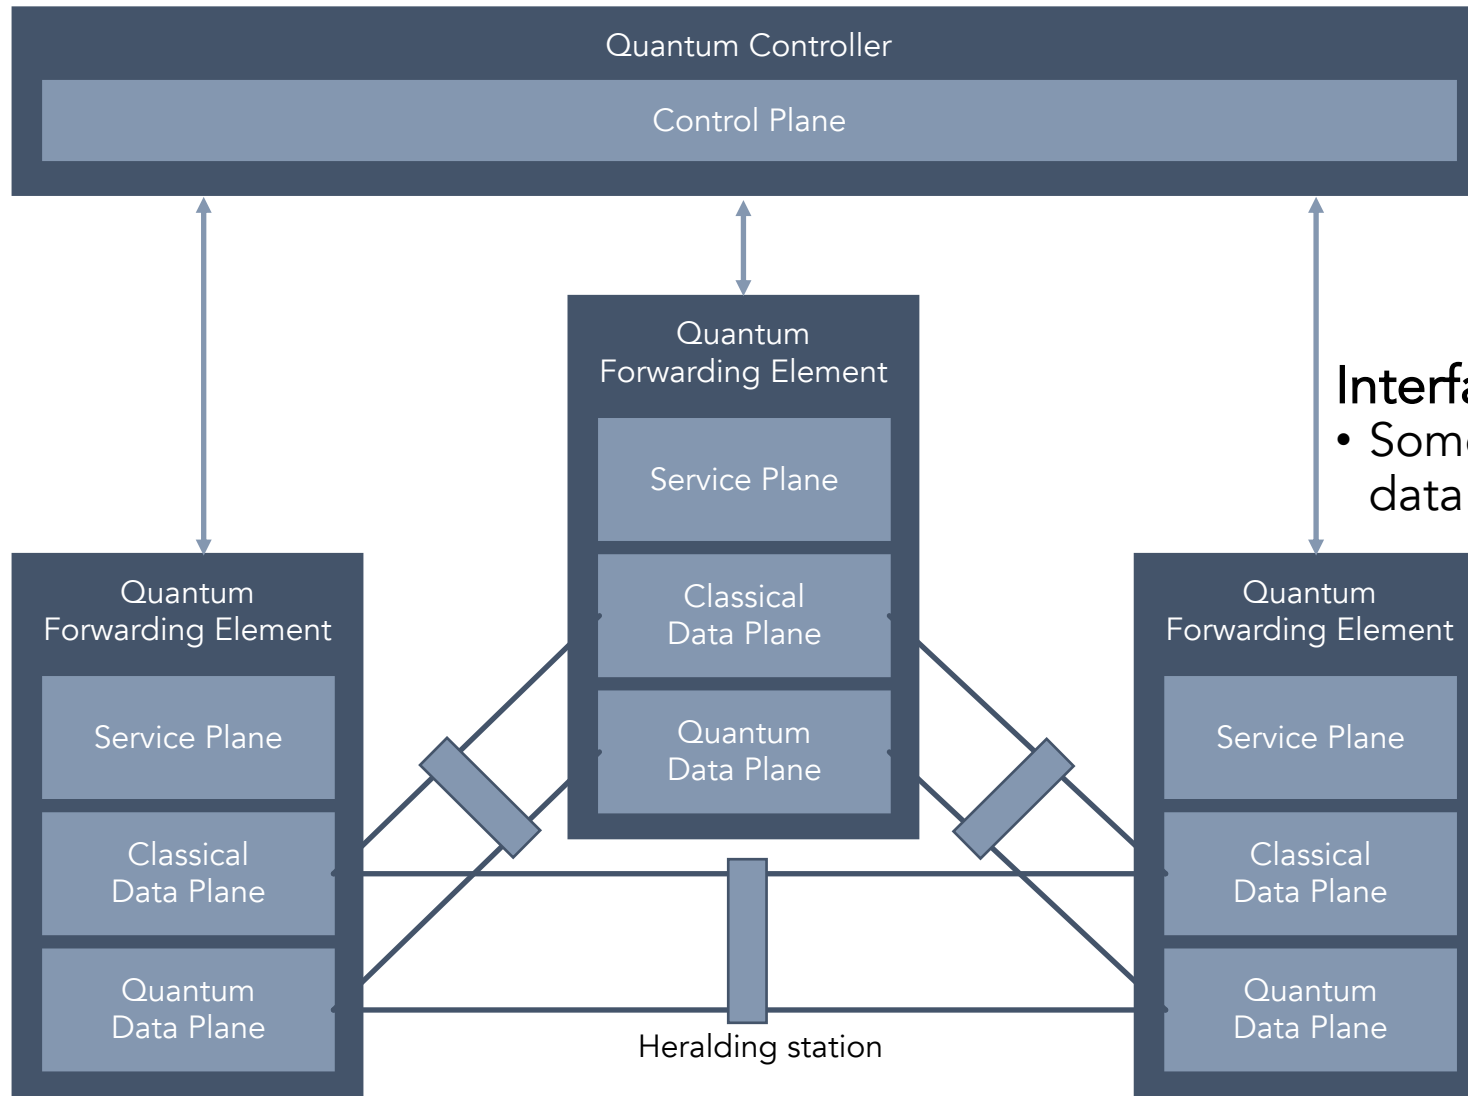


Section 5

Programmable quantum data plane

Quantum planes (centralized control plane)

Helpful to think of RWA in OTN networks as an analogy



Aware of end-to-end network service

- Aware of full traffic demand matrix (rate)
- Aware of required fidelity
- Aware of technology model
- Computes optimal paths and schedules
- Translates to network primitives


Interface

- Some protocol to program primitives into data plane

Implements local behavior

- Expressed in simple primitives
- Two data planes: classical and quantum
- Data plane will probably be modeled using state machines and scheduled timeslots to drive entanglement attempts and local operations (e.g. swap attempts)
- Something akin to "service plane" probably needed to implement stateful datalink protocols

Quantum planes (distributed control plane)



Key questions for the quantum data plane

- What are the right abstractions for the quantum data plane?
 - Abstract enough to abstract away the details of the underlying technology
 - Rich enough to expose all capabilities of the underlying technology
- What service primitives does the quantum data plane offer?
 - Generate point-to-point entangled Bell Pairs, swaps, error correction, distillation, resource discovery and management (e.g. number of memory qubits and characteristics)
- What is the interface between the quantum control plane and the quantum data plane?
- How do we make the quantum data plane programmable?
- Can we re-use / generalize existing mechanisms (e.g. P4) for the quantum data plane API / programming?

Some additional interesting observations


- Qubits do not have headers nor payload
 - Protocol “headers” must be in associated classical messages
 - But received qubits do have implicit meta-data: incoming port and timeslot
- Qubits can only be stored for a very short time (decoherence)
 - The network layer protocol needs to optimize for using entanglement as soon as possible
 - Any involvement of the control plane is out of the question (local and datalink operations only)
- There are many widely diverging candidate technologies for the quantum data plane
 - Need a very strong abstraction layer for the data plane
- Quantum data plane protocols are very stateful (as opposed to stateless IP/Ethernet)
 - Match-action tables and registers are probably too simple
 - Probably need an NPU akin to the service layer to process datalink protocols. “Punt to service”
- The control plane must be aware of the distributed nature of Bell Pairs and the physical model of the technology: noise models, loss models, etc.
 - But the data plane can probably get away without know any of this (analogy with OTN)
- The quantum data plane is inherently uni-directional
 - Quantum repeaters only send qubits.
 - Heralding stations only receive qubits.
 - Entanglement success is reported in a classical message.

Primitive operations in the quantum data plane

- Classical data plane
 - Re-use existing classical architecture (parser, match-action tables, de-parser)
 - Probably need an NPU to implement datalink protocols (akin to service plane)
- Quantum data plane
 - Primitives to send and receive qubits for local entanglement attempts
 - Primitives for local operations
 - What is the right level of abstraction for local operations?
 - Unitary gates?
 - Or higher level abstractions such as swap attempts, distillation attempts, etc.
- Where on the spectrum between controllable and automated are the quantum nodes?
 - Routers vs repeaters vs heralding stations?

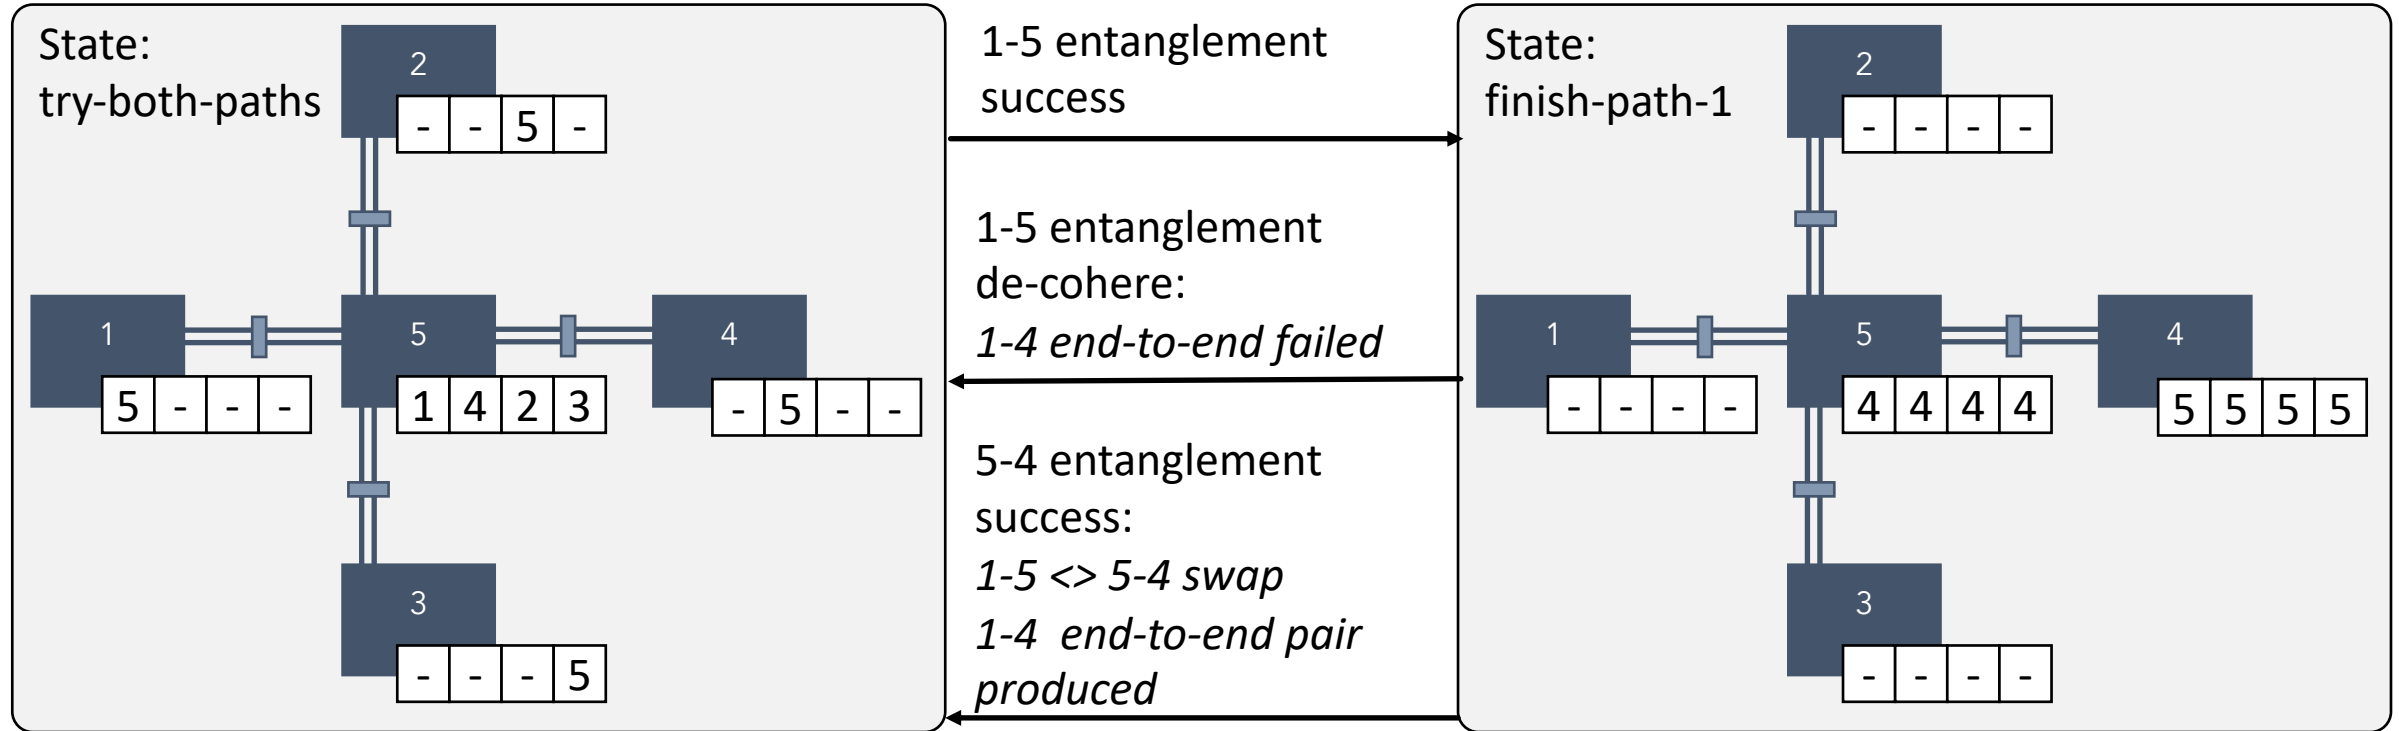
Some initial thoughts on a possible abstraction

Timeslot schedules (similar to GMPLS + TDM)



A possible optimization to minimize lost opportunities due to decoherence

Timeslot schedules and state machines



P4 seems promising for quantum data plane

- P4 is extremely flexible extensible
 - We probably can use P4 for the quantum data plane with no or little language syntax changes
- Use concept of "architecture packages" to represent unique characteristics of quantum router hardware pipeline
 - "Programmable Quantum Switch Architecture" equivalent of PSA
- Concept of "externals" to represent unique quantum hardware components
 - Swap attempts, distillation attempts, move to/from storage qubit, decoherence events, possibly individual gates, ...
- Represent received qubit as empty packet with only implicit meta-data (port and timeslot)

Thank you.