



Protocol Design Considerations

Bruno Rijsman

Version 1, Presented at QuTech (Delft) on 22 November 2019

A few true stories...

Designing correct protocols is hard

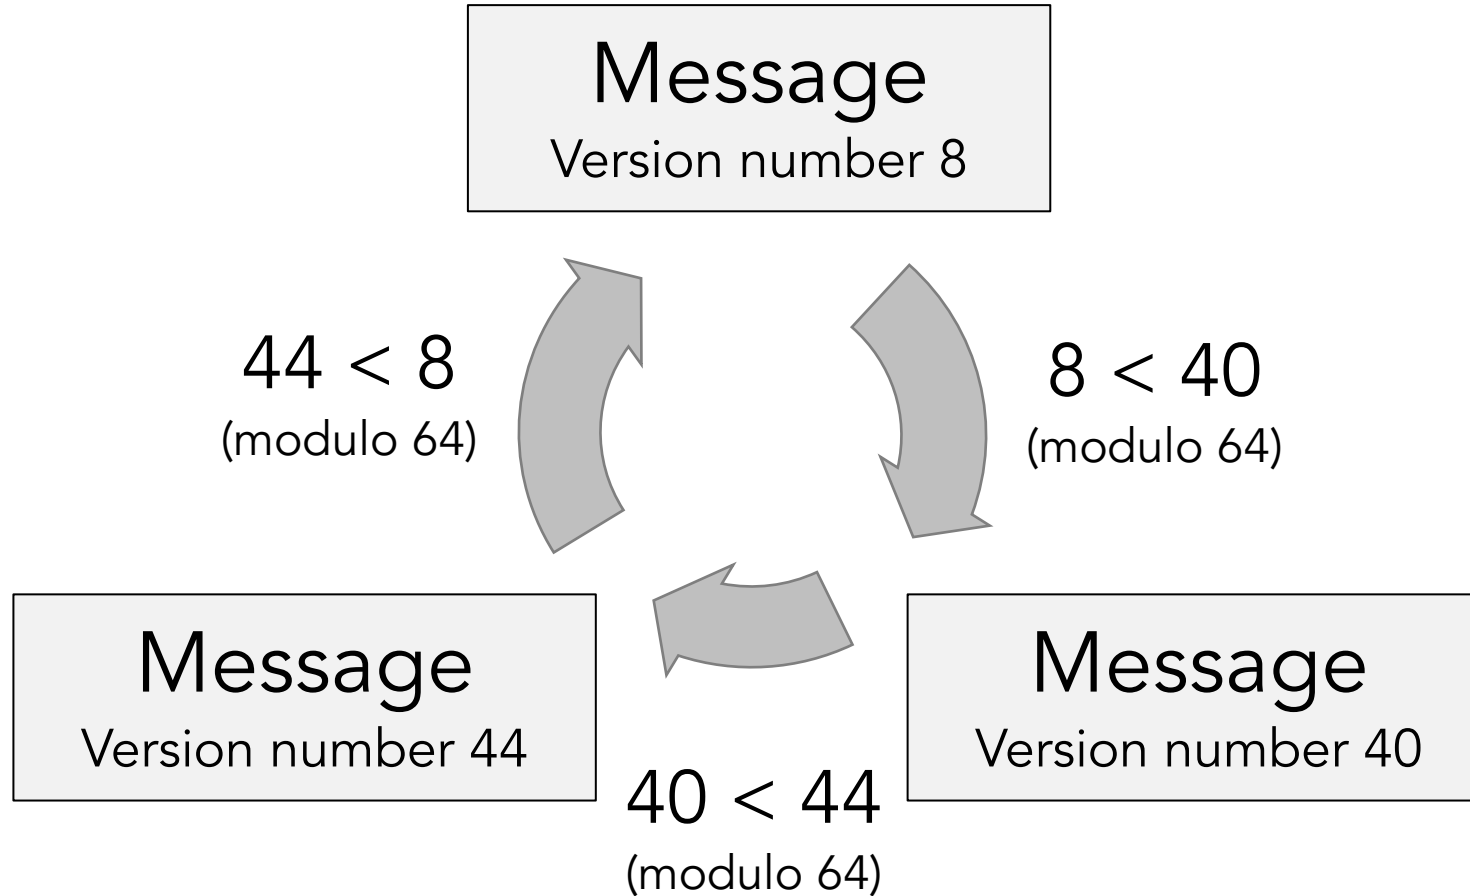
On 27 October 1980, *all* routers in the entire Internet had to be rebooted *simultaneously* to recover from a protocol design weakness.

Version numbers to pick newest message



Note: this mechanism is still used today, for example in OSPF and ISIS link state flooding

Modulo 64 comparison creates "version loop"



Version field was unsigned 6-bit integer, so math is modulo 64

Competing message versions kept replacing each other forever, eating up all CPU

The versions were introduced to a double bit error, which was not detected because check-sums were disabled.

We have not quite yet learned our lessons

A similar problem exists in the BGP specification and is observed in the Internet *to this day*.

See BGP persistent route oscillations ([RFC3345](#))

Implementing protocols correctly is hard

A bug in a widely used BGP implementation allowed an attacker to advertise a BGP route that caused many BGP routers *on the other side of the world* to crash.

Sending a specially constructed 4-octet AS-path caused certain routers to crash
CVE-2014-3818 <https://securitytracker.com/id/1031009>

Protocol security is hard

A BGP vulnerability enabled hackers to rob a bitcoin bank.

<https://www.wired.com/2014/08/isp-bitcoin-theft/>


Planning for evolution

Multi-version, multi-feature, multi-vendor networks

Your protocol will evolve over time


- Typical versioning scheme:
 - Minor version bump = backwards compatible
 - Major version bump = not backwards compatible
 - In practice, this does not work!
- Examples of doing it wrong:
 - IPv4 to IPv6
 - OSPFv2 to OSPFv3
- Examples of doing it right
 - ISIS
 - BGP
 - Evolve features without breaking backwards compatibility
- If you don't design your protocol with the future in mind, every new feature will be a major version bump...

The problem with major version bumps



- It's extremely hard avoiding running different major versions in your network if major bumps happen frequently.
- Temporarily during upgrade. You cannot upgrade your devices all at once; it takes a lot of time.
- Or semi-permanently due to operational considerations (e.g. different vendors support different versions of the protocol). You want to be able to have multiple vendors in the same network.
- Your network will be partitioned. Having a partitioned network makes finishing the upgrade nearly impossible.

A better approach based on features

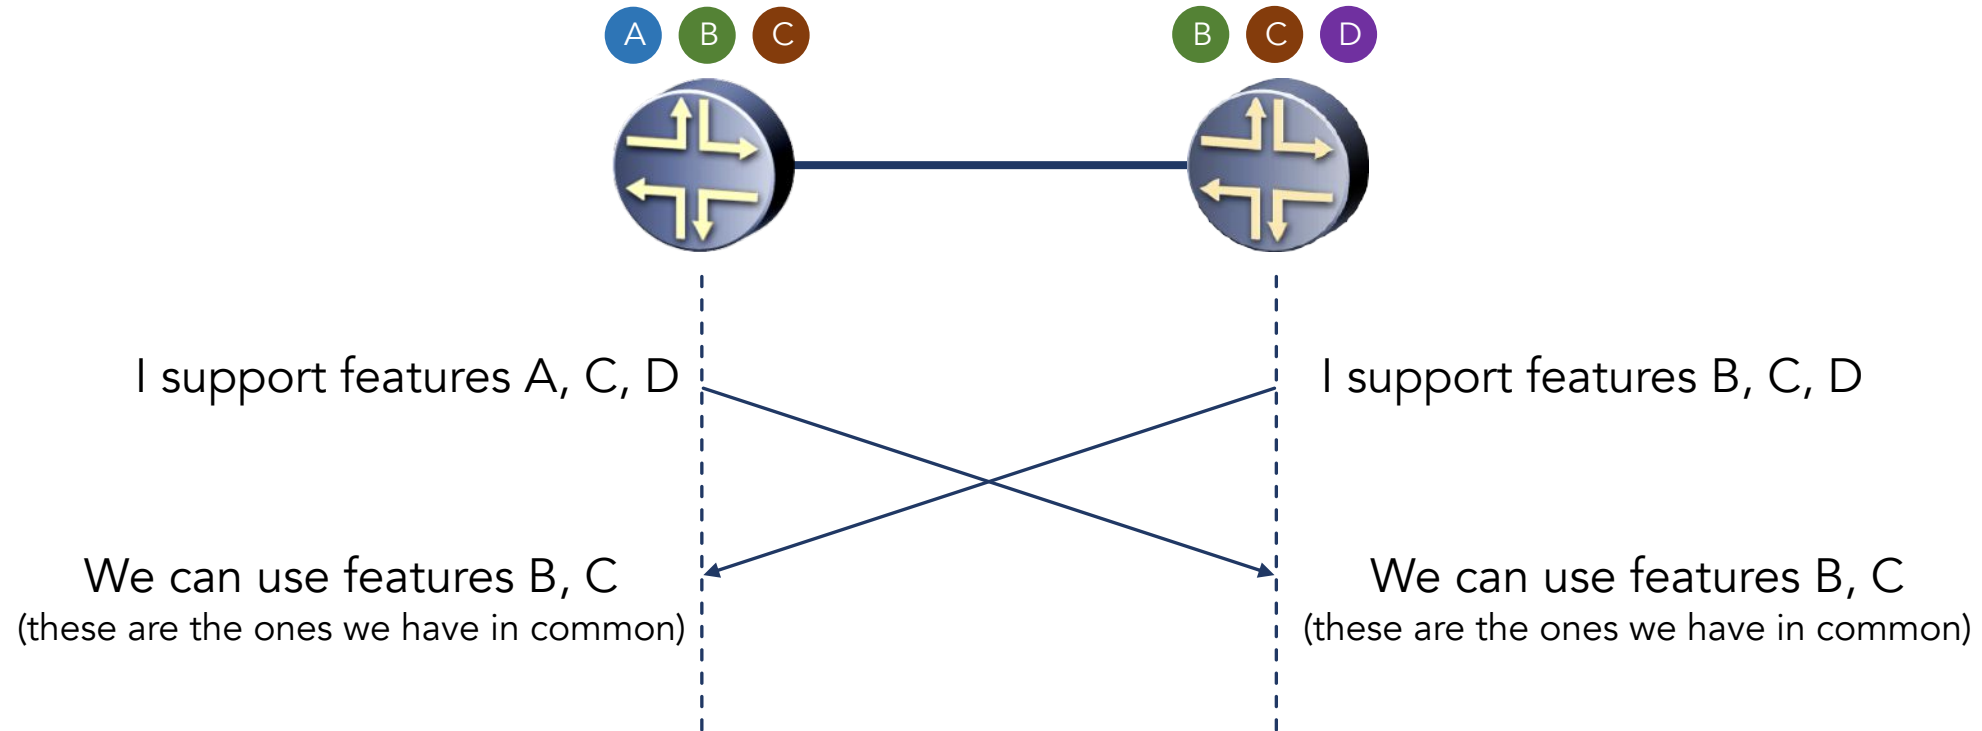


Protocol features
(aka **capabilities**)



- Protocol implementation that has feature A must interoperate with protocol implementation that does not have feature A.
- Different vendors and versions of code have different feature sets.
- Don't assume the feature set can be derived from the protocol version. **Allows protocol extensions to evolve and be deployed independently.**
- Don't assume the feature set is the same. Feature set A+B must interoperate with B+C.

Capability announcement ("negotiation")



(Things get a bit more complicated when more than 2 devices are involved, e.g. flooding)

Example: BGPv4 capabilities

Capability Codes

Reference

[\[RFC5492\]](#)

Available Formats



CSV

Range	Registration Procedures	Note
1-63	IETF Review	
64-127	First Come First Served	
128-255	Reserved for Private Use	IANA does not assign

Value	Description	Reference
0	Reserved	[RFC5492]
1	Multiprotocol Extensions for BGP-4	[RFC2858]
2	Route Refresh Capability for BGP-4	[RFC2918]
3	Outbound Route Filtering Capability	[RFC5291]
4	Multiple routes to a destination capability (deprecated)	[RFC8277]
5	Extended Next Hop Encoding	[RFC5549]
6	BGP Extended Message	[RFC8654]
7	BGPsec Capability	[RFC8205]
8	Multiple Labels Capability	[RFC8277]
9	BGP Role (TEMPORARY - registered 2018-03-29, extension registered 2019-03-18, expires 2020-03-29)	[draft-ietf-idr-bgp-open-policy]
10-63	Unassigned	
64	Graceful Restart Capability	[RFC4724]
65	Support for 4-octet AS number capability	[RFC6793]
66	Deprecated (2003-03-06)	
67	Support for Dynamic Capability (capability specific)	[draft-ietf-idr-dynamic-cap]
68	Multisession BGP Capability	[draft-ietf-idr-bgp-multisession]
69	ADD-PATH Capability	[RFC7911]
70	Enhanced Route Refresh Capability	[RFC7313]
71	Long-Lived Graceful Restart (LLGR) Capability	[draft-uttaro-idr-bgp-persistence]
72	Unassigned	
73	FQDN Capability	[draft-walton-bgp-hostname-capability]
74-127	Unassigned	
128-255	Reserved for Private Use	[RFC5492]

Number	Description	Reference	Registration Date
0	Reserved		
1	IP (IP version 4)		
2	IP6 (IP version 6)		
3	NSAP		
4	HCLOC (8-bit multihop)		
5	BBN 1822		
6	802 (includes all 802 media plus Ethernet "canonical format")		
7	E.163		
8	E.164 (SMDS, Frame Relay, ATM)		
9	F.80 (Telex)		
10	X.121 (X.25, Frame Relay)		
11	IPX		
12	Aspiracle		
13	Decnet IV		
14	Banyan Vines		
15	E.164 with NSAP format subaddress	[ATM Forum UNI 3.1, October 1995]	
16	DNS (Domain Name System)		
17	Distinguished Name	[Charles Lint]	
18	AS Number	[Charles Lint]	
19	XTP over IP version 4	[Mika Sau]	
20	XTP over IP version 6	[Mika Sau]	
21	XTP native mode XTP	[Mika Sau]	
22	Fibre Channel World-Wide Port Name	[Mark Rabe]	
23	Fibre Channel World-Wide Node Name	[Mark Rabe]	
24	GWID	[Sutra Hood]	
25	API for L2VPN information	[RFC4761]	
26	MPLS-TP Section Endpoint Identifier	[RFC7212]	
27	MPLS-TP LSP Endpoint Identifier	[RFC7212]	
28	MPLS-TP Pseudowire Endpoint Identifier	[RFC7212]	
29	MT IP: Multi-Topology IP version 4	[RFC7307]	
30	MT IPv6: Multi-Topology IP version 6	[RFC7307]	
31-16383	Unassigned		
16384	EGRP Common Service Family	[Donnie Savage]	2008-05-13
16385	EGRP IPv4 Service Family	[Donnie Savage]	2008-05-13
16386	EGRP IPv6 Service Family	[Donnie Savage]	2008-05-13
16387	LSRP Canonical Address Format (CAF)	[David Ward]	2008-11-12
16388	BGP-LS	[RFC7293]	2013-03-20
16389	48-bit MAC	[RFC7294]	2013-05-06
16390	64-bit MAC	[RFC7294]	2013-05-06
16391	CU	[RFC7391]	2013-09-25
16392	MAC/24	[RFC7391]	2013-09-25
16393	MAC/48	[RFC7391]	2013-09-25
16394	IPv6v4	[RFC7391]	2013-09-25
16395	RBridge Port ID	[RFC7391]	2013-09-25
16396	THILL Nickname	[RFC7391]	2013-09-25
16397	Universally Unique Identifier (UUID)	[Nischal Sheth]	2019-11-04
16398-65534	Unassigned		
65535	Reserved		


Range	Registration Procedures	Note
1-63	Standards Action	
64-127	First Come First Served	
128-240	Some recognized assignments below, others Reserved	
241-254	Reserved for Private Use	Not to be assigned

Value	Description	Reference
0	Reserved	[RFC4761]
1	Network Layer Reachability Information used for unicast forwarding	[RFC4761]
2	Network Layer Reachability Information used for multicast forwarding	[RFC4761]
3	Reserved	[RFC4761]
4	Network Layer Reachability Information (NLR) with MPLS Labels	[RFC6277]
5	MC-ASTNVR	[RFC6277]
6	Network Layer Reachability Information used for Dynamic Placement of Multi-Segment Pseudowires	[RFC6277]
7	Encapsulation SAFI	[RFC6277]
8	MC-ASTNVR	[RFC6277]
9-63	Unassigned	
64	Tunnel SAFI	[draft-ietf-idr-bgp-open-policy]
65	Virtual Private LAN Service (VPLS)	[RFC4434]
66	BGP MDI SAFI	[RFC6037]
67	BGP 4over6 SAFI	[RFC6247]
68	BGP 6over4 SAFI	[RFC6247]
69	Layer 1 VPN auto-discovery information	[RFC6190]
70	BGP EVPN	[RFC7432]
71	BGP-LS	[RFC7293]
72	BGP-LS-VPN	[RFC7293]
73	SR TE Policy SAFI	[draft-ietf-idr-bgp-open-policy]
74	SD-WAN Capabilities	[draft-ietf-idr-bgp-open-policy]
75-127	Unassigned	
128	MPLS-labeled VPN address	[RFC4365]
129	Multicast for BGP/MPLS IP Virtual Private Networks (VPLS)	[RFC4365]
130-131	Reserved	
132	Route Target constraints	[RFC6864]
133	IPv4 dissemination of flow specification rules	[RFC6864]
134	VPNv4 dissemination of flow specification rules	[RFC6864]
135-136	Reserved	
137	VPN auto-discovery	[draft-ietf-idr-bgp-open-policy]
140	Reserved	
141-246	Reserved	

Type Length Value (TLV) Encoding


- Makes it easy for old versions of the protocol to interoperate with new versions of the protocol.
- Also avoid bugs in the code by making it easier to write parsers.
- Used in almost all modern protocols (OSPF, ISIS, BGP, RSVP, ...)

Type Length Value (TLV) Encoding



Very easy to parse; less chance of bugs in parses.

Hierarchical TLVs



TLVs

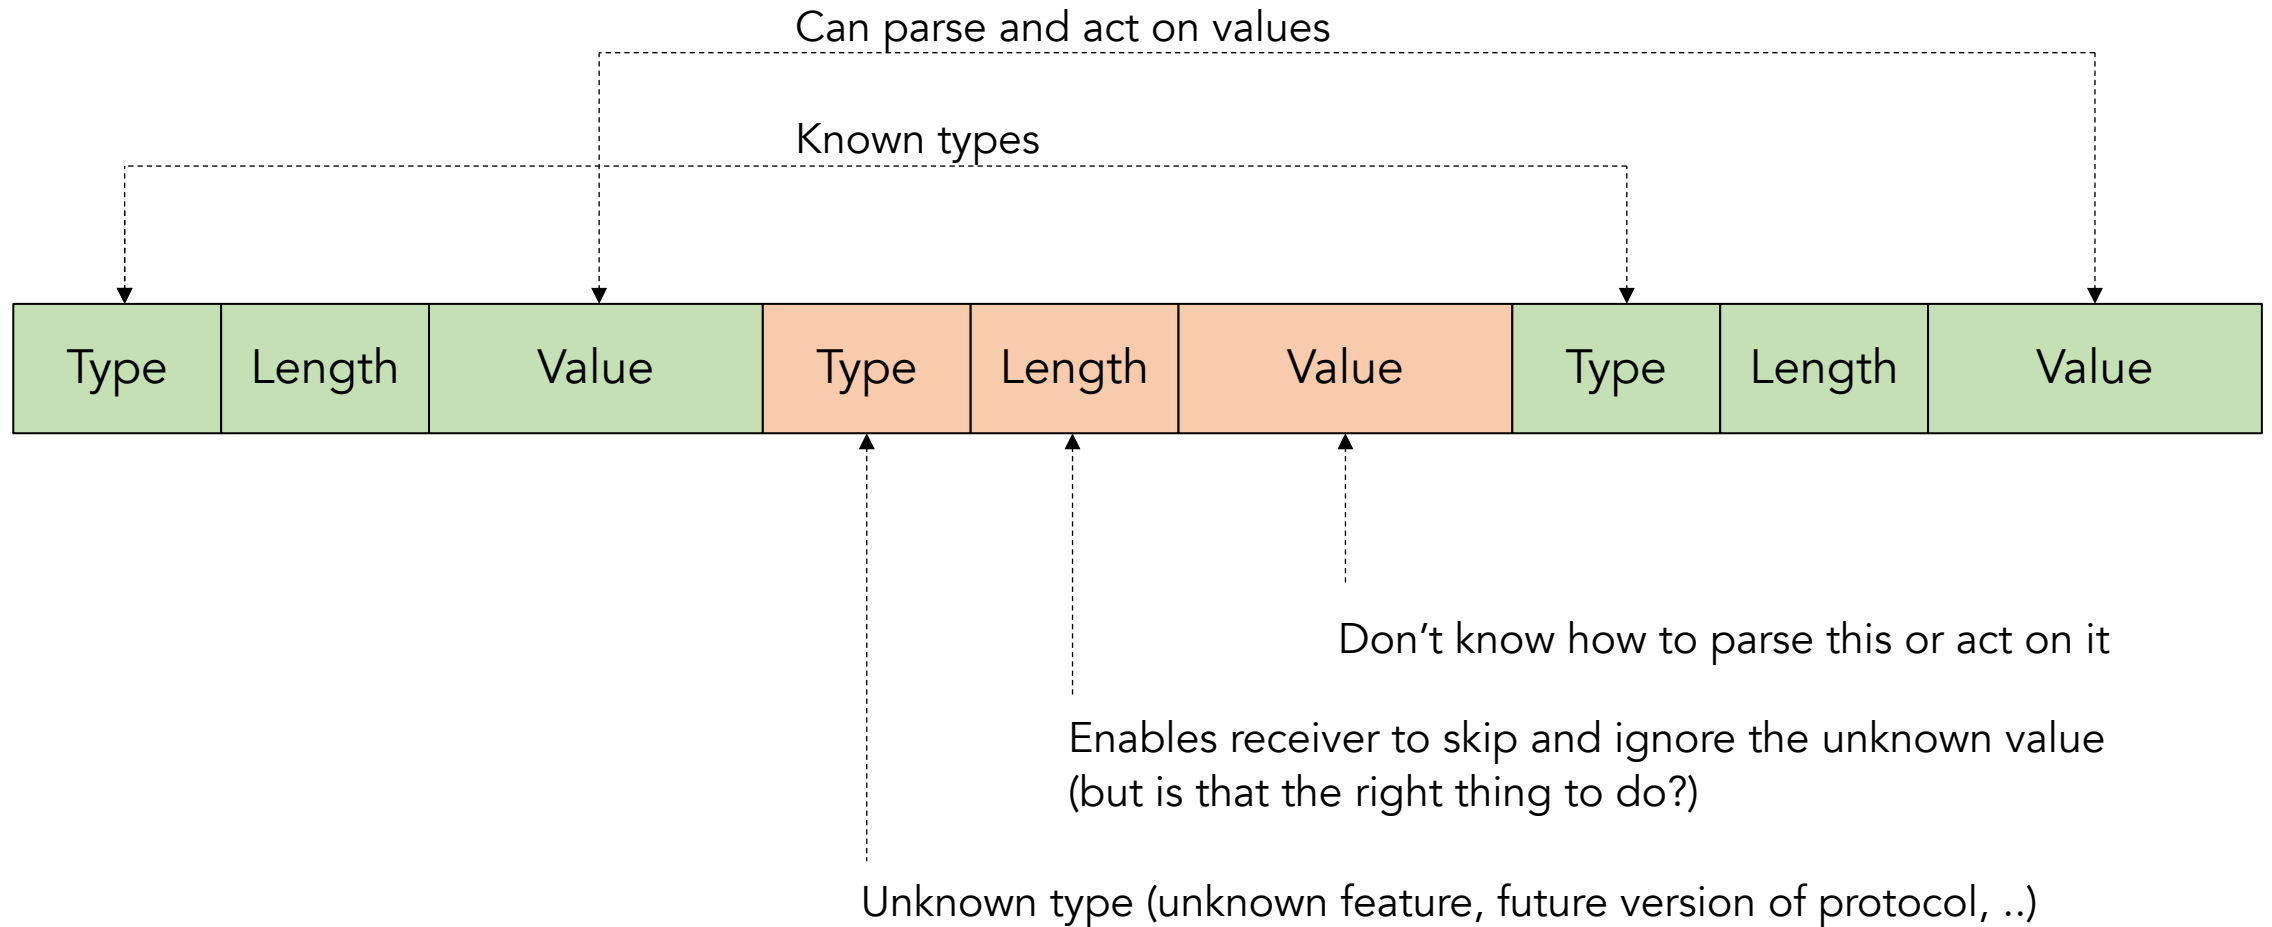


Sub-TLVs



Sub-Sub-TLVs

Easy to skip unknown TLVs



Example: IANA ISIS TLV code point registry

TLV Codepoints Registry

Registration Procedure(s)

Expert Review

Expert(s)

Chris Hopps, Hannes Gredler, Les Ginsberg

Reference

[\[RFC3563\]](#)[\[RFC6233\]](#)[\[RFC7356\]](#)

Note

IETF SHALL keep JTC1/SC6 informed of TLV codepoint values allocated, and JTC1/SC6 SHALL refer allocation requests arising within JTC1 constituencies to the IANA registry process.

Note

Codepoints greater than 255 can only be used in PDUs designated to support extended TLVs.

Available Formats




CSV

Value	Name	IIH	LSP	SNP	Purge	Status/Reference
0	Reserved					
1	Area Addresses	y	y	n	n	[ISO 10589, "Intermediate System to Intermediate System Intra- Domain Routeing Exchange Protocol for use in Conjunction with the Protocol for Pro Service (ISO 8473)", International Standard 10589: 2002, Second Edition, 2002.]
2	IIS Neighbors	n	y	n	n	[ISO 10589, "Intermediate System to Intermediate System Intra- Domain Routeing Exchange Protocol for use in Conjunction with the Protocol for Pro Service (ISO 8473)", International Standard 10589: 2002, Second Edition, 2002.]
3	ES Neighbors	n	y	n	n	[ISO 10589, "Intermediate System to Intermediate System Intra- Domain Routeing Exchange Protocol for use in Conjunction with the Protocol for Pro Service (ISO 8473)", International Standard 10589: 2002, Second Edition, 2002.]
4	Part. DIS	n	y	n	n	[ISO 10589, "Intermediate System to Intermediate System Intra- Domain Routeing Exchange Protocol for use in Conjunction with the Protocol for Pro Service (ISO 8473)", International Standard 10589: 2002, Second Edition, 2002.]
5	Prefix Neighbors	n	y	n	n	[ISO 10589, "Intermediate System to Intermediate System Intra- Domain Routeing Exchange Protocol for use in Conjunction with the Protocol for Pro Service (ISO 8473)", International Standard 10589: 2002, Second Edition, 2002.]
6	IIS Neighbors	y	n	n	n	[ISO 10589, "Intermediate System to Intermediate System Intra- Domain Routeing Exchange Protocol for use in Conjunction with the Protocol for Pro Service (ISO 8473)", International Standard 10589: 2002, Second Edition, 2002.]
7	Instance Identifier	y	y	y	y	[RFC8202]
8	Padding	y	n	n	n	[ISO 10589, "Intermediate System to Intermediate System Intra- Domain Routeing Exchange Protocol for use in Conjunction with the Protocol for Pro Service (ISO 8473)", International Standard 10589: 2002, Second Edition, 2002.]
9	LSP Entries	n	n	y	n	[ISO 10589, "Intermediate System to Intermediate System Intra- Domain Routeing Exchange Protocol for use in Conjunction with the Protocol for Pro Service (ISO 8473)", International Standard 10589: 2002, Second Edition, 2002.]
10	Authentication	y	y	y	y	[ISO 10589, "Intermediate System to Intermediate System Intra- Domain Routeing Exchange Protocol for use in Conjunction with the Protocol for Pro Service (ISO 8473)", International Standard 10589: 2002, Second Edition, 2002.] [RFC6233]
11	ESN TLV	y	n	y	n	[RFC7602]
12	Opt. Checksum	y	n	y	n	[RFC3358]
13	Purge Originator Identification	n	y	n	y	[RFC6232]
14	LSPBufferSize	n	y	n	n	[ISO 10589, "Intermediate System to Intermediate System Intra- Domain Routeing Exchange Protocol for use in Conjunction with the Protocol for Pro Service (ISO 8473)", International Standard 10589: 2002, Second Edition, 2002.]
15	Router-Fingerprint	y	y	n	y	[RFC8196]
16	Reverse Metric	y	n	n	n	[RFC8500]
17	IS-IS Area Node IDs TLV (TEMPORARY - registered 2019-08-08, expires 2020-08-08)	n	y	n	n	[draft-ietf-lsr-dynamic-flooding]
18	IS-IS Flooding Path TLV (TEMPORARY - registered 2019-08-08, expires 2020-08-08)	n	y	n	n	[draft-ietf-lsr-dynamic-flooding]

How to handle unrecognized TLVs?

- What should you do when you receive a TLV with an unknown type?
- Happens when another device is running a newer version of the protocol.
- Options:
 - Ignore and propagate the TLV
 - Ignore and remove the TLV
 - Declare an error
 - Note: ignoring is possible because the length L is known and the value can be skipped
- The originator of the TLV (who by definition understands it) sets some extra bits in the TLV to instruct the receiver what to do.

Example: Label Distribution Protocol (LDP)



U-bit

Unknown TLV bit. Upon receipt of an unknown TLV, if U is clear (=0), a notification **MUST** be returned to the message originator and the entire message **MUST** be ignored; if U is set (=1), the unknown TLV **MUST** be silently ignored and the rest of the message processed as if the unknown TLV did not exist. The sections following that define TLVs specify a value for the U-bit.

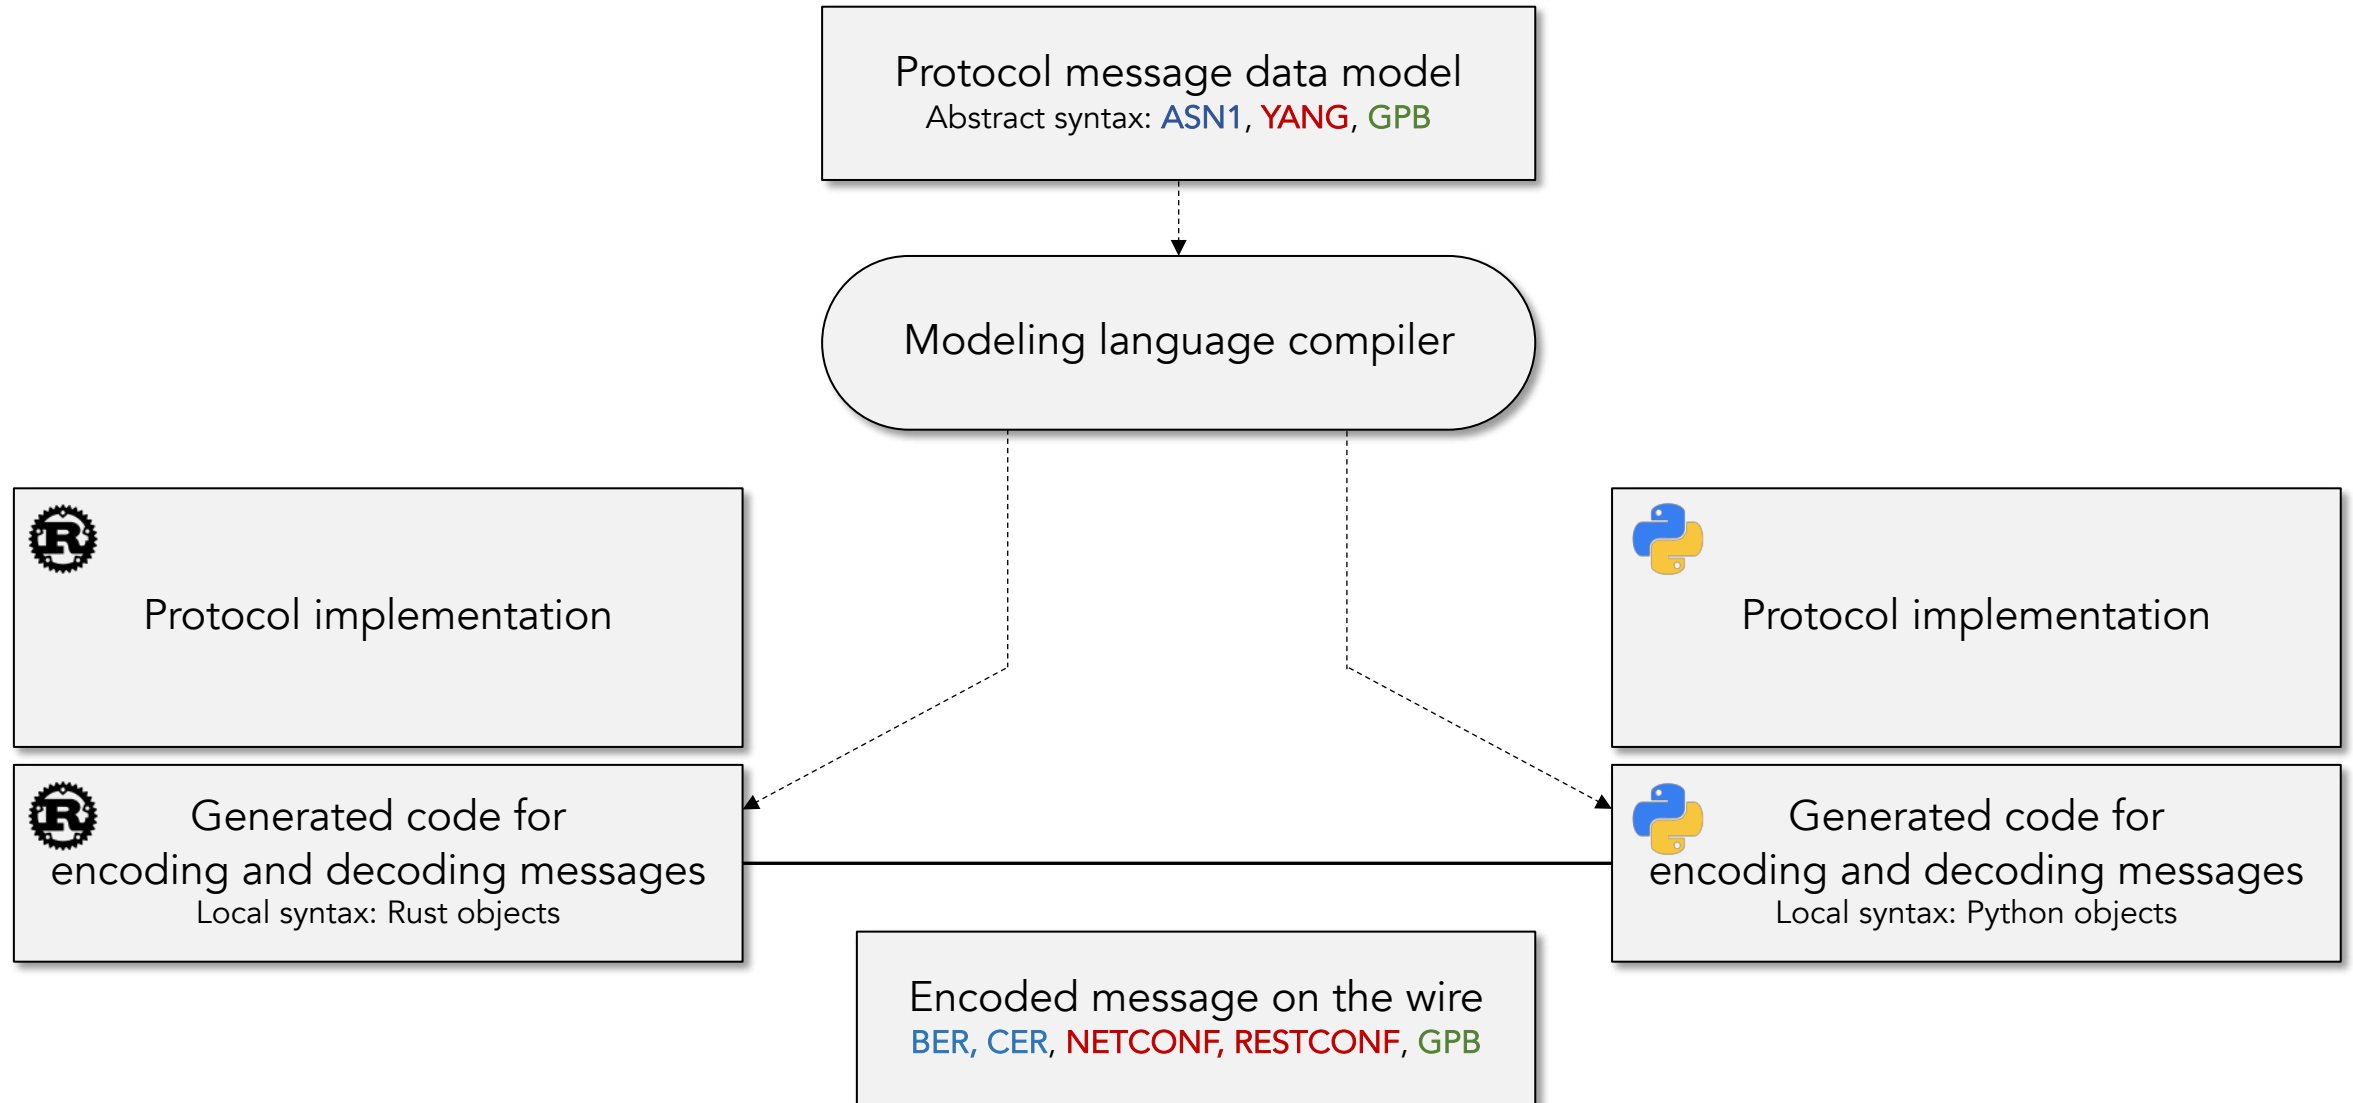
F-bit

Forward unknown TLV bit. This bit applies only when the U-bit is set and the LDP message containing the unknown TLV is to be forwarded. If F is clear (=0), the unknown TLV is not forwarded with the containing message; if F is set (=1), the unknown TLV is forwarded with the containing message. The sections following that define TLVs specify a value for the F-bit. By setting both the U- and F-bits, a TLV can be propagated as opaque data through nodes that do not recognize the TLV.

TLV encoding is also used in:

- BGP
- OSPF
- ISIS
- RSVP
- Many more...

Beyond TLVs: formal modeling languages



Examples

Modeling language	Protocols
ASN.1	Simple Network Management Protocol (SNMP)
ASN.1	Light-weight Directory Access Protocol (LDAP)
YANG	Network Configuration Protocol (NETCONF)
Protobuf	GRPC Network Management Interface (gNMI)
Thrift	Routing In Fat Trees (RIFT)

Examples

ASN.1 model for LDAP protocol

```
SearchRequest ::= [APPLICATION 3] SEQUENCE {
    baseObject      LDAPDN,
    scope           ENUMERATED {
        baseObject      (0),
        singleLevel     (1),
        wholeSubtree    (2),
        ... },
    derefAliases    ENUMERATED {
        neverDerefAliases (0),
        derefInSearching  (1),
        derefFindingBaseObj (2),
        derefAlways       (3) },
    sizeLimit       INTEGER (0 .. maxInt),
    timeLimit       INTEGER (0 .. maxInt),
    typesOnly       BOOLEAN,
    filter          Filter,
    attributes      AttributeSelection }
```

YANG model for interface management

```
/*
 * Configuration data nodes
 */

container interfaces {
    description
        "Interface configuration parameters.";

    list interface {
        key "name";

        description
            "The list of configured interfaces on the device.

            The operational state of an interface is available in the
            /interfaces-state/interface list. If the configuration of a
            system-controlled interface cannot be used by the system
            (e.g., the interface hardware present does not match the
            interface type), then the configuration is not applied to
            the system-controlled interface shown in the
            /interfaces-state/interface list. If the configuration
            of a user-controlled interface cannot be used by the system,
            the configured interface is not instantiated in the
            /interfaces-state/interface list.";

        leaf name {
            type string;
            description
                "The name of the interface.

                A device MAY restrict the allowed values for this leaf,
                possibly depending on the type of the interface."
        }
    }
}
```

Thrift model for RIFT protocol

```
union TIEElement {
    /** used in case of enum common.TITypeType.NodeTIType */
    1: optional NodeTIEElement node;
    /** used in case of enum common.TITypeType.PrefixTIType */
    2: optional PrefixTIEElement prefixes;
    /** positive prefixes (always southbound)
     * It MUST NOT be advertised within a North TIE and ignored otherwise
     */
    3: optional PrefixTIEElement positive_disaggregation_prefixes;
    /** transitive, negative prefixes (always southbound) which
     * MUST be aggregated and propagated
     * according to the specification
     * southwards towards lower levels to heal
     * pathological upper level partitioning, otherwise
     * blackholes may occur in multiplane fabrics.
     * It MUST NOT be advertised within a North TIE.
     */
    5: optional PrefixTIEElement negative_disaggregation_prefixes;
    /** externally reimported prefixes */
    6: optional PrefixTIEElement external_prefixes;
    /** positive external disaggregated prefixes (always southbound).
     * It MUST NOT be advertised within a North TIE and ignored otherwise
     */
    7: optional PrefixTIEElement positive_external_disaggregation_prefixes;
    /** Key-Value store elements */
    9: optional KeyValueTIEElement keyvalues;
}

/** TIE packet */
struct TIEPacket {
    1: required TIEHeader header;
    2: required TIEElement element;
}

/** content of a RIFT packet */
union PacketContent {
    1: optional LIEPacket lie;
    2: optional TIDEPacket tide;
    3: optional TIREDPacket tire;
    4: optional TIEPacket tie;
}
```

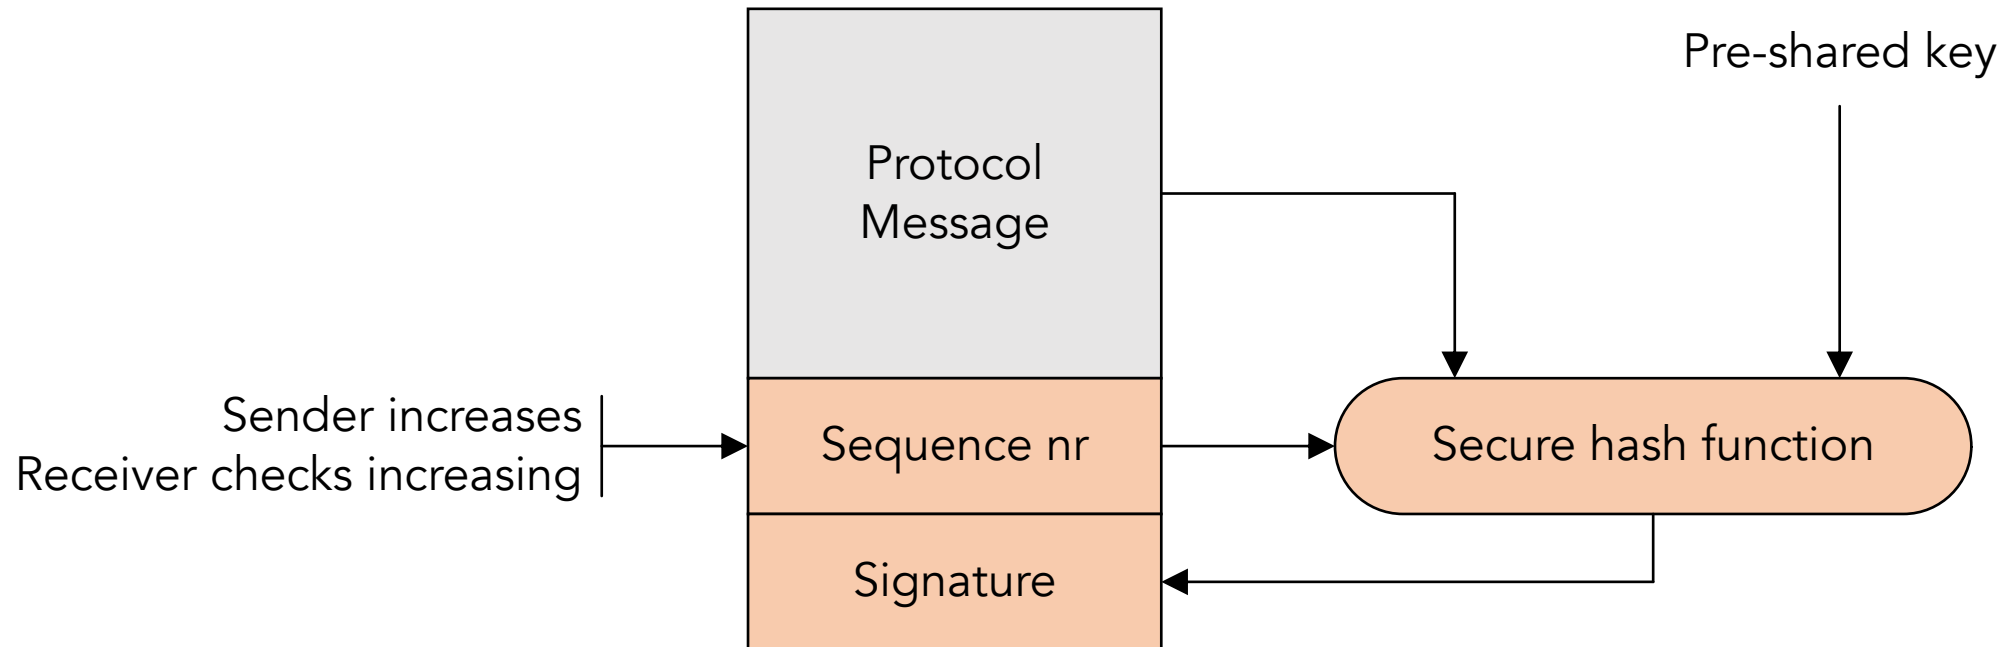
Control protocol security

Control plane security

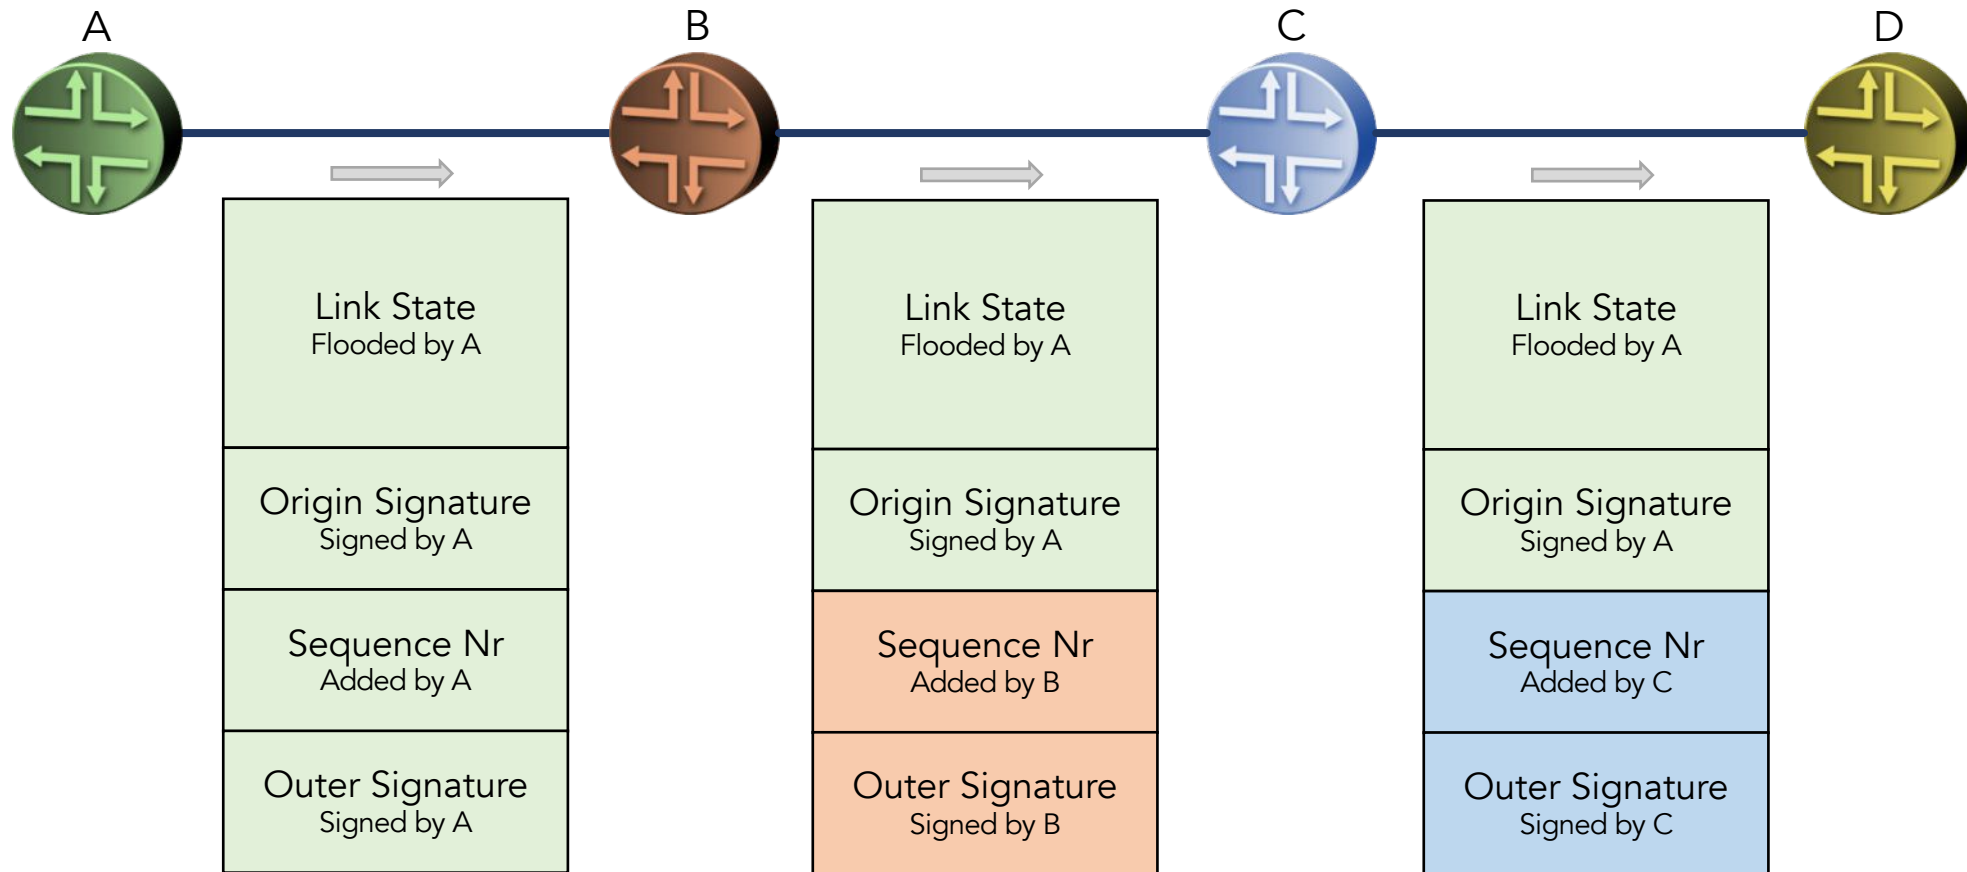
- Authentication
 - Is the router that I am talking to really the router it claims to be?
- Integrity
 - Are the routing messages modified or spoofed by an adversary?
 - Special case: replay attack
- Confidentiality?
 - Encrypt the protocol messages (often **not** a requirement)
- Authorization
 - Is this router allowed to have an adjacency with me?
 - Is this router allowed to advertise this route?
 - If it originates the route, can it prove it owns the prefix?
 - If it propagated the route, can it prove that the advertised path is correct?

Neighbor authentication and integrity

- Pre-shared key configured on each router
- Sign each message with secure hash using pre-shared key
- Add sequence number to prevent replay attacks
- Message often not encrypted

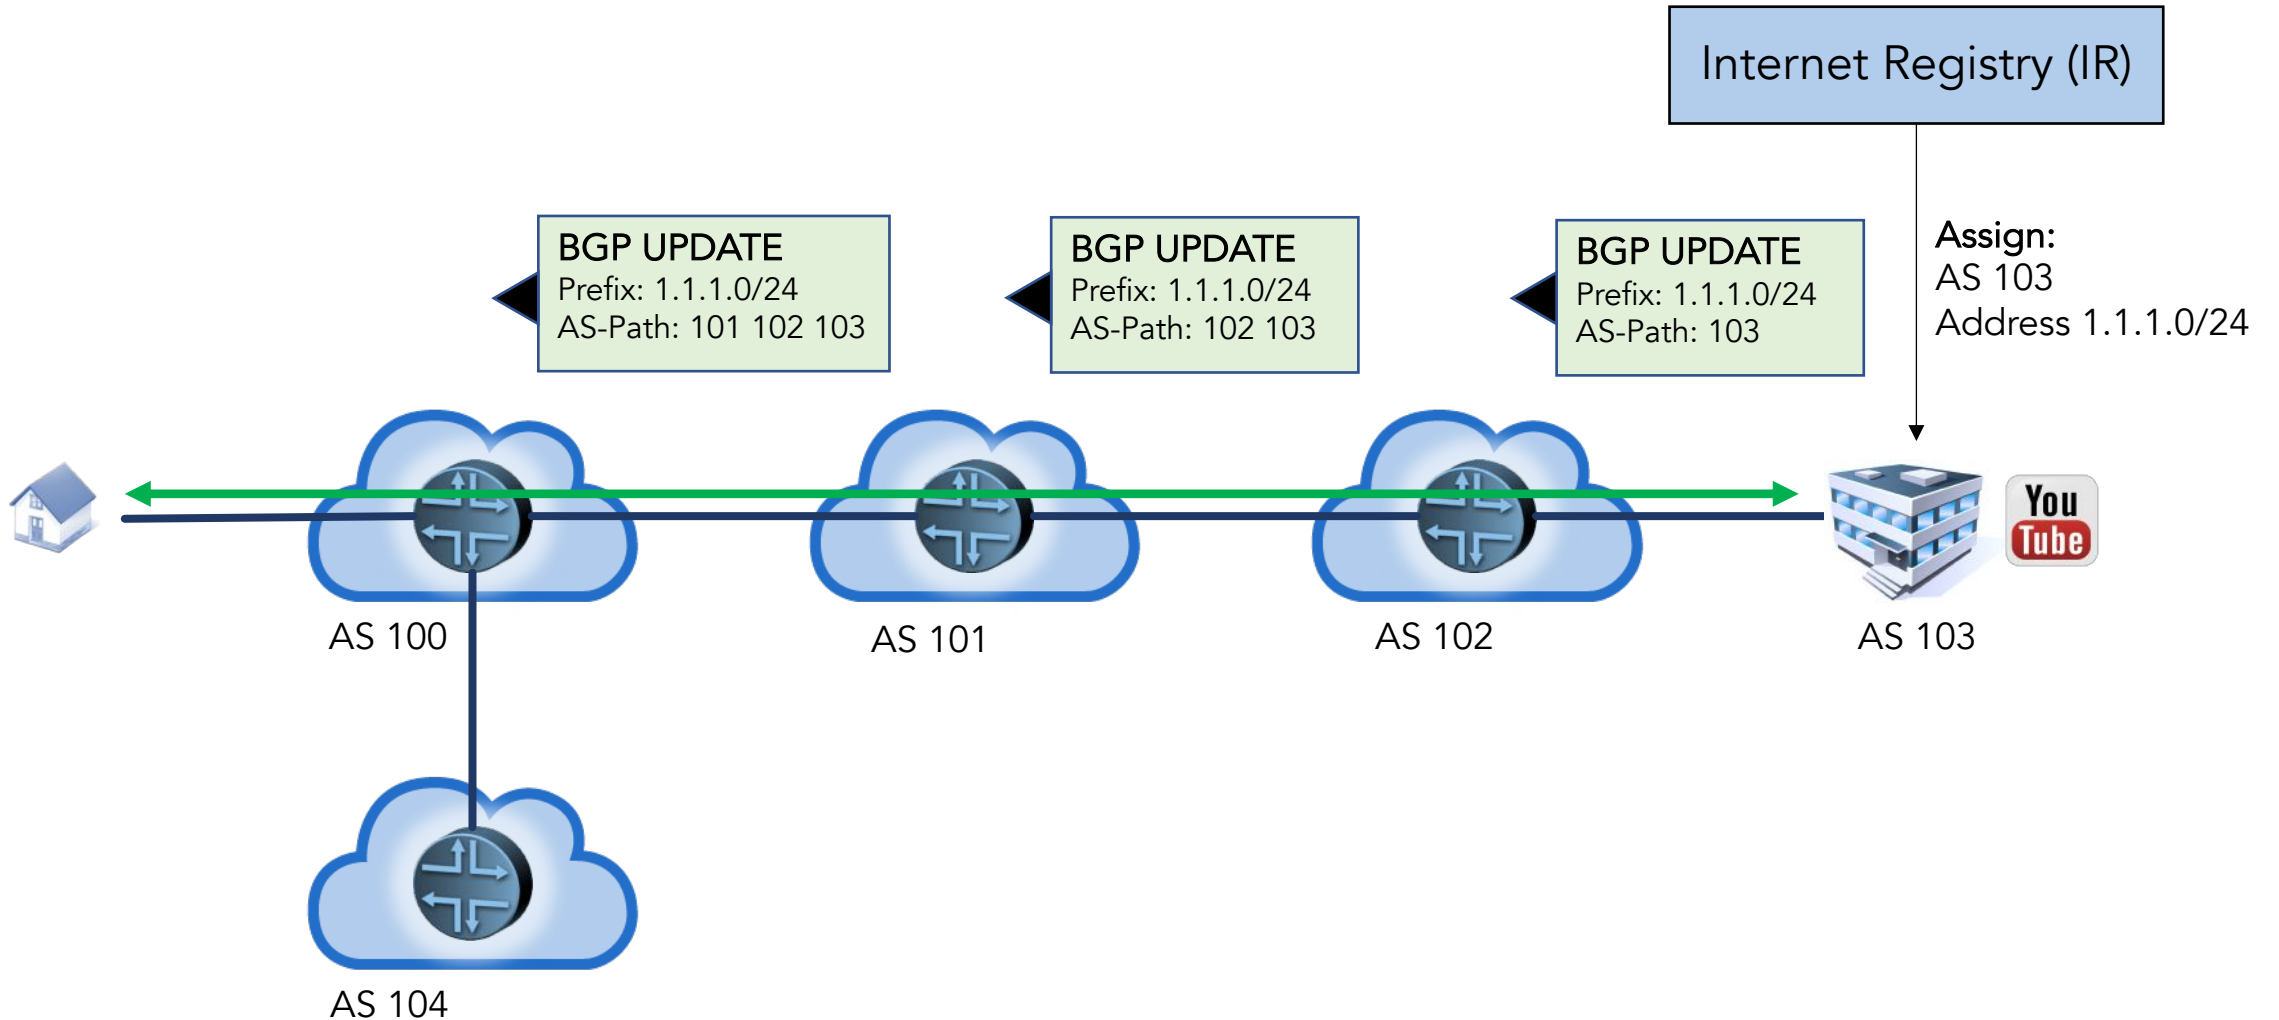


Link state flooding: origin and outer key

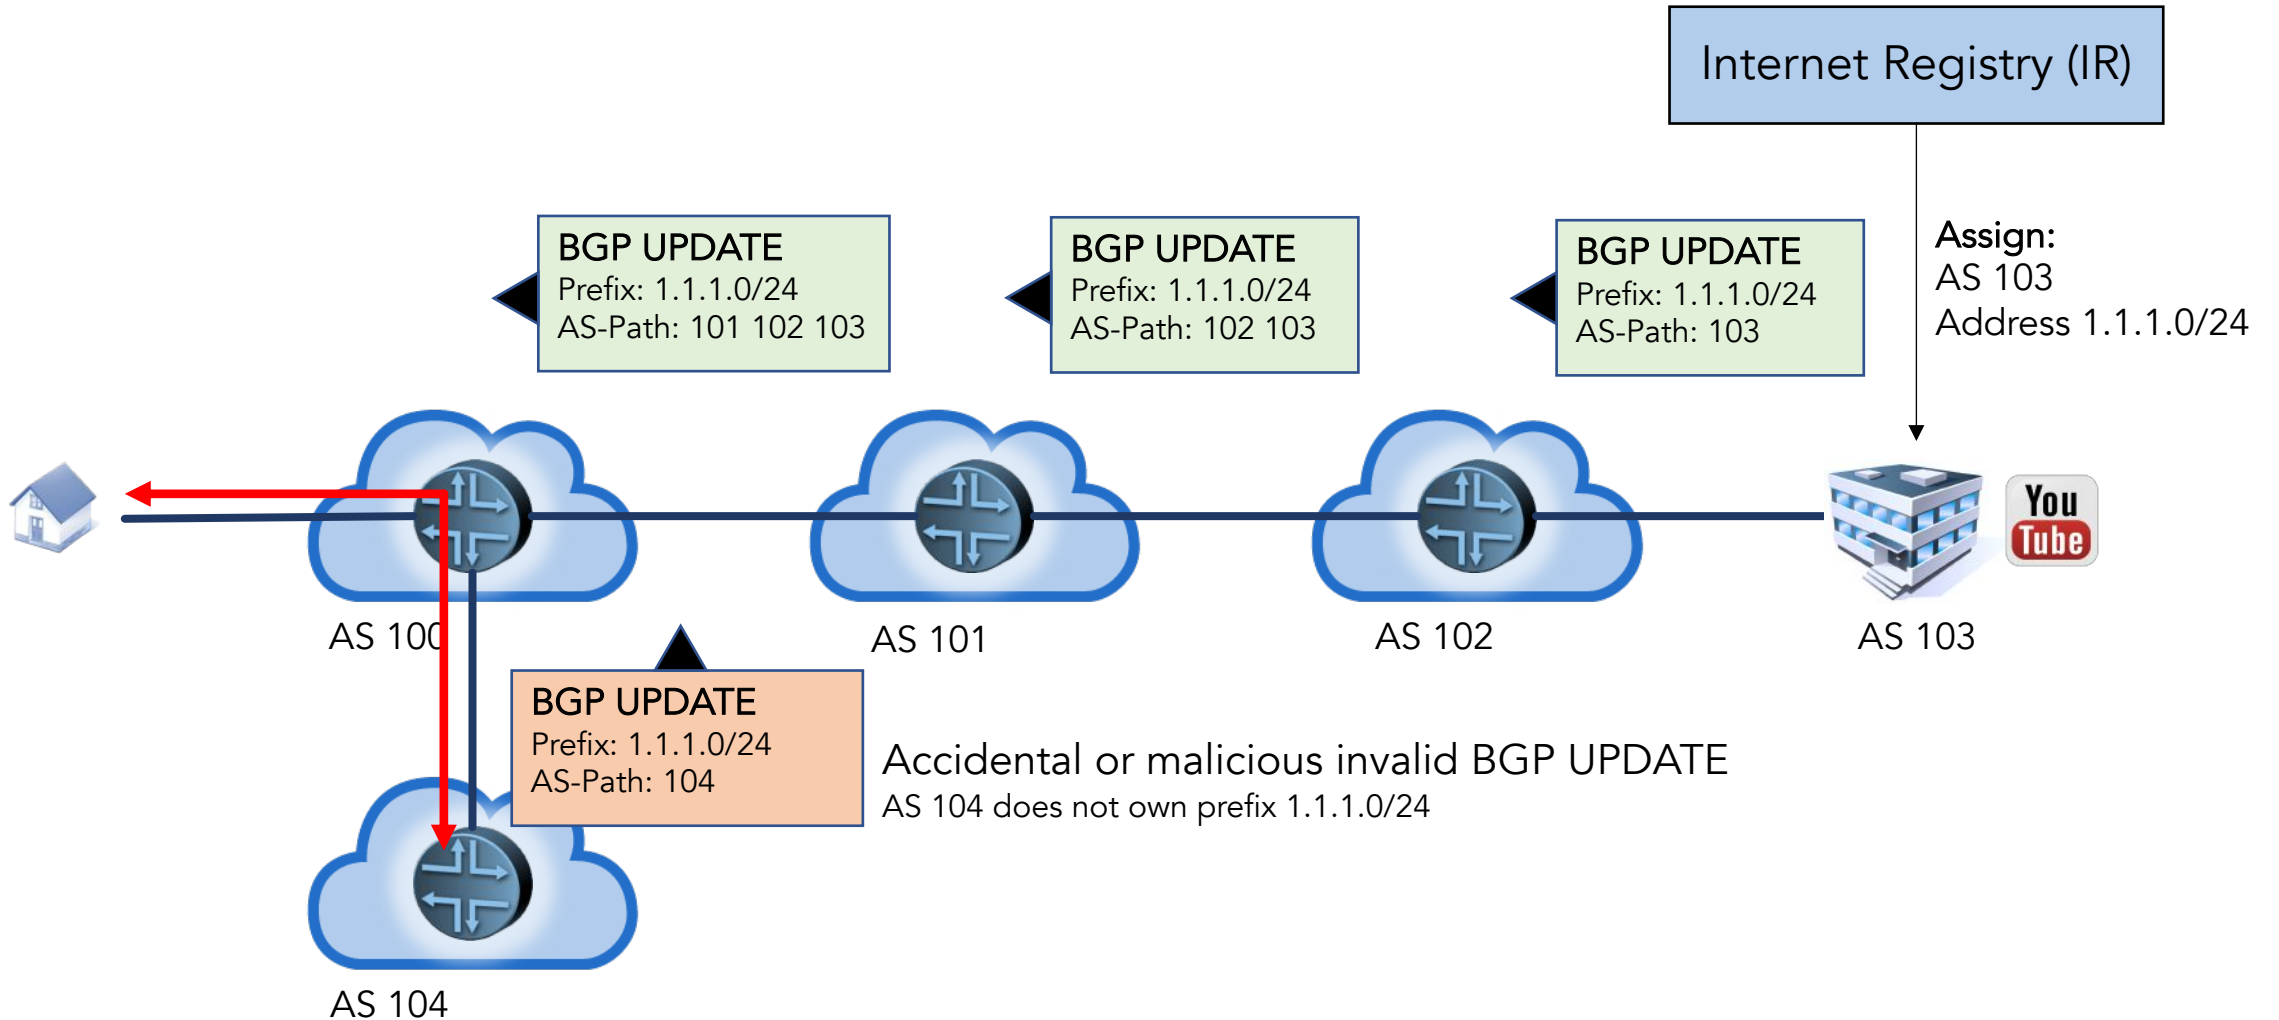


Every router in the networks must have the same pre-shared origin key (*not feasible for IDR*)
Pre-shared outer keys can be different per link, but are the same in practice
Pre-shared key roll-over is very painful; often does not happen in practice

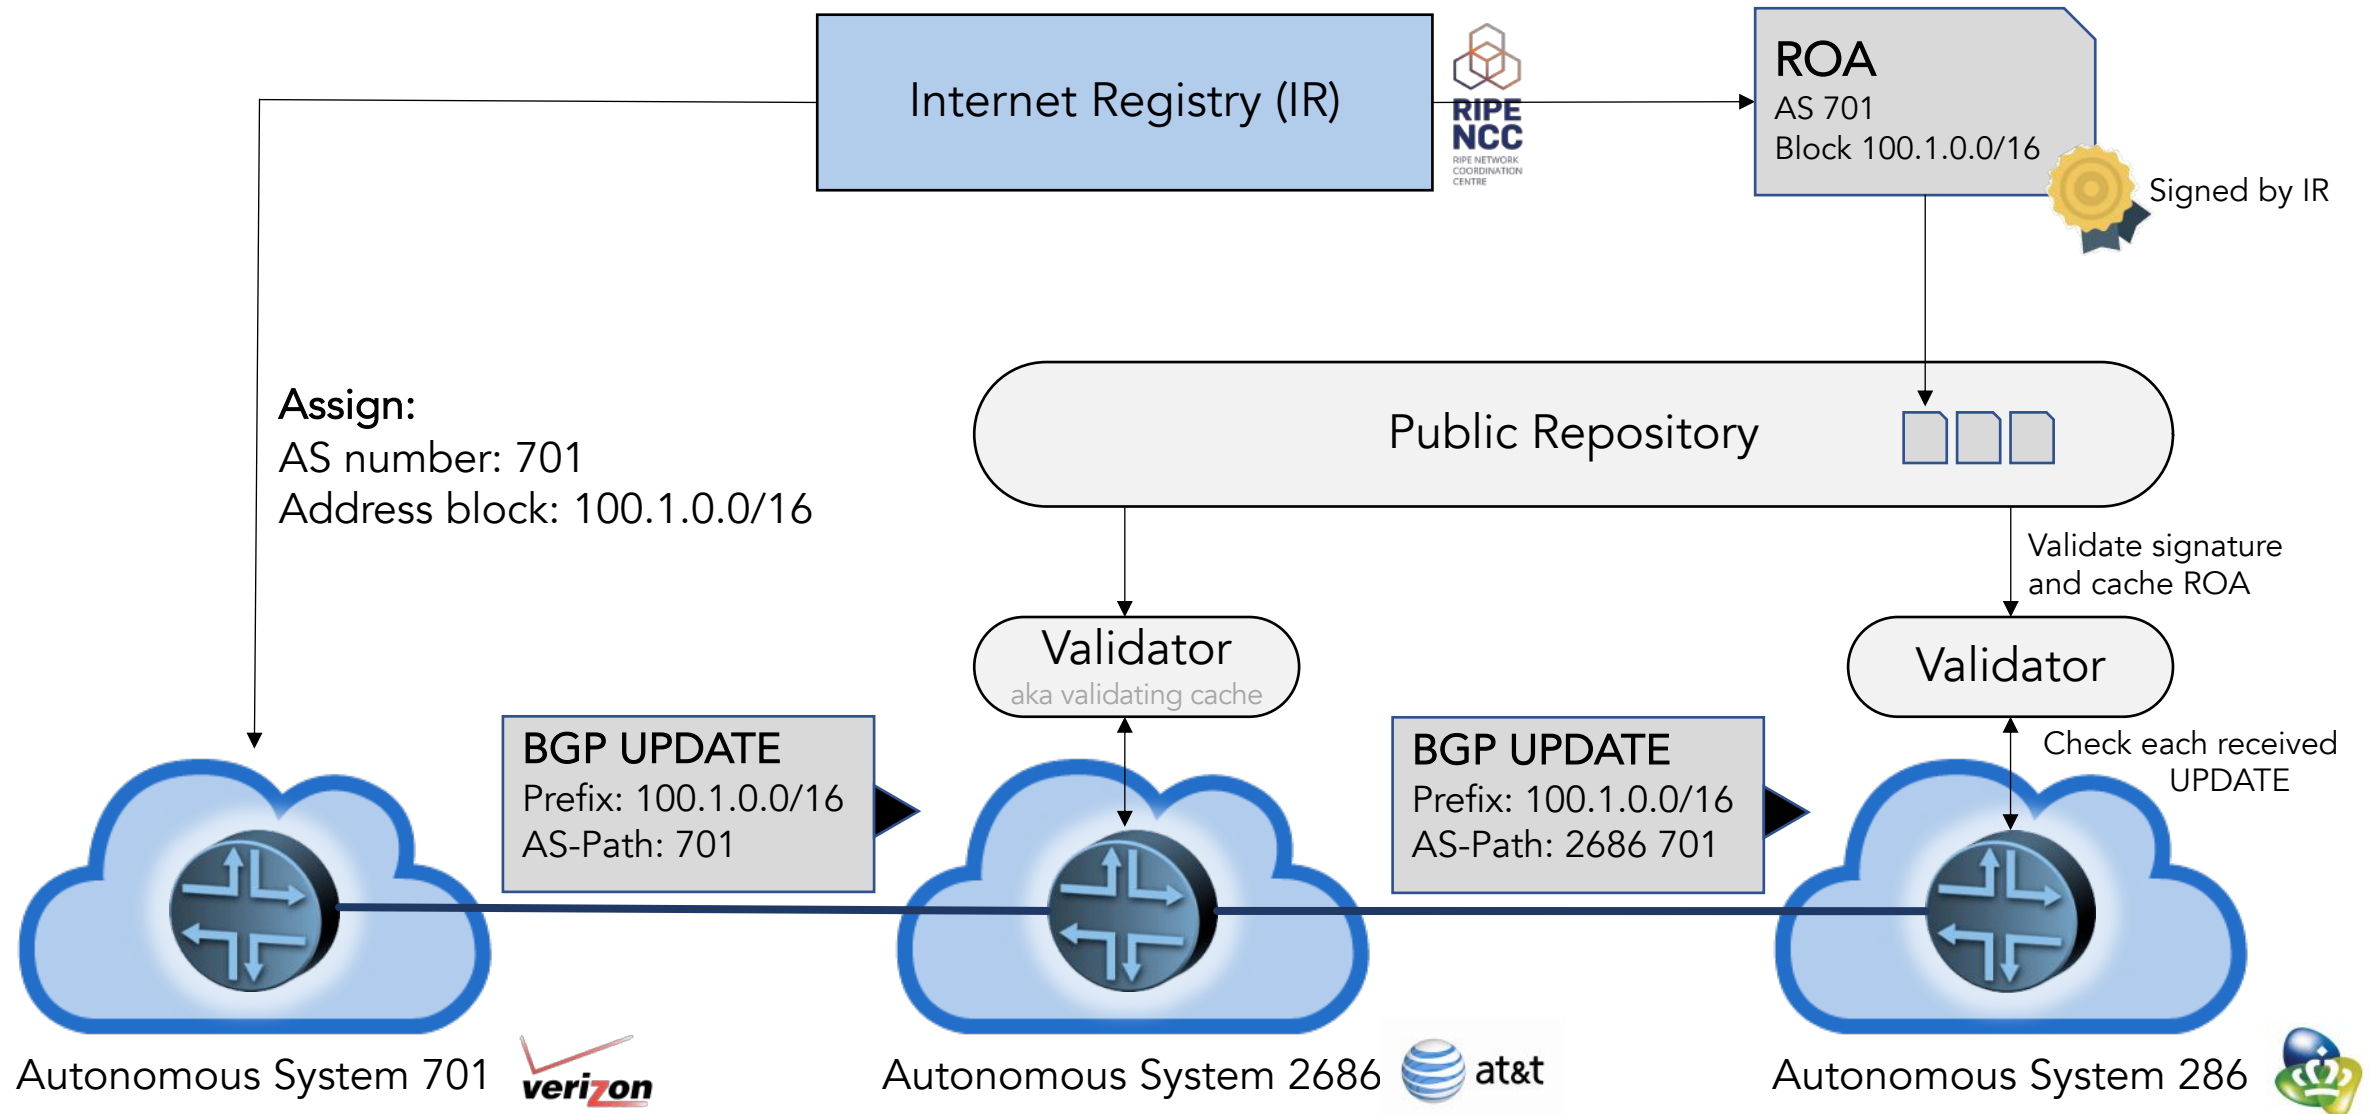
BGP prefix hijacking




BGP prefix hijacking



BGP route origin authentication (ROA) (Simplified)



Also need BGP path validation



Accidental violation of valley-free routing

Y should not do transit through customer B

B should not advertise A to Y


Y should not accept A from B

(There are also malicious flavors of this, e.g. BGP poisoning)

Upgrades, crashes

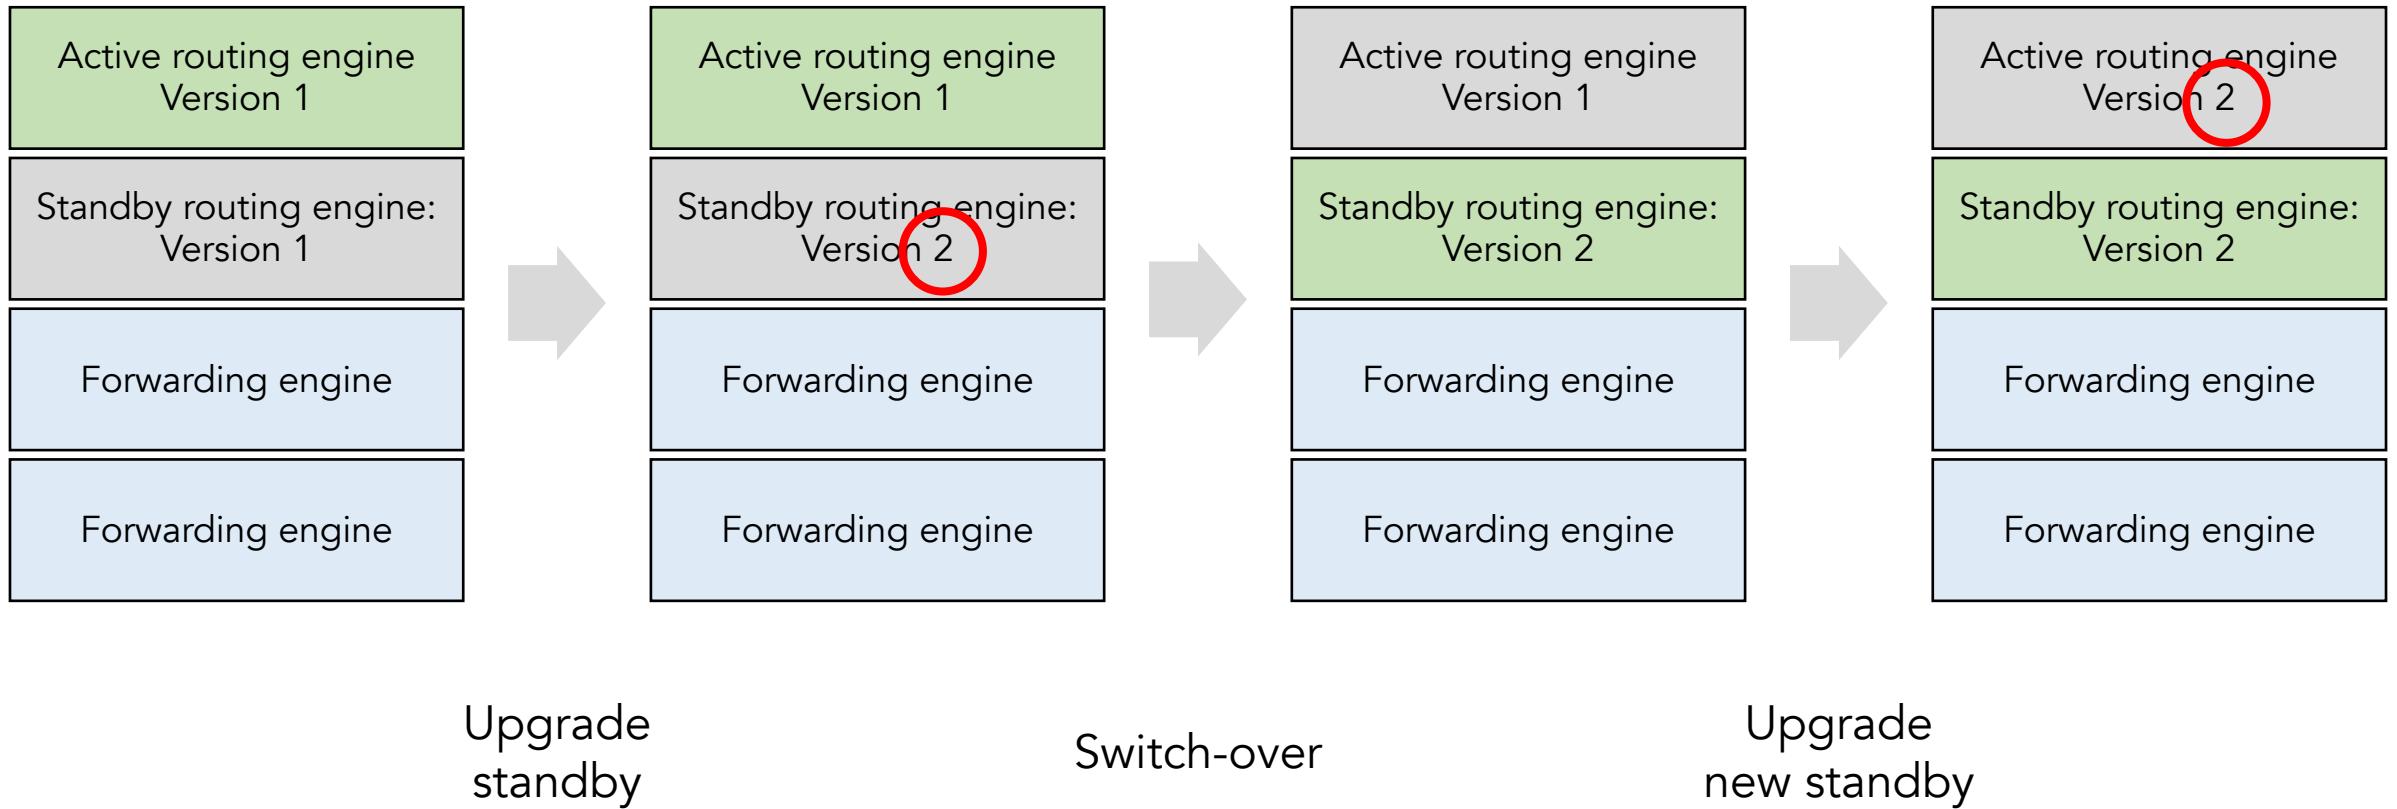
There is no maintenance window for the Internet

Chassis devices



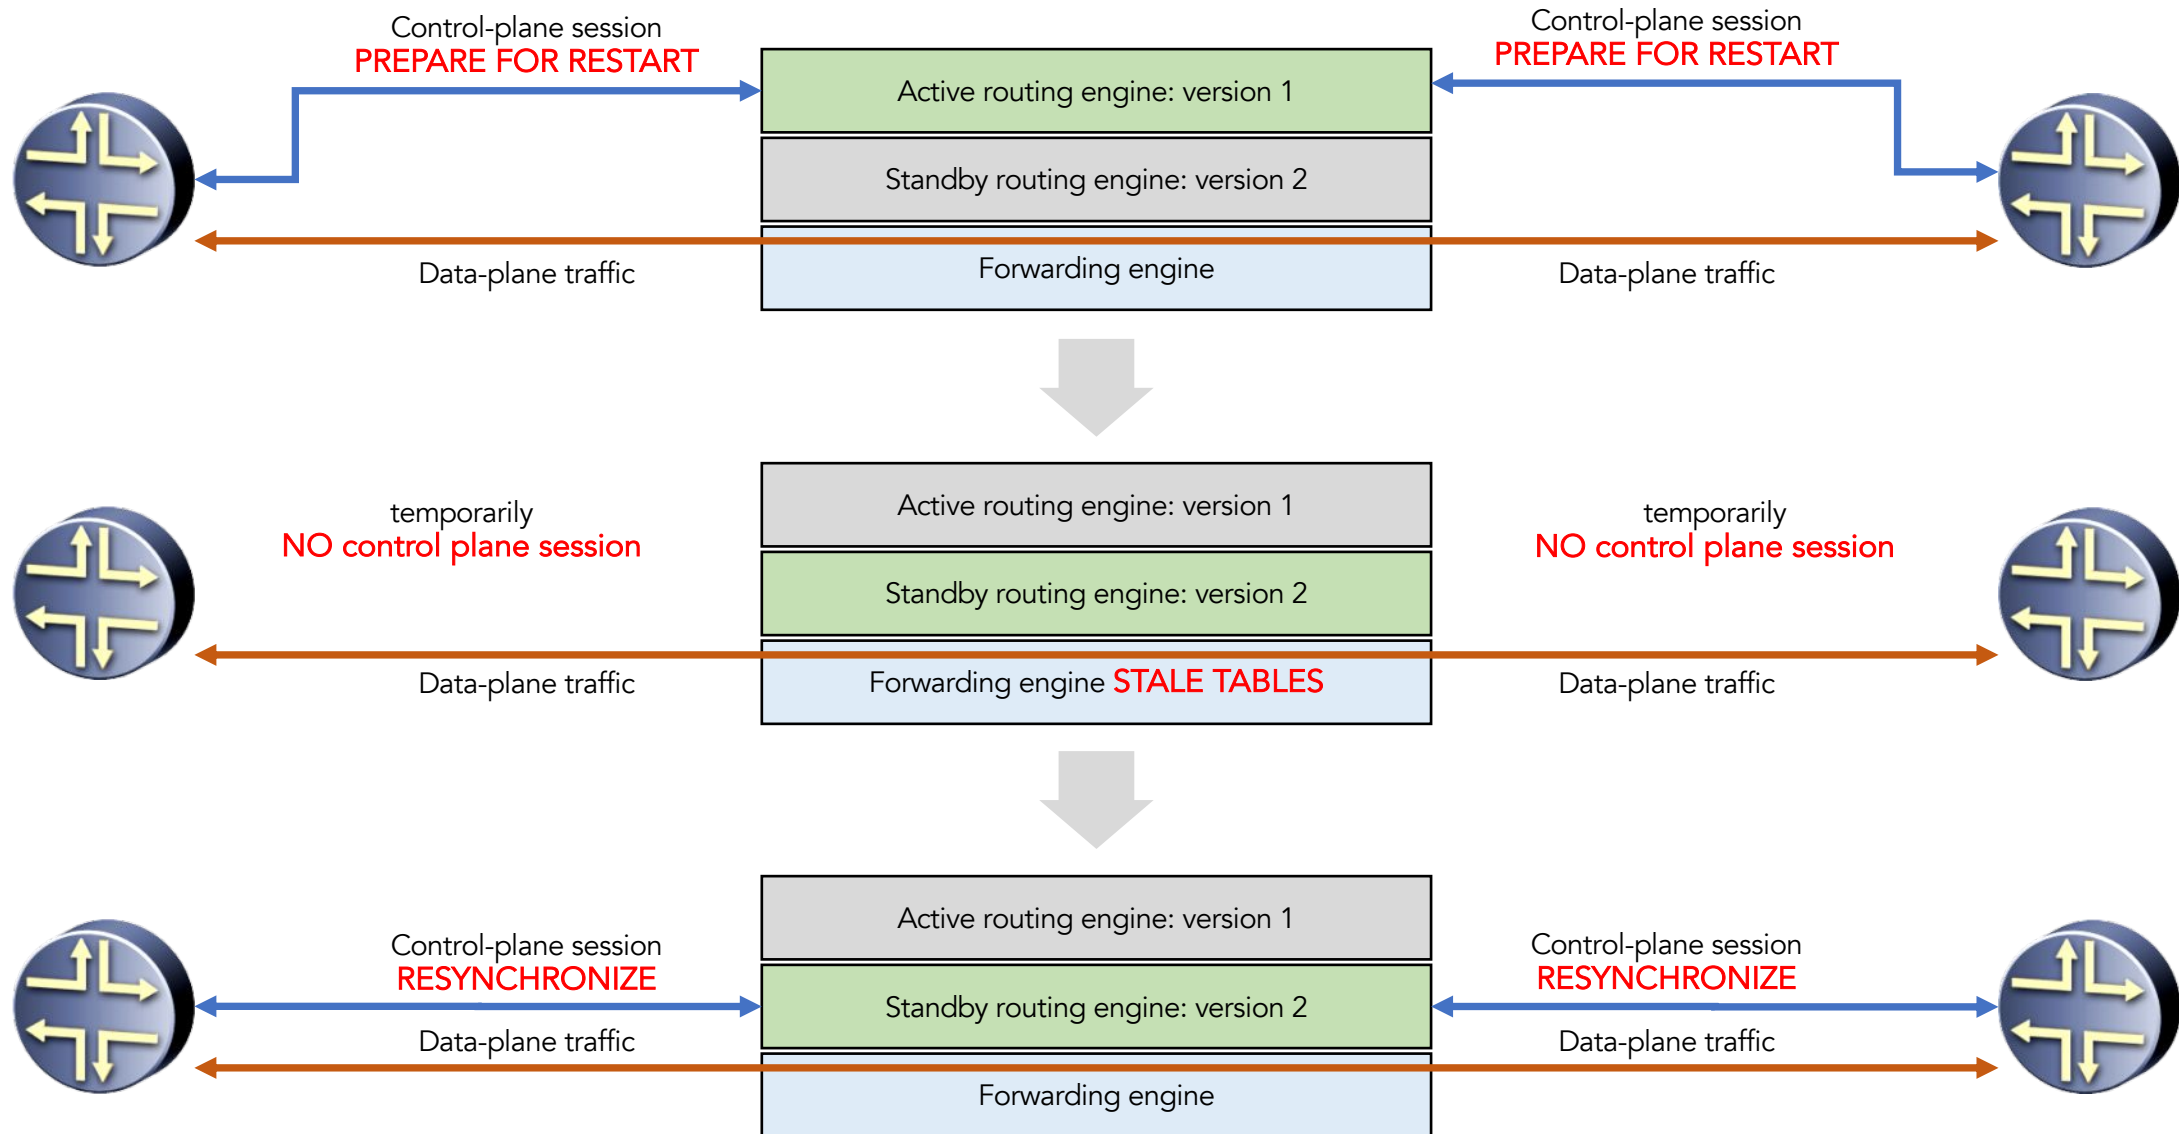
Everything (routing engine, switch fabric, power supplies, fans, ...) redundant and in-service replaceable

In-Service Software Upgrade (ISSU)

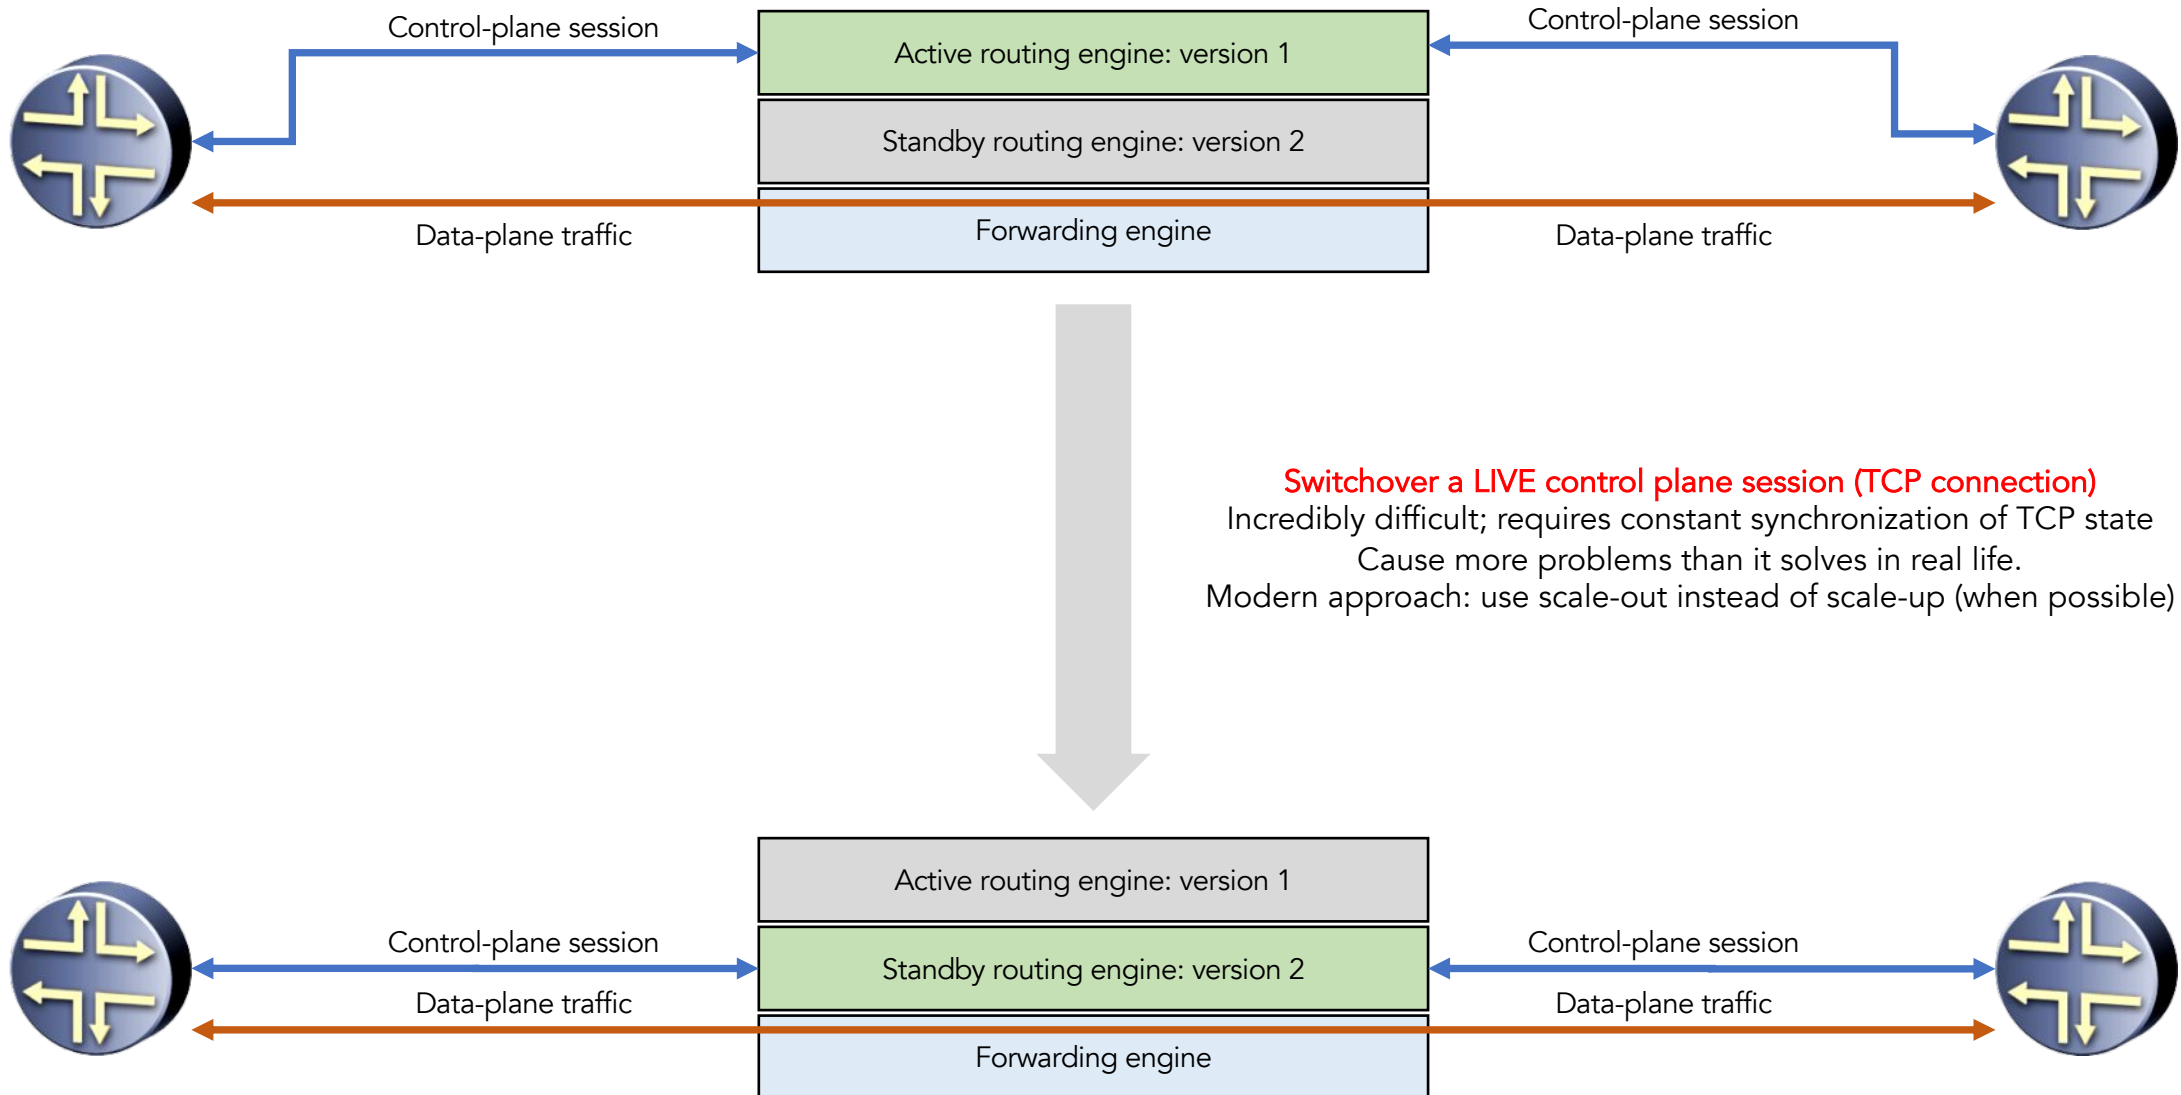


Non-Stop Forwarding (NSF)

Also known as Graceful Restart (GR)



Non-Stop Routing (NSR)



Pizza-box devices




Does not have redundant routing hardware routing engines.

(Sometimes in redundant software routing engines using virtual machines or containers.)


Cheap commodity hardware

Scale-out instead of scale-up



Scale-up

- One large chassis switch
- Internal switch fabric
- If switch fails, all is lost
- NSF or NSR for upgrade
- Everything must be redundant
- Expensive
- Complex



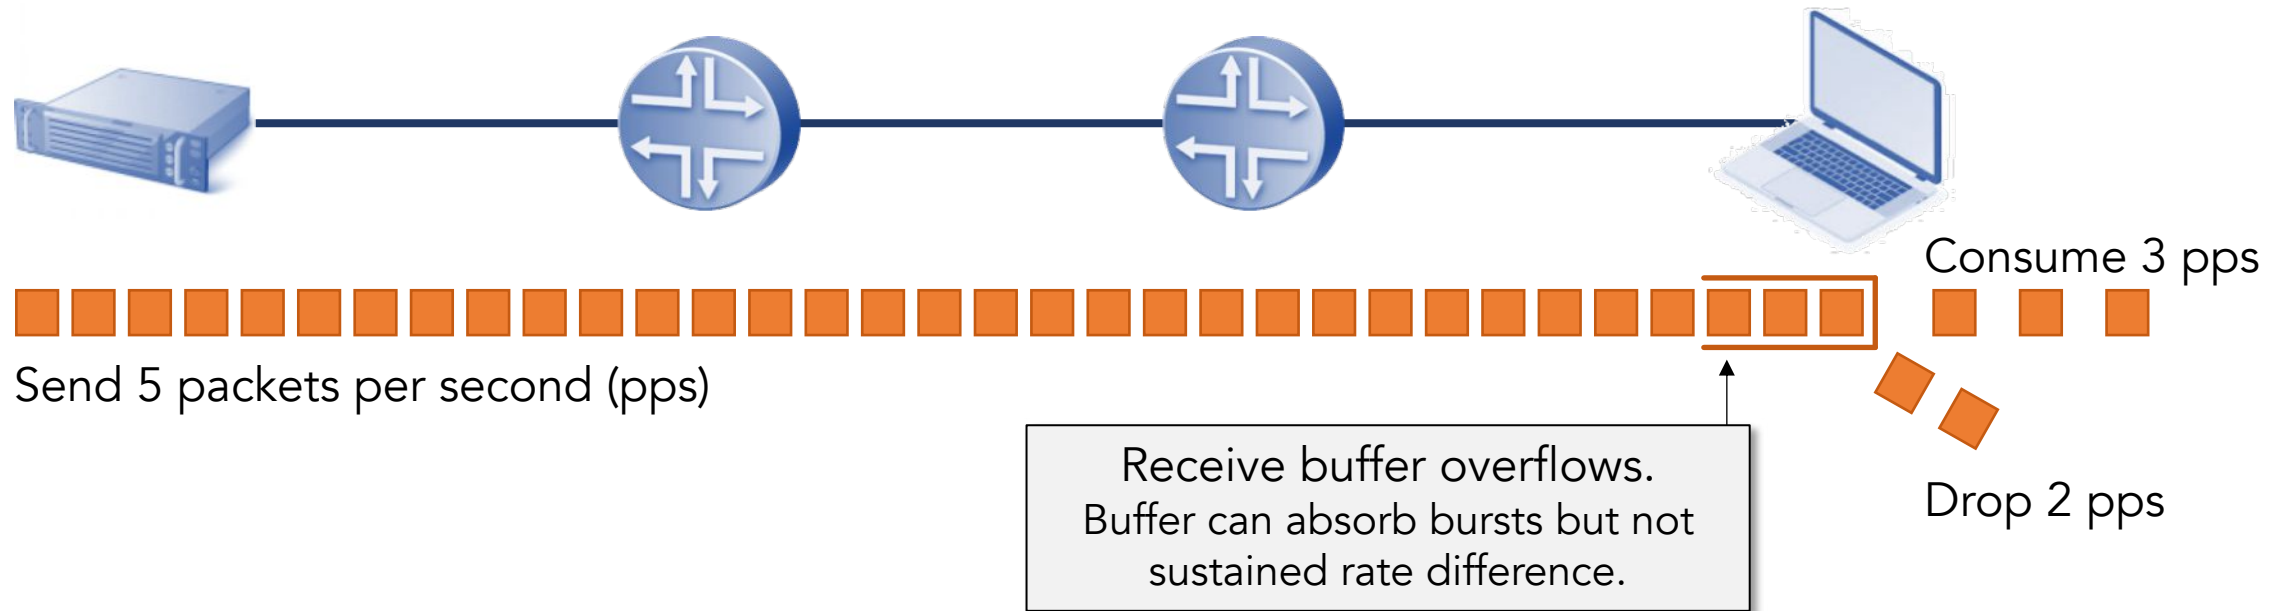
Scale-out

- Multiple small pizza-box switches
- Clos fabric
- If one switch fails, only 1/Nth of capacity lost
- Rolling upgrade
- No redundancy needed inside switch
- Cheap
- Simple

Flow control
Congestion control
Error control
in the data plane

Flow control

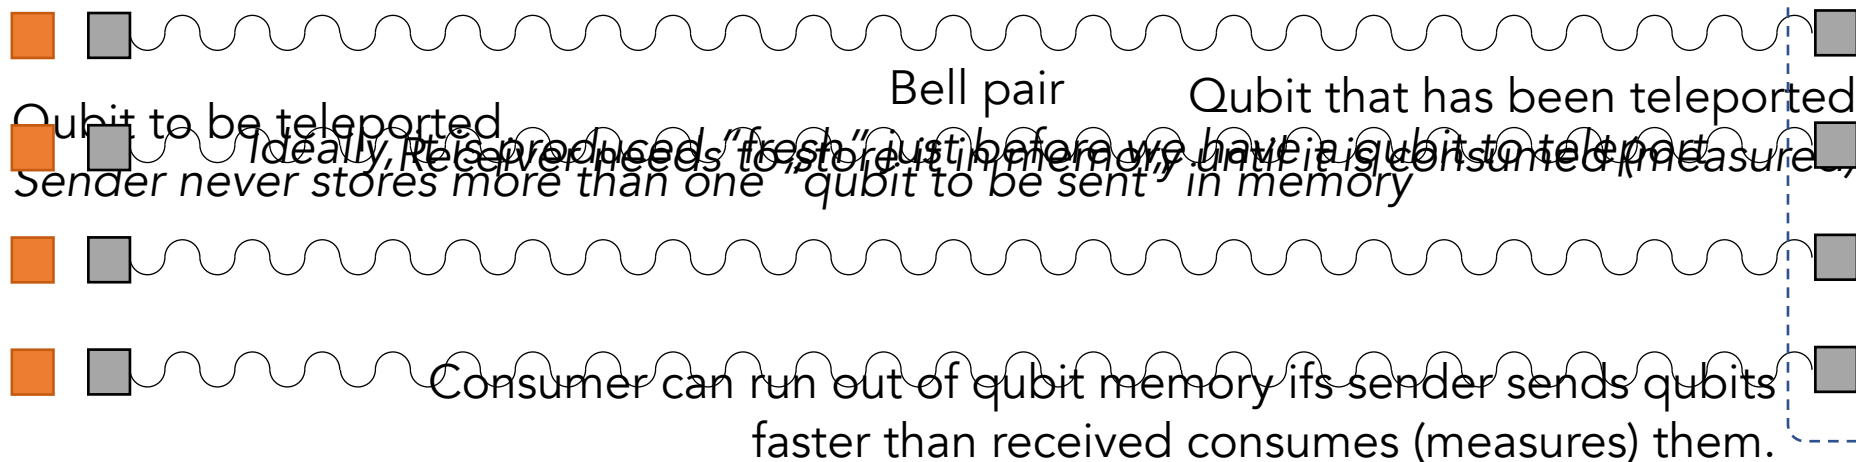
Make sure the sender does not produce data faster than the receiver can consume it.



Note: here it is the receiver that is not keeping up; the network can keep up fine.

Flow control in quantum network

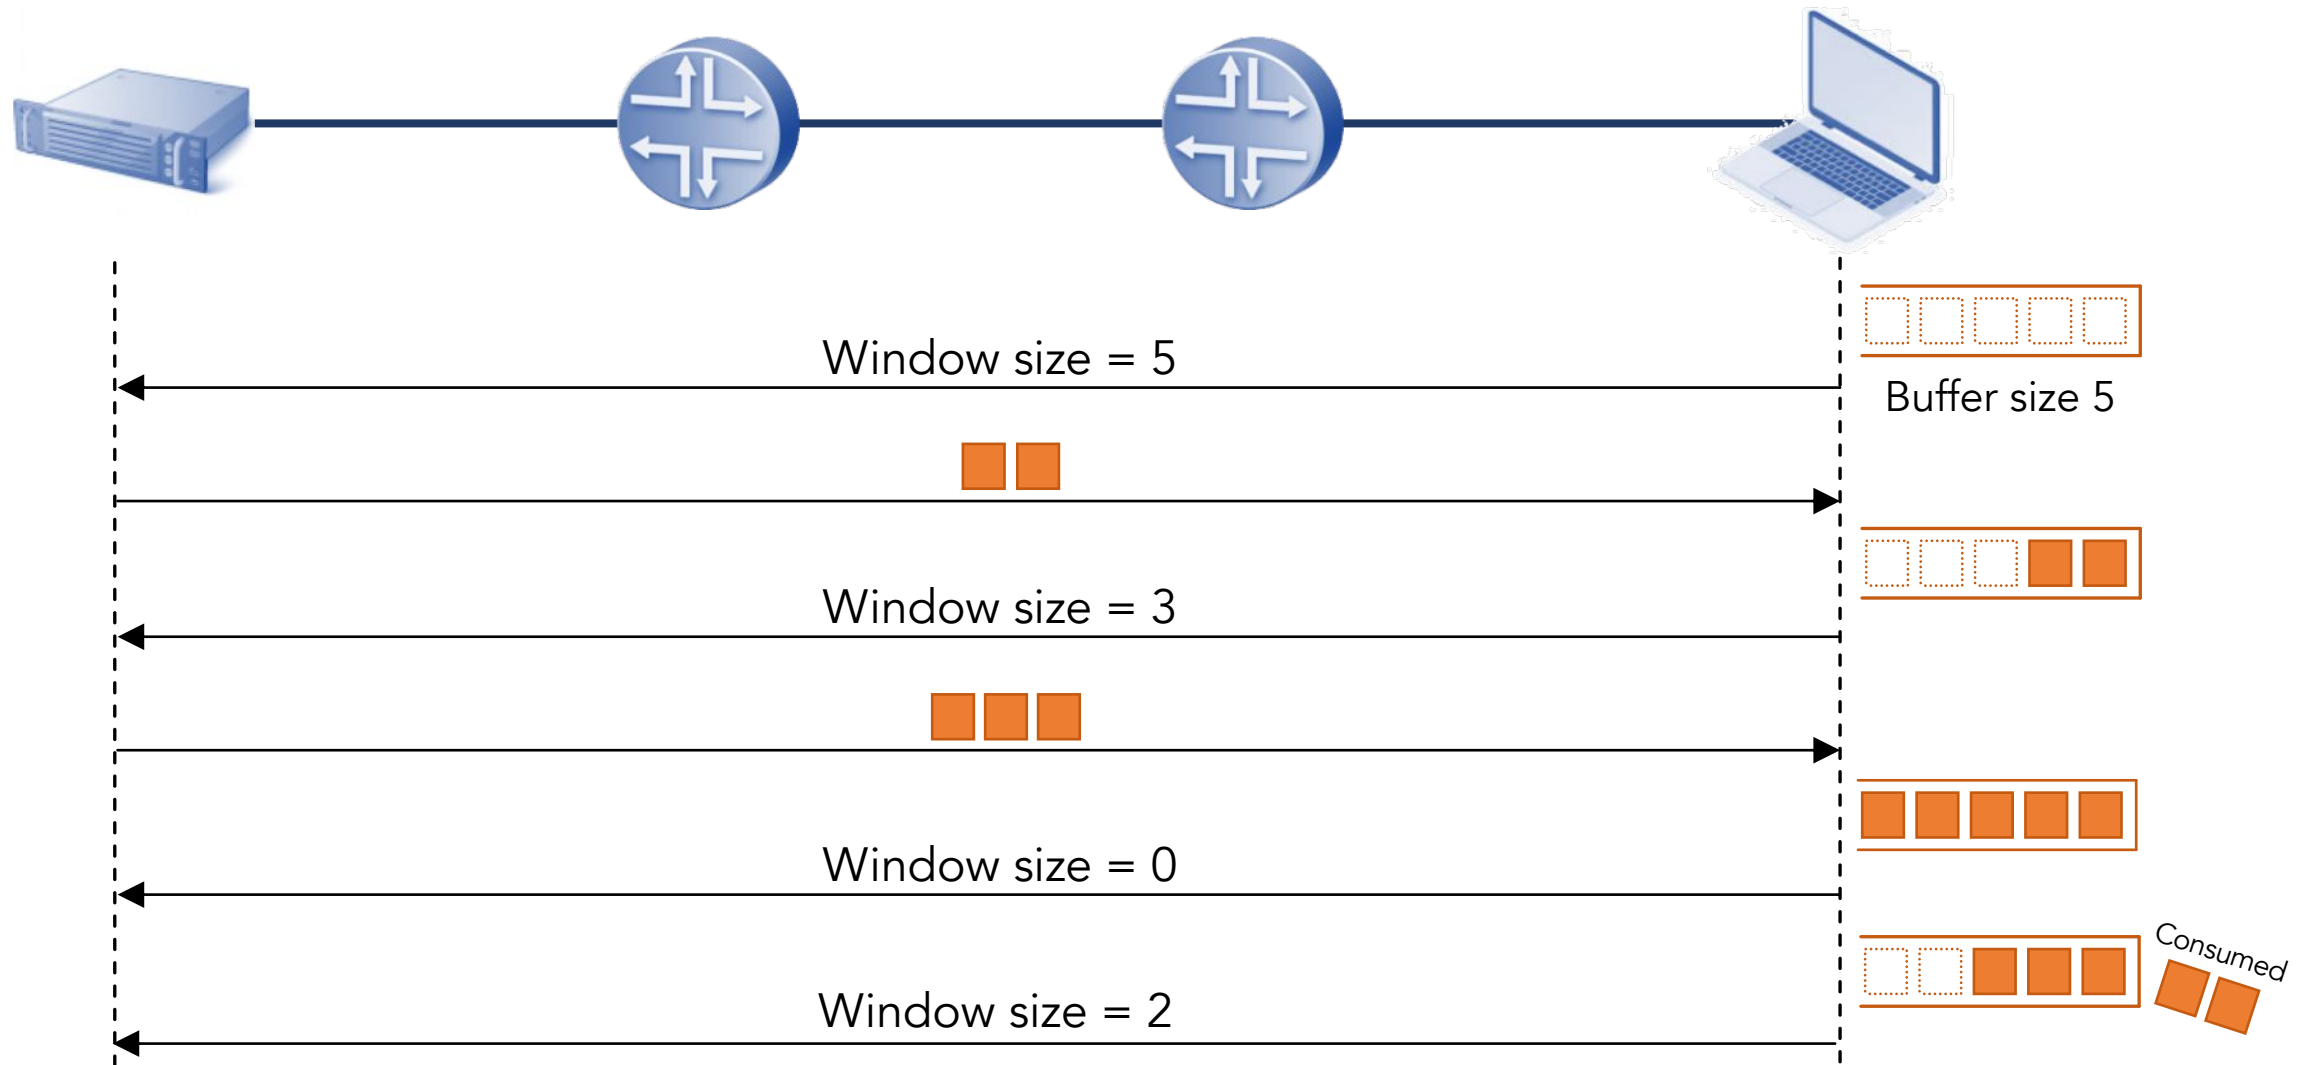
Make sure sender does not produce qubits faster than receiver can consume them.




Flow control approaches

- General principle: make the data arrive more slowly
- Approach 1: **Sliding window**
 - Receiver makes the sender slow down by delaying permission to send
 - It is okay to receive data later (delay in-sensitive traffic)
 - Web pages, file transfer, e-mail, ...
- Approach 2: **Adaptive encoding**
 - The sender sends less data by switching to more lossy compression
 - It is not okay to receive data later (delay-sensitive traffic)
 - Voice, live video, video on demand, ...
- Related approach: **Time slot scheduling**
 - Used in Time Division Multiplexing (TDM) and wireless networks
 - Avoids congestion from happening in the first place
- Sender must be able to back-pressure application

Sliding window



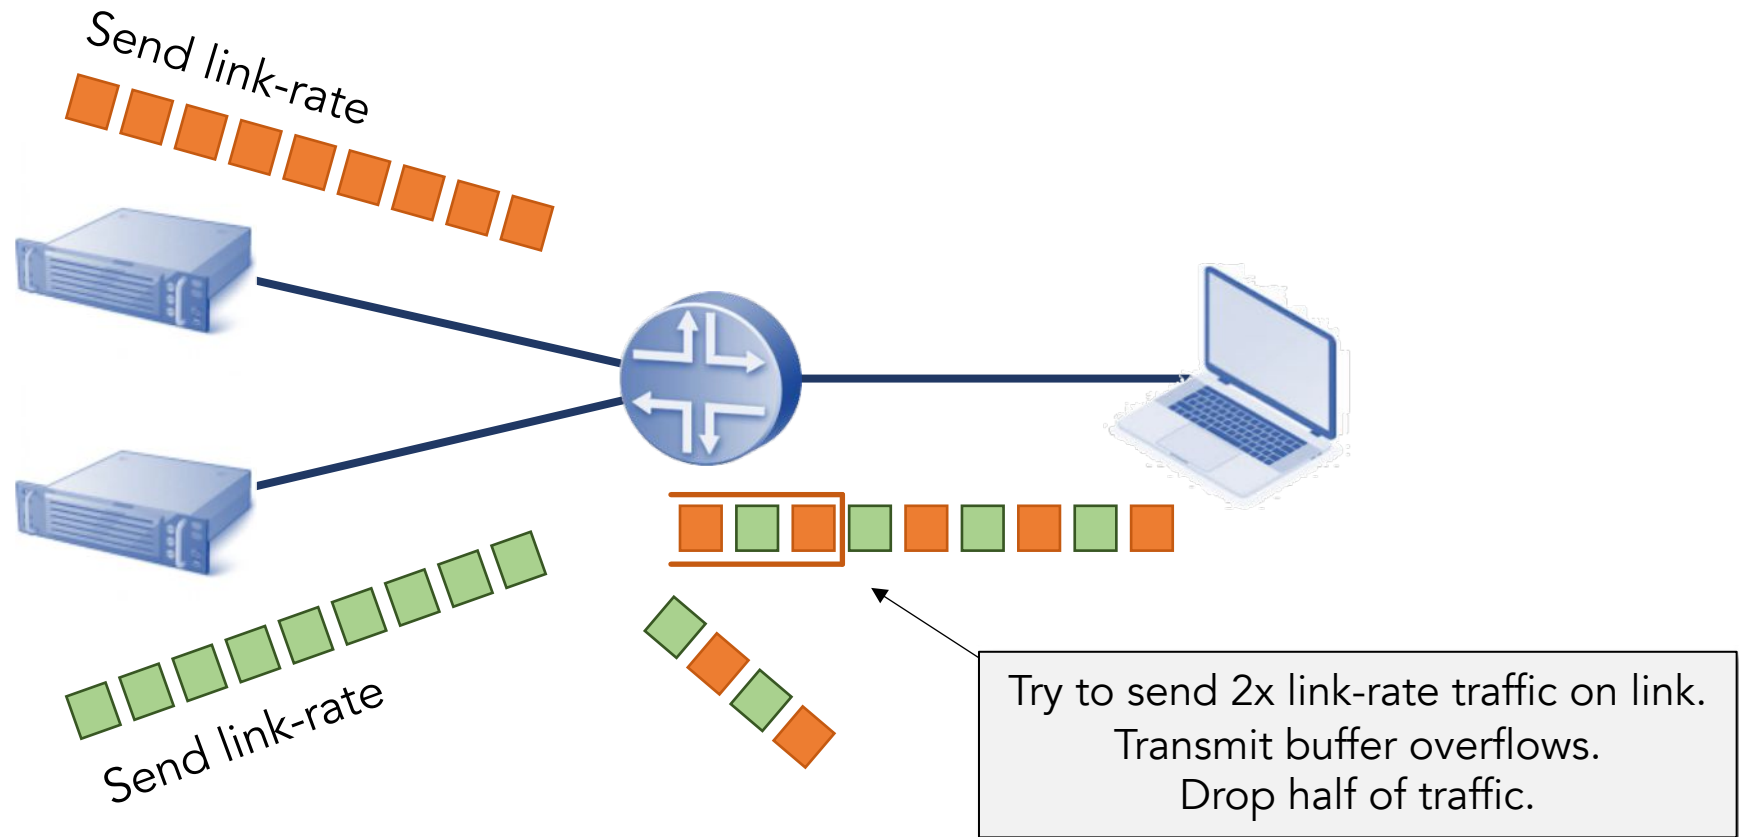
Bandwidth-delay product



Window size must be \geq bandwidth * round-trip delay to fully utilize available bandwidth.
In this example window size (= receive buffer size) is 100 MB = 800 Mbit
Only 32% of available bandwidth is used.

Congestion

Make sure senders do not produce data faster than network can carry it.



Congestion control

- Congestion **avoidance**: avoid congestion in the first place
- Congestion **mitigation**: deal with congestion when it does occur
- Worst case scenario: congestive **collapse**
 - Effective capacity of the network to carry useful traffic drops to zero
 - Retransmissions due to excessive drops and queueing delays overwhelm network
 - More likely to happen when capacity is low compared to demand (as it will be in early quantum networks!)
 - This actually happened in the early internet: in 1986 the capacity collapsed from 32 Kbps to 40 bps (*not* Kbps)

<https://blog.acolyer.org/2015/05/21/congestion-avoidance-and-control/>

<https://tools.ietf.org/html/rfc896>


Congestion avoidance

- Congestion avoidance: mechanisms to avoid congestion in the first place
- General principle: send traffic more slowly when congestion occurs
- **TCP congestion control**
 - TCP tries to estimate maximum rate before congestion occurs
 - TCP sender dynamically congestion window based on presence or absence of drops
 - Extremely difficult to get right: Reno, Vegas, BIC, CUBIC, and many more algorithms
 - Assumes everyone plays fair
- **Backwards Explicit Congestion Notification (BECN)**
 - Router sends message to source to slow down when it observes congestion
- **Forward Explicit Congestion Notification (FECN)**
 - Set bits in forwarded IP packet: Congestion Encountered (CE)
 - Requires cooperation from the transport layer to slow down

Congestion mitigation


- Active Queue Management (AQM):
 - Mechanism for dealing with congestion when it does occur
 - Queueing only happens if there is contention = congestion
- Decide which packets are most important
 - Classification at the edge
 - Mark classification result into QoS / precedence bits in packet header
- Decide which packet to service next in the send queue
 - Hierarchical schedulers
 - Shapers
- Decide when to start dropping packets and which packet to drop
 - Tail-drop
 - Random Early Drop (RED)
 - Drop out-of-profile packet before in-profile packets

Hierarchical queue scheduler



"Congestion control" in quantum networks

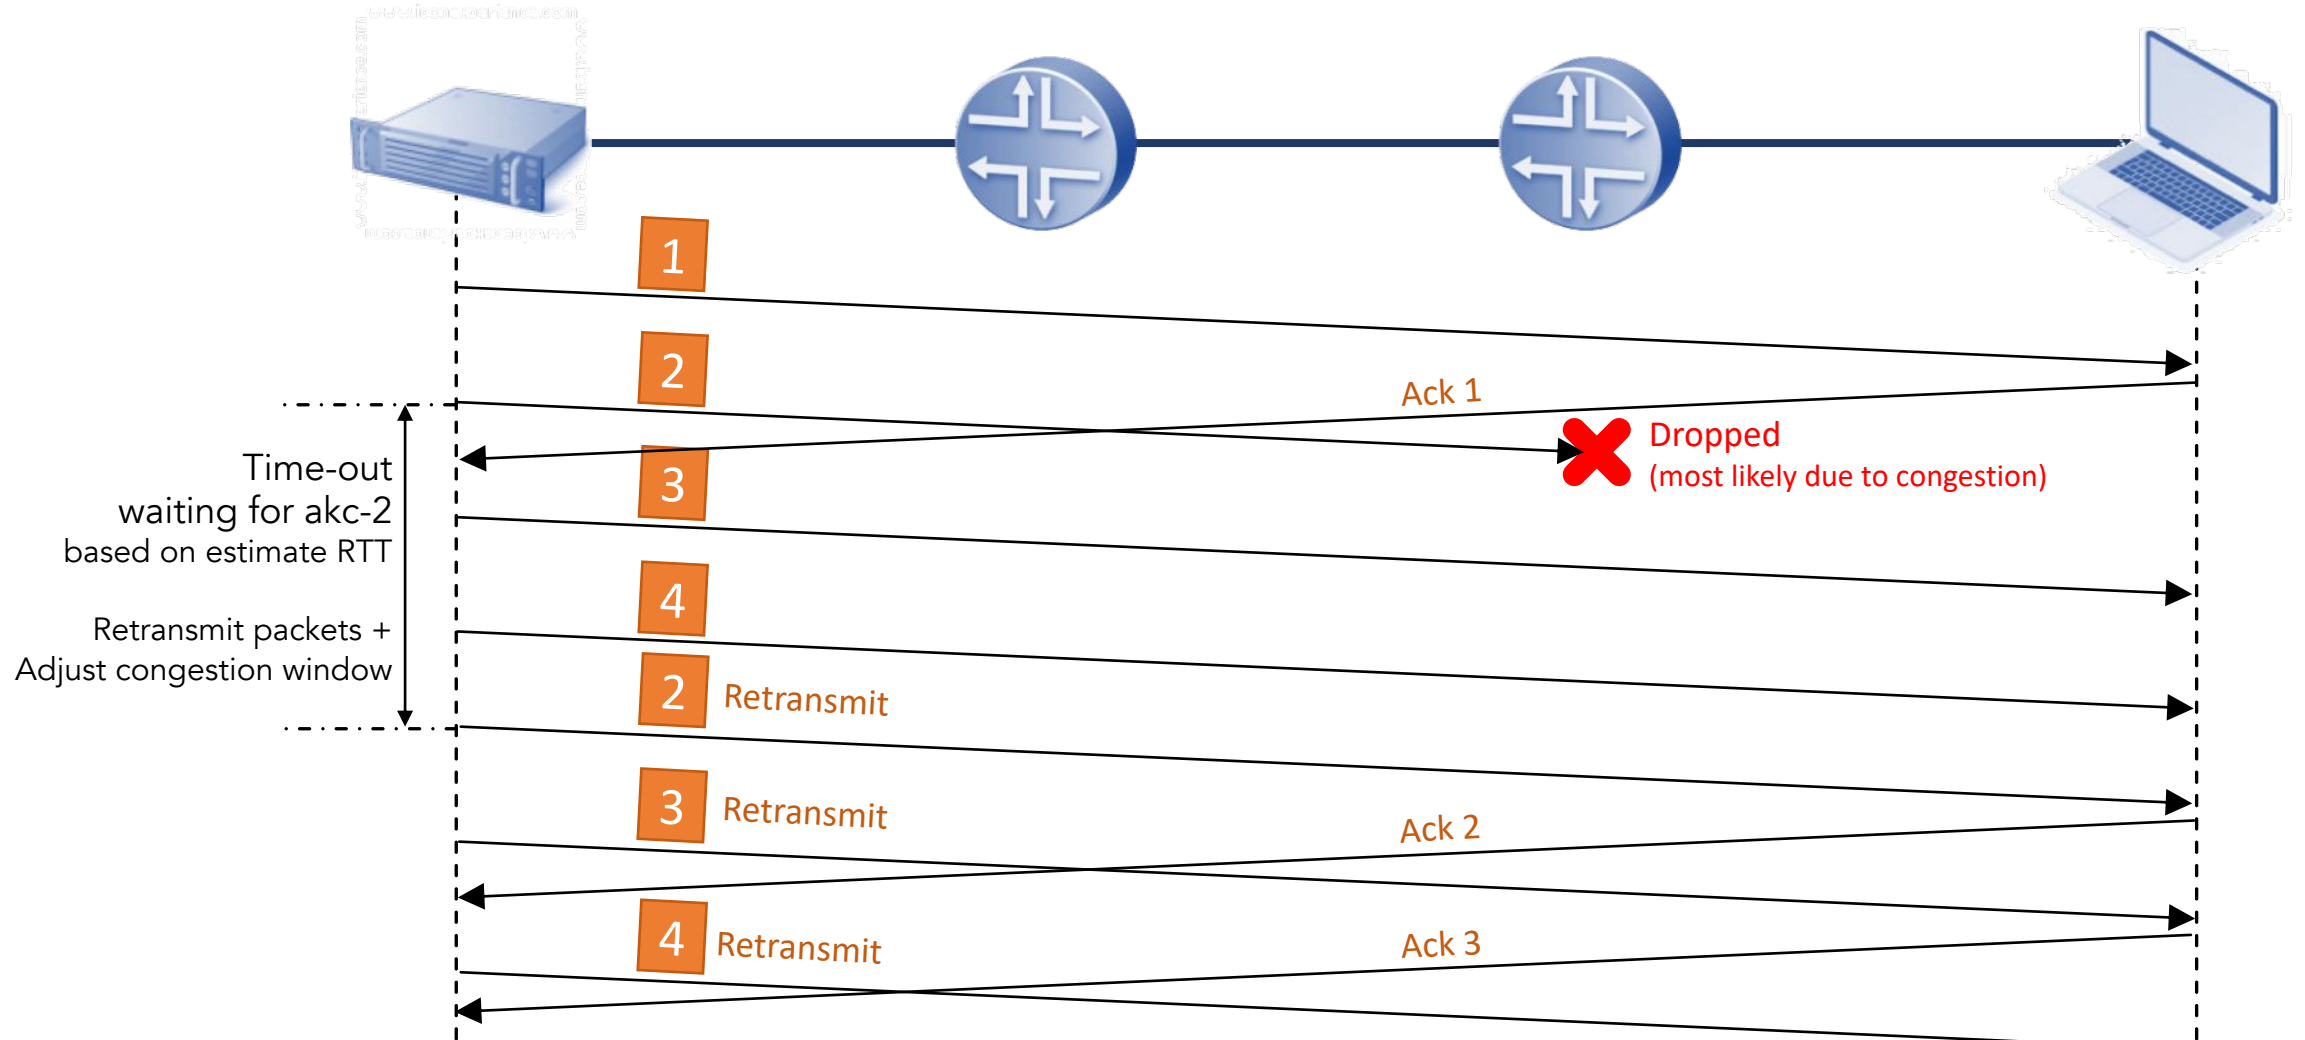
End-points cannot get more end-to-end Bell pairs per second (Bpps) than network can produce them based on bottleneck link or router.



Bottleneck link.
Cannot achieve desired end-to-end entanglements if $X+Y > 10$
(or even less considering non-deterministic swaps and distillation)

Reliable transport: error detection and recovery

This is an example of the go-back-N error recovery



Error detection and recovery mechanisms

- Checksums
 - To detect transmission errors
 - Classical networks are so reliable that the vast majority of drops are due to congestion
- Automatic Repeat Request (ARQ)
 - Receiver sends an acknowledgement (ack) when it has correctly received the data
 - The sender re-transmits the data when it does not receive the ack
 - Not an option for quantum networks (requires that the sender keeps a copy until the data is acknowledged, in case the data needs to be retransmitted).
 - Go-back-N: resend all data starting at the dropped packet
 - Selective repeat: only resend the data that was dropped (requires selective ack)
- Forward Error Correction (FEC)
 - When error rate is too high (noisy links)
 - When retransmission delay is not acceptable (e.g. long latency satellite links)
 - The only option in quantum networks (quantum error correcting codes)

What is so difficult about this?

- Flow control, congestion control, and error control are separate but closely related concepts
- Most transport protocols don't have clean separation of concerns
- For example, in TCP absence of ACK can mean many things:
 - Flow control: receiver wants us to slow down
 - Error: data packet or ACK was dropped
 - Congestion: data packet or ACK was delayed in queue
- Over time TCP has been enhanced:
 - Selective acknowledgement (SACK)
 - Timestamps

Flow control
Congestion control
Error control
in the control plane

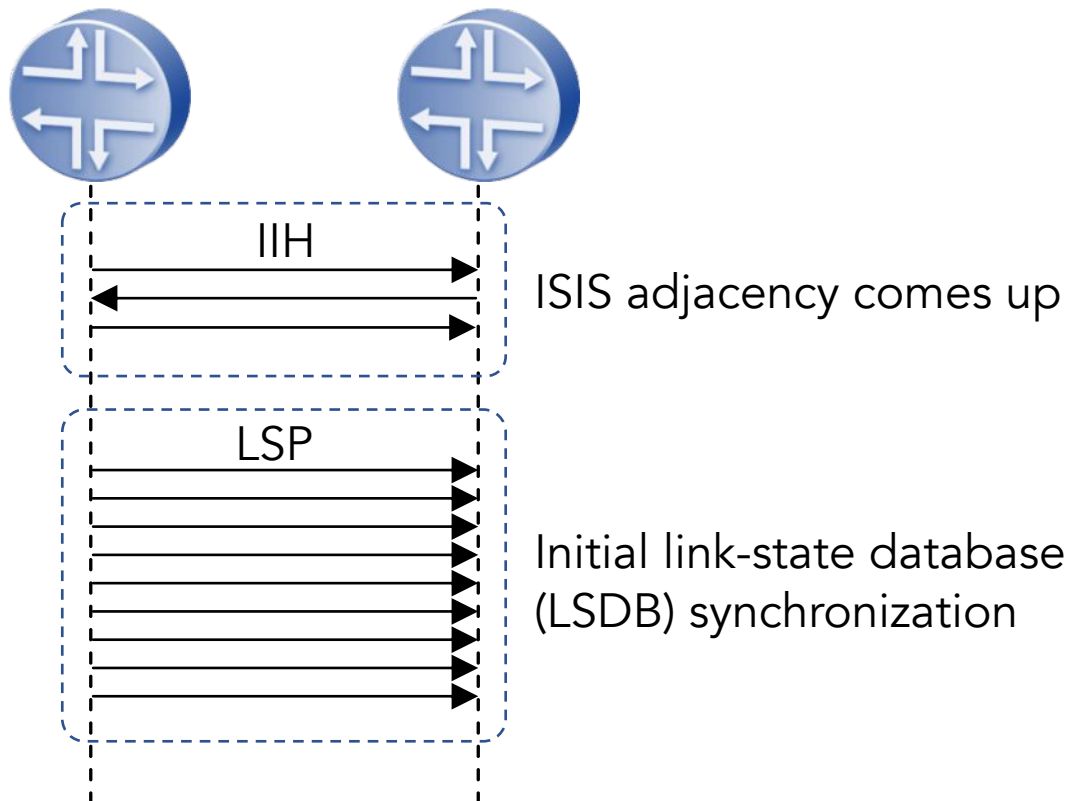
Flow control matters in the control plane

- Routing and signaling protocols exchange large volumes of data
- Sender must not overwhelm receiver
- Flow control approaches in routing and signaling protocols:
 - Fixed pacing
 - Flow control mechanism built into the protocol itself
 - Rely on TCP flow control

Example: ISIS flow control

Initial database synchronization sends large volume of link state packets (LSPs)

Potentially tens of thousands of LSPs



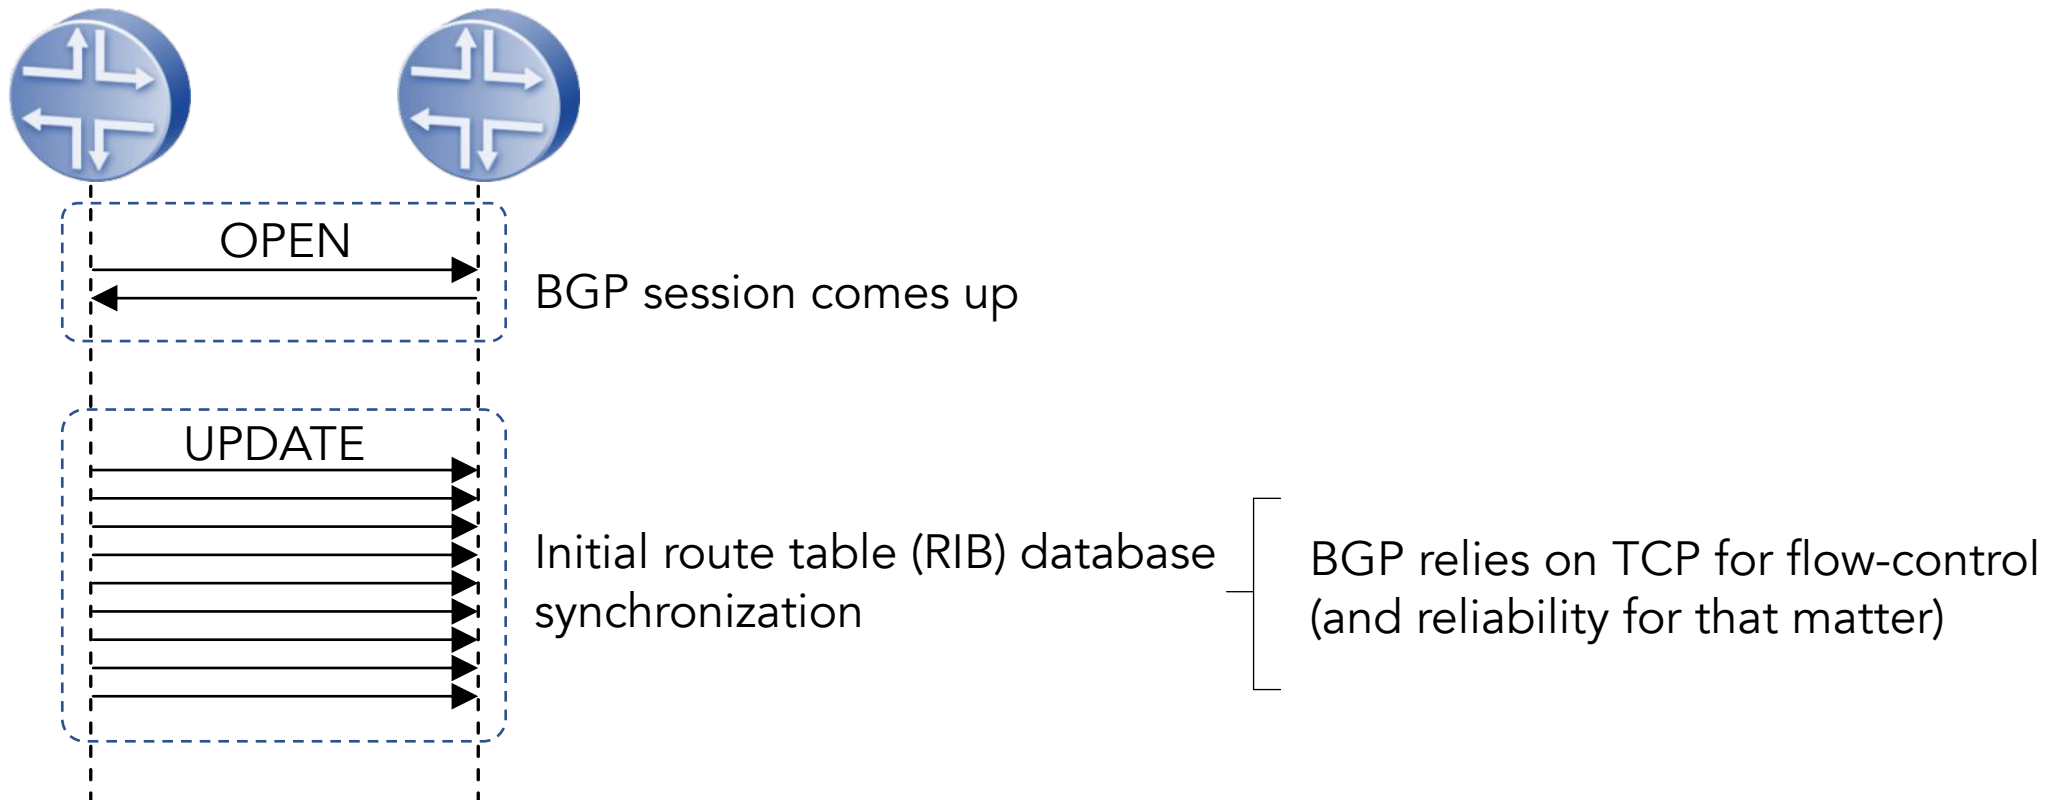
Pacing parameters in standard,
e.g. `minimumLSPTransmissionInterval`

Recent proposals to dynamically control
intervals in ISIS protocol.

Example: BGP flow control

Initial database synchronization sends large volume of routes (UPDATEs)

Often well above 1 million routes



Congestion control matters in the control plane

- Control packets have strict priority over user packets
- Large volume of control packets can still cause congestion
- Congestive collapse of control protocols is not rare
 - Flapping ISIS adjacencies self-reinforcing feedback loop
- Solutions:
 - Prioritize important control plane packets (mainly hellos)
 - Multiple streams (QUIC, SCTP)
 - Off-load liveness detection to separate protocol (BFD)

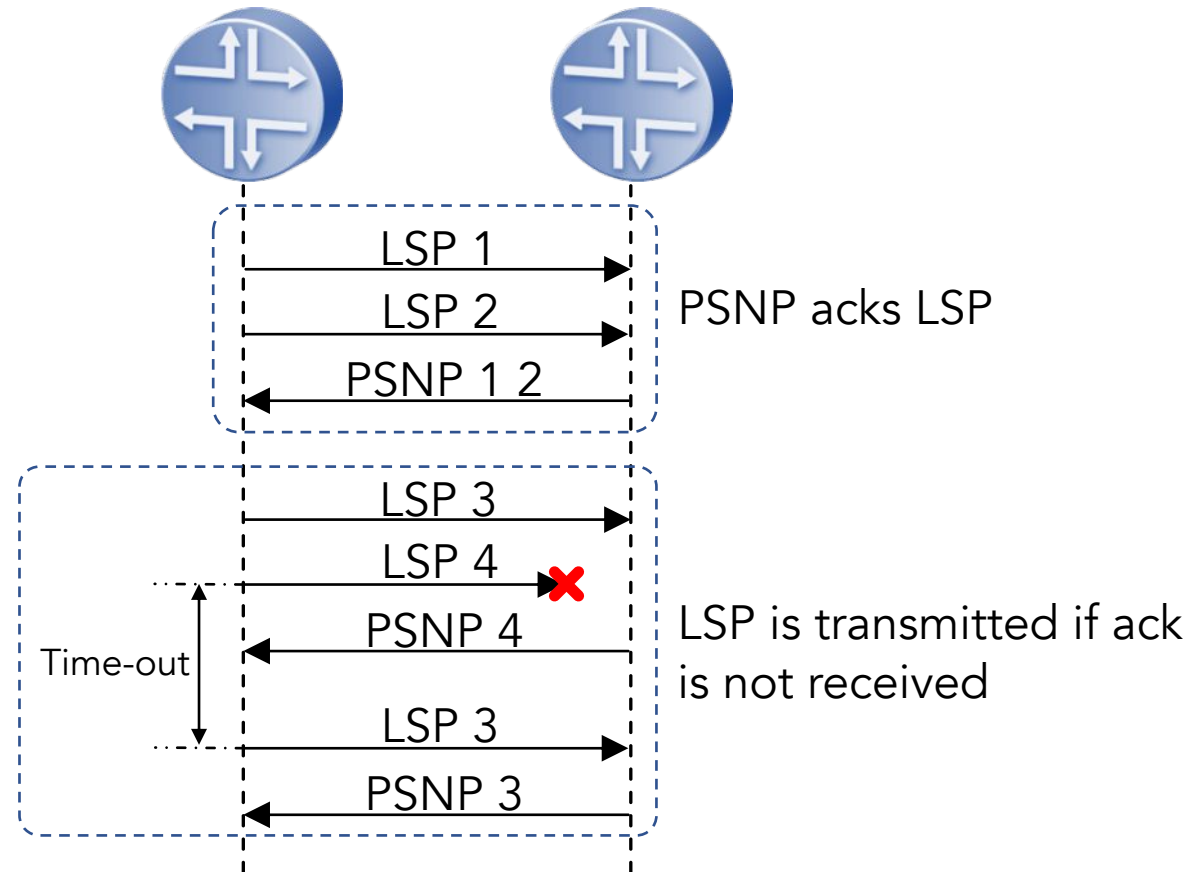
Error control matters in the control plane

- Control plane packets also get dropped
- Need reliability mechanism (error control) in routing and signaling protocols
- Approaches:
 - Separate mechanism built-in to protocol itself: ISIS, OSPF, ...
 - Rely on transport (TCP) for reliability: BGP, LDP, ...

Example: ISIS error control

Partial Sequence Number PDU (PSNP) is used to ack LSP

There is also a Complete Sequence Number PDU (CSNP) that we don't discuss here



Take-aways for quantum control protocols

- You need to worry about flow control
- You need to worry about congestion control
- You need to worry about error control
- Not just in the data plane, but also in the control plane
- Avoid re-inventing the wheel: rely on layer 4 transport when possible

This is somewhat contentious; it is my opinion that not everyone agrees with.

Soft state versus hard state

Soft state vs hard state

Sender wants to advertise some data (e.g. a route) to receiver.

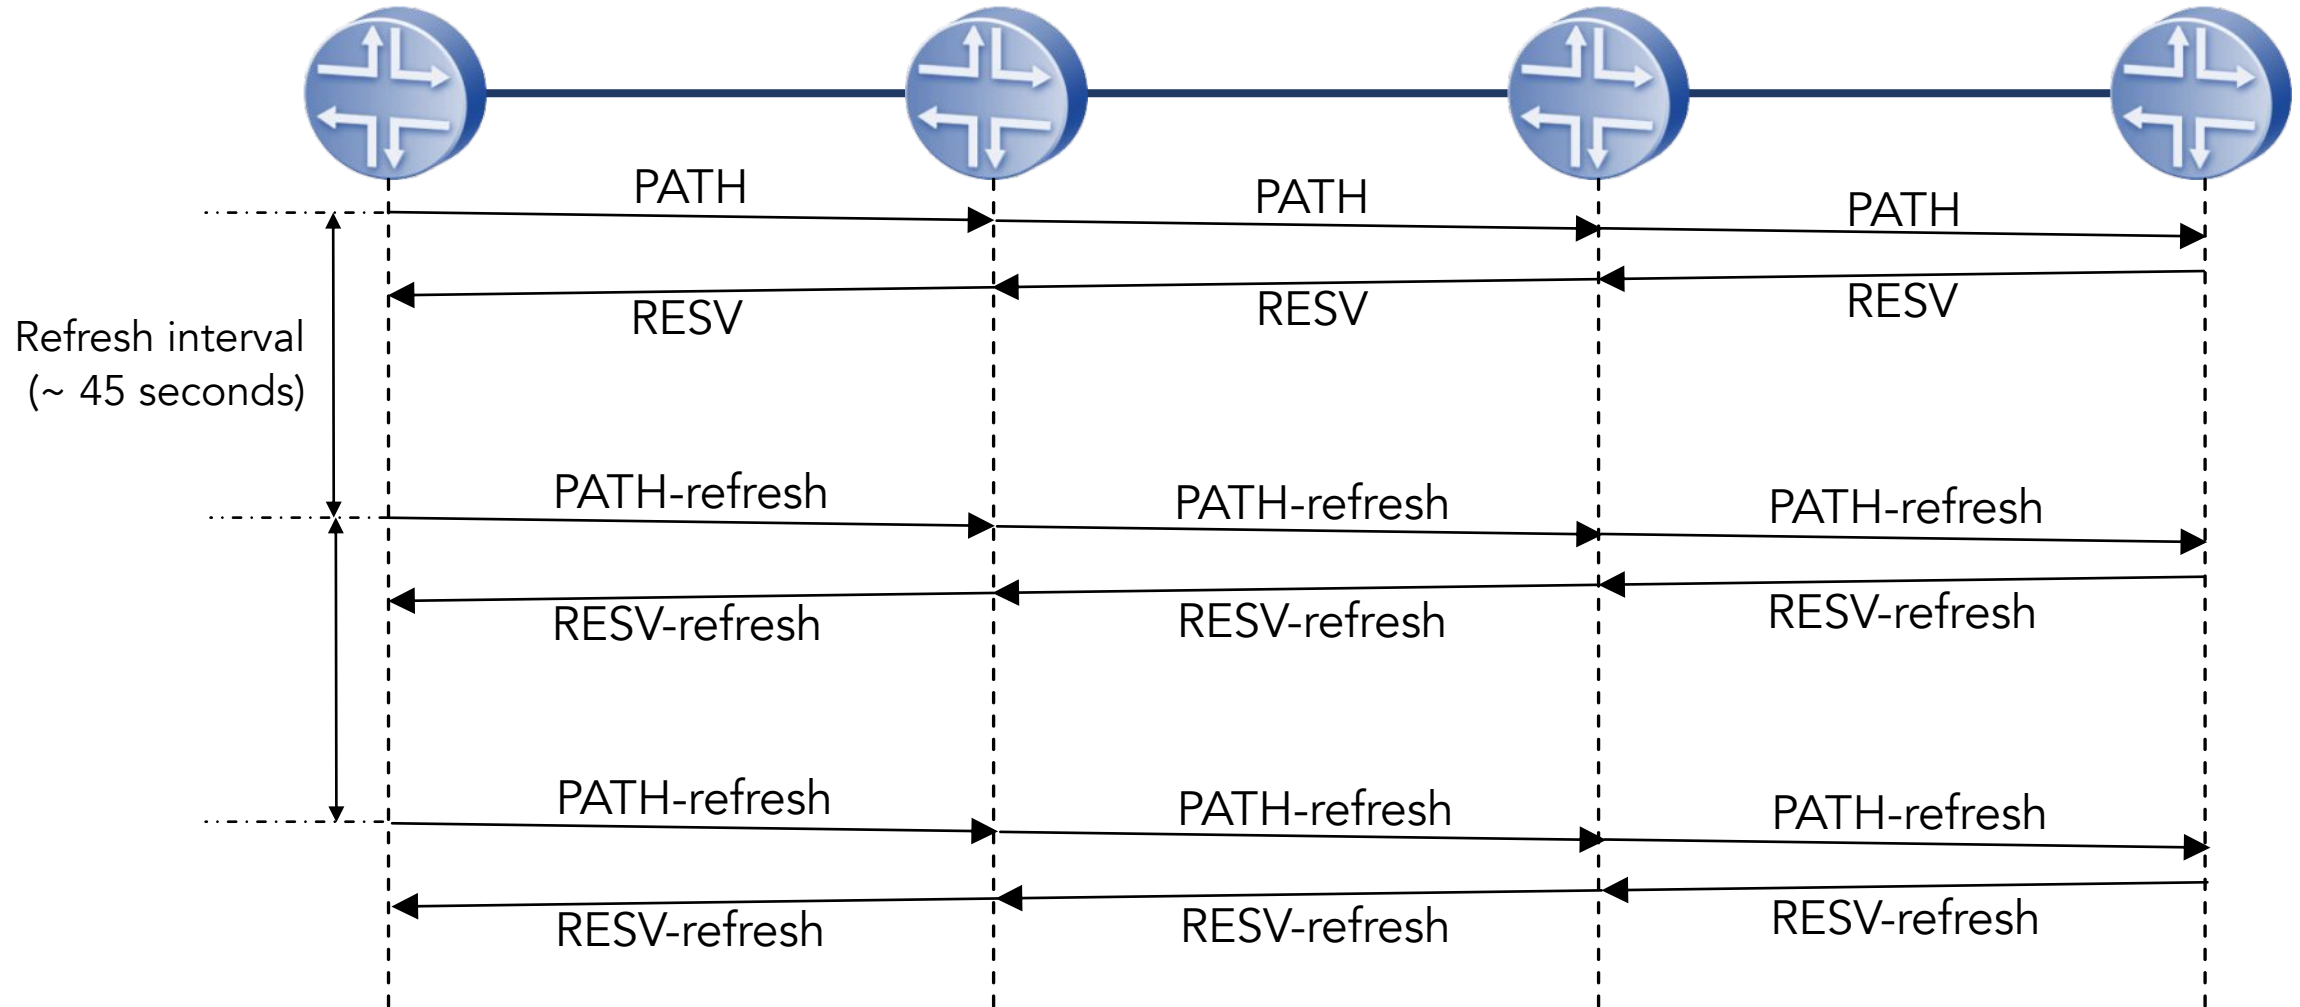
Soft state protocol

- A.k.a. Periodic protocol
- Sender initially sends data.
- Receiver puts data in its database.
- Sender periodically resends data (refresh).
- If receiver doesn't receive data anymore, it removes data from its database after time-out.
- Heavy load due to periodic transmission.
- Fallacy: simpler cleanup when sender crashes.


Hard state state protocol

- Runs protocol over reliable transport connection.
- Sender initially sends data.
- Receiver puts data in its database.
- No periodic re-transmissions of data.
- Receiver keep data in database until (a) sender explicitly withdraws data or (b) transport connection is disconnected.
- No traffic if no changes in the database.
- Needs reliable transport connection and keep-alive mechanism (typically TCP)

Example: soft state in RSVP



The problem with soft state



Typical network topology




Core router



Edge router

Dual-homed to core router

The problem with soft state



Typical network topology



Core router



Edge router

Dual-homed to core router


— Full mesh of LSPs

From every edge router to every other edge router

Traffic engineered according to bandwidth-demand matrix

This diagram show logical topology of LSPs (not physical topology)

The problem with soft state



Core router

Number of LSPs is $O(N^2)$

When N is number of edge routers

Refresh frequency is $O(N^2)$

Real-life example

Number of edge routers: 500

Number of LSPs: $500 \times 499 = 249,500$

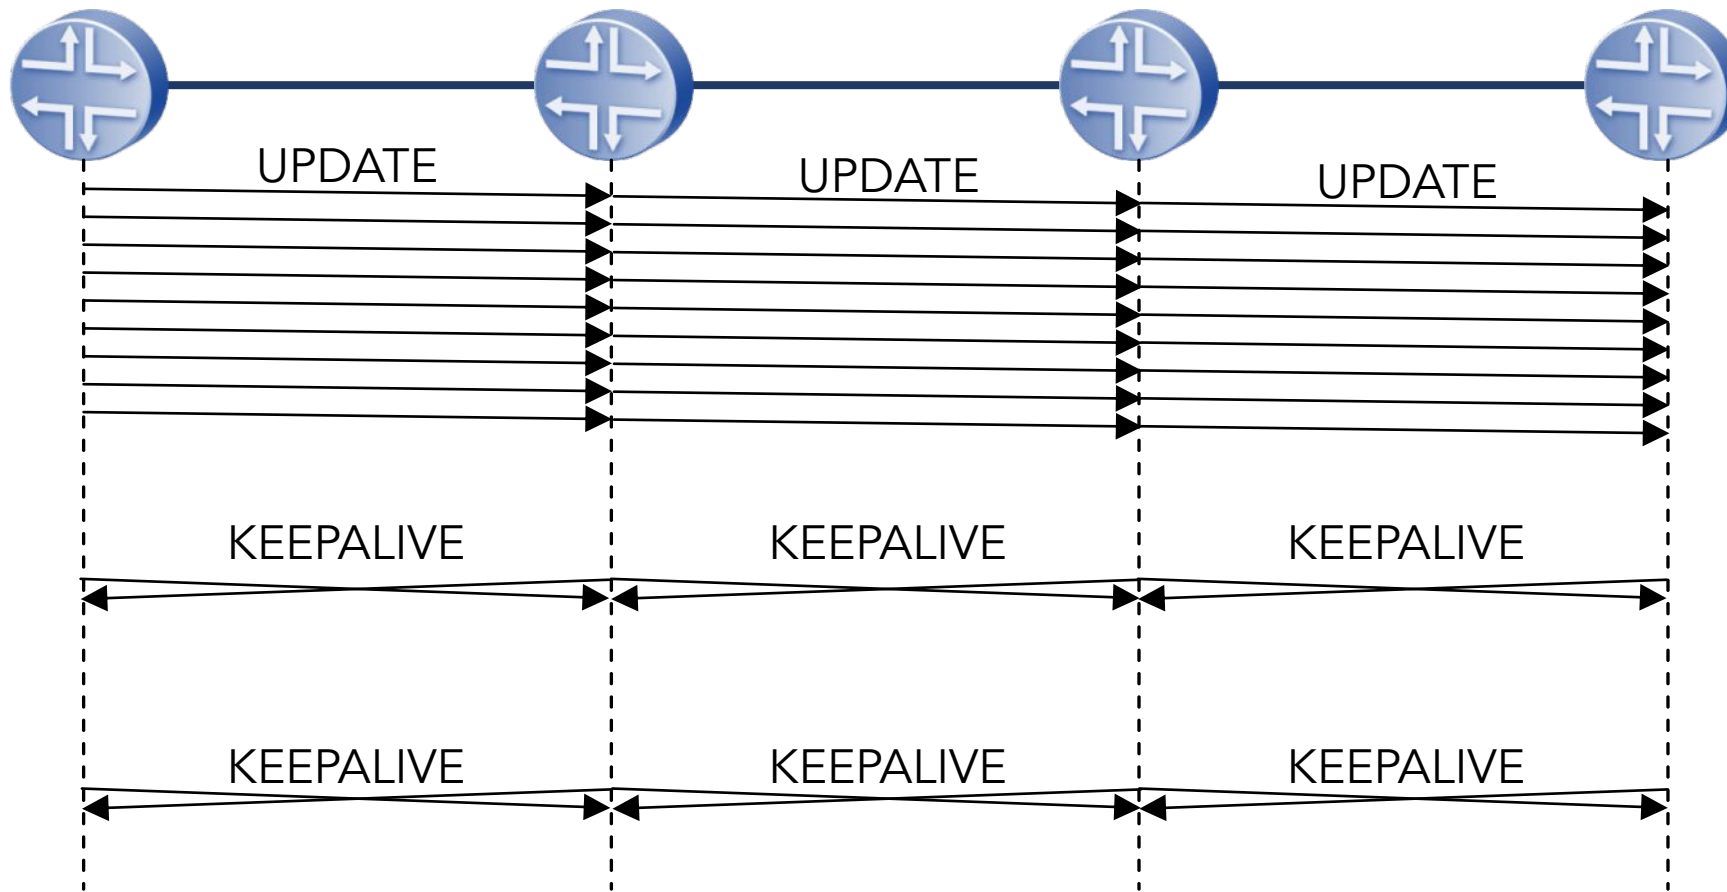
Refresh interval = 45 seconds

Refresh traffic: 11,088 messages/sec

(Simplifying assumption that all LSPs pass through the same core router)

This is not feasible in the control plane

Example: hard state in BGP



UPDATE creates **hard state**

Can be > 1 million routes

Periodic refresh not feasible

Route stays valid until:

- Explicitly **WITHDRAW** or
- TCP connection breaks

Periodic **KEEPALIVE**

Needed to check TCP connection

Just a single message

Independent of number of routes

Network meltdown due to
Run-away replication

Packet replication loops

Packet replication

Making multiple copies of a packet when you forward it (e.g. multicast)



+

Forwarding loop



=

Network meltdown



Packet replication loops

Packet replication

Making multiple copies of a packet when you forward it (e.g. multicast)



+

Forwarding loop



+

No time to live (TTL) mechanism

Packets can loop forever without ever being removed due to TTL expiry

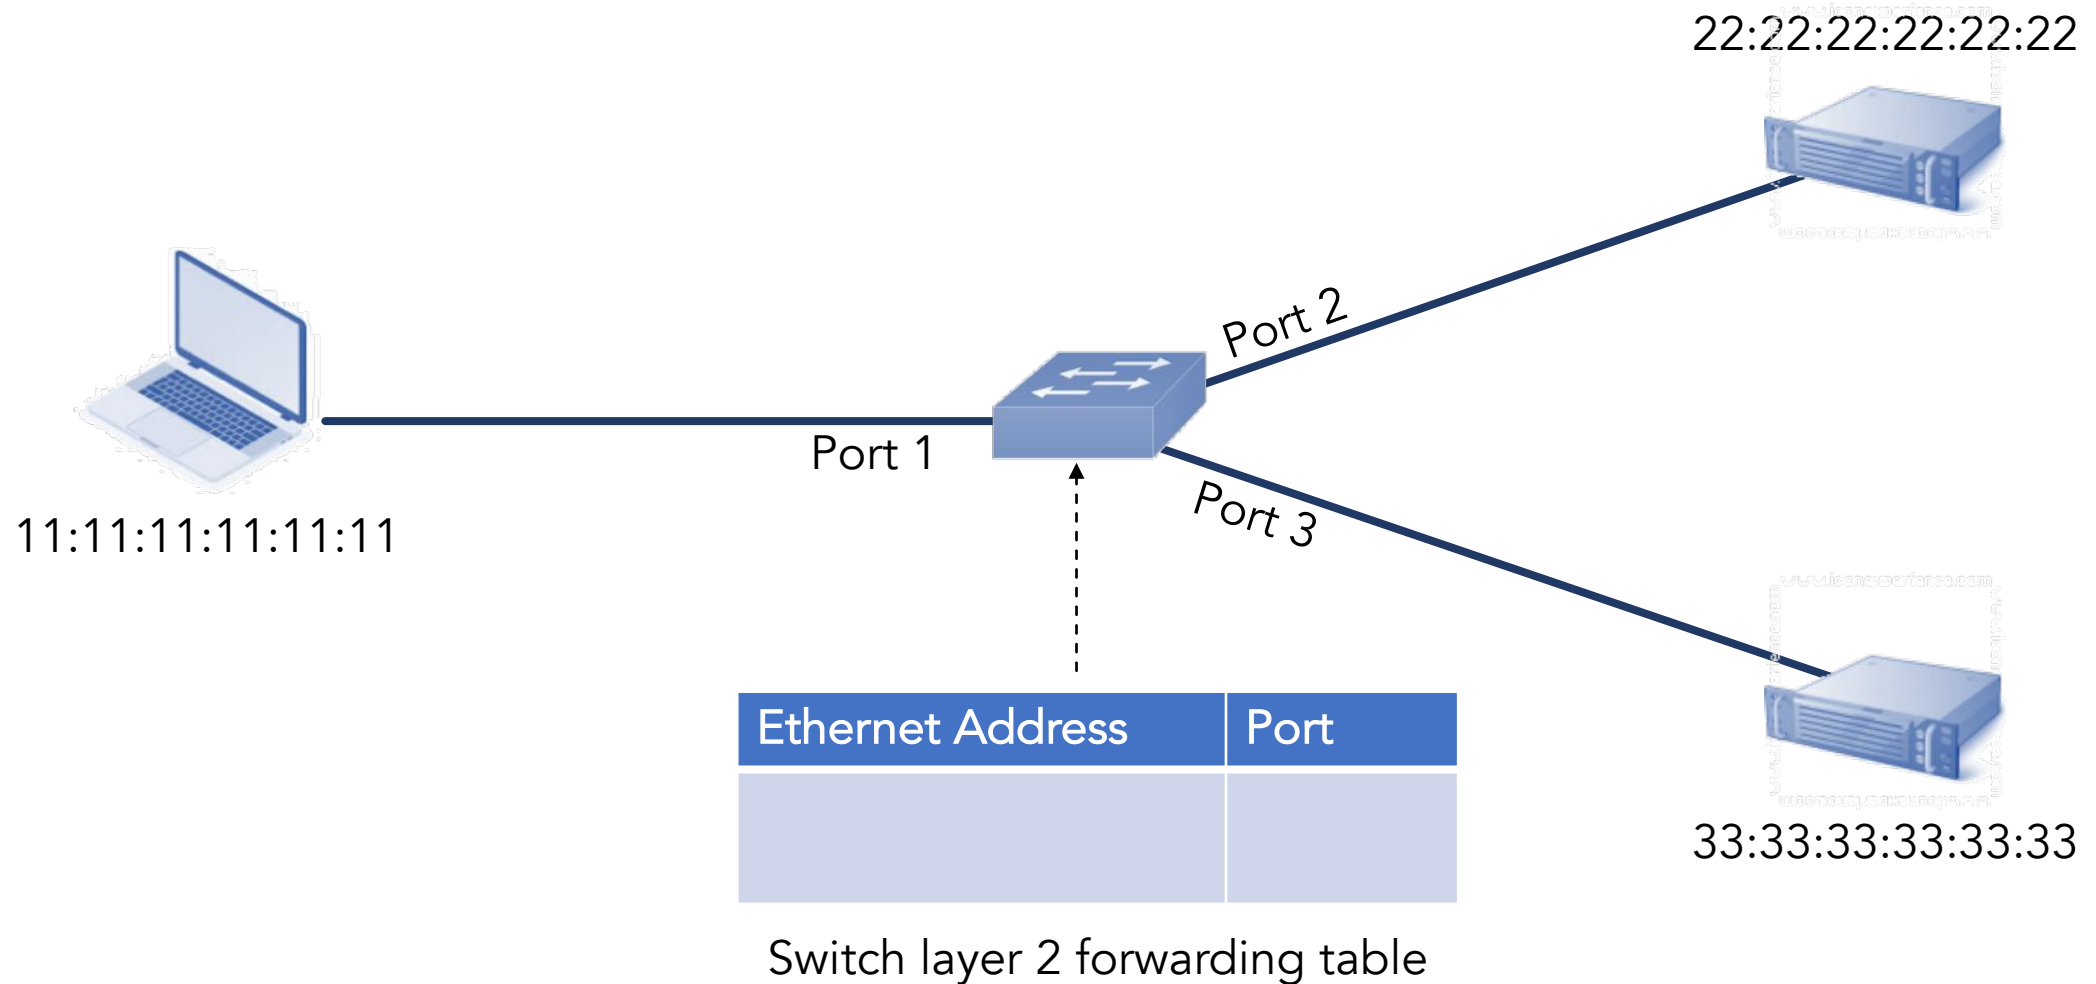


=

Catastrophic network meltdown

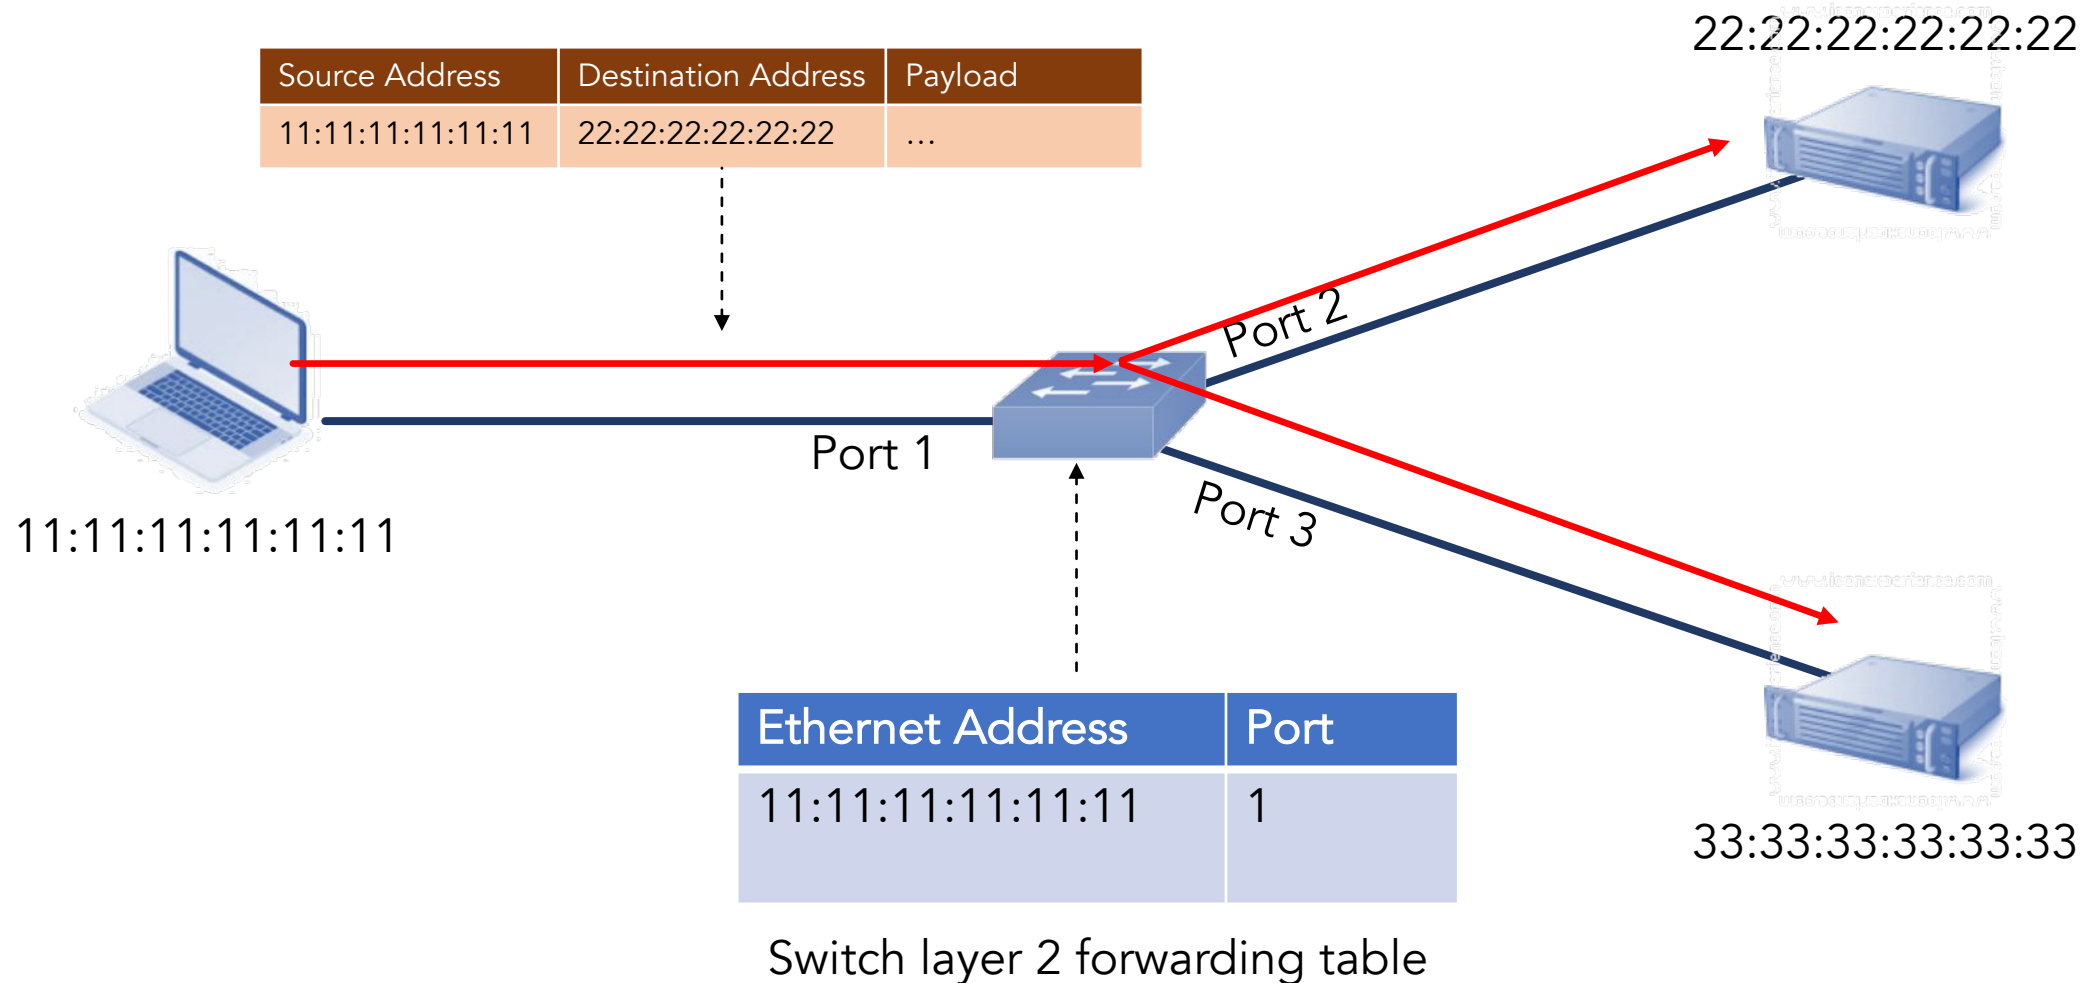


Ethernet address learning



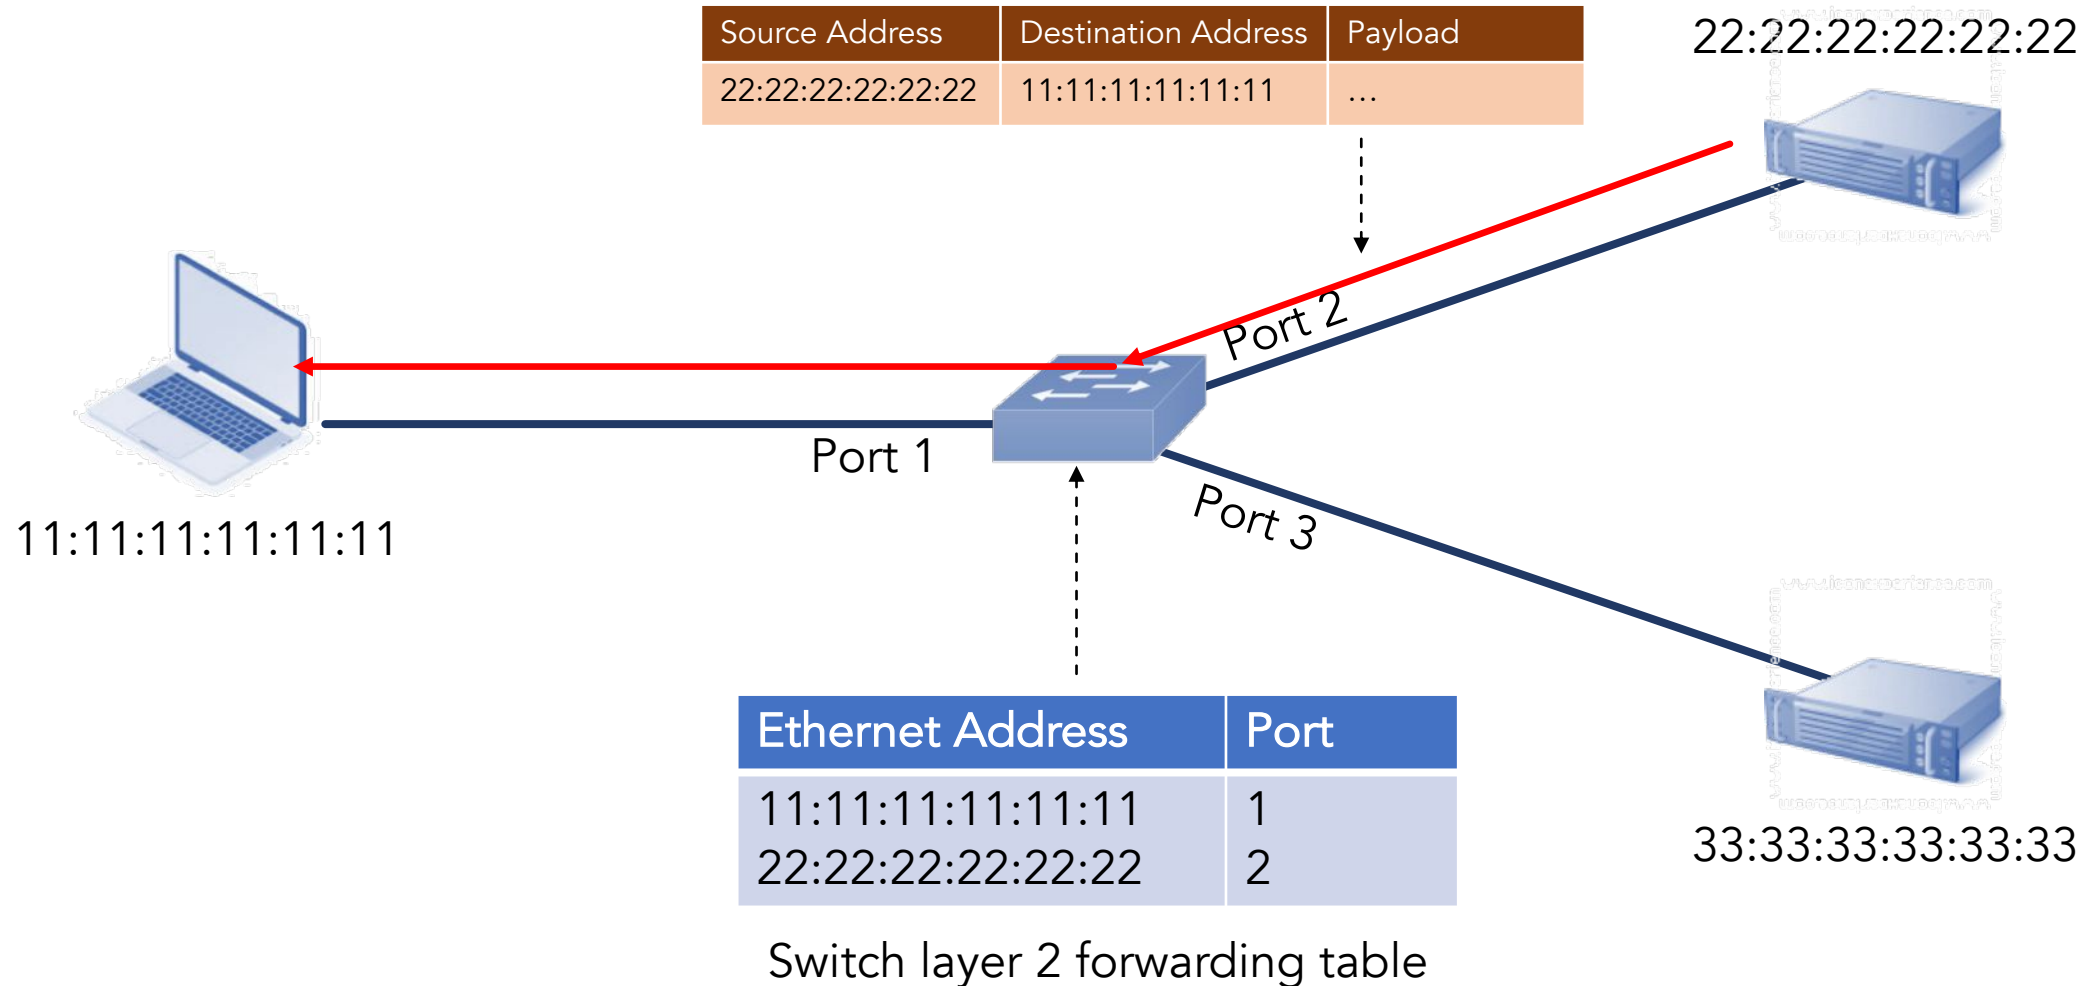
Ethernet address learning

Unknown destination address: flood and learn source




Ethernet address learning

Known destination address: unicast and learn source




Ethernet meltdown due to forwarding loop

Packet loop around the network at line rate
Number of packets grows exponentially



Spanning Tree Protocol (STP)

STP removes loops by blocking links to form a spanning tree.

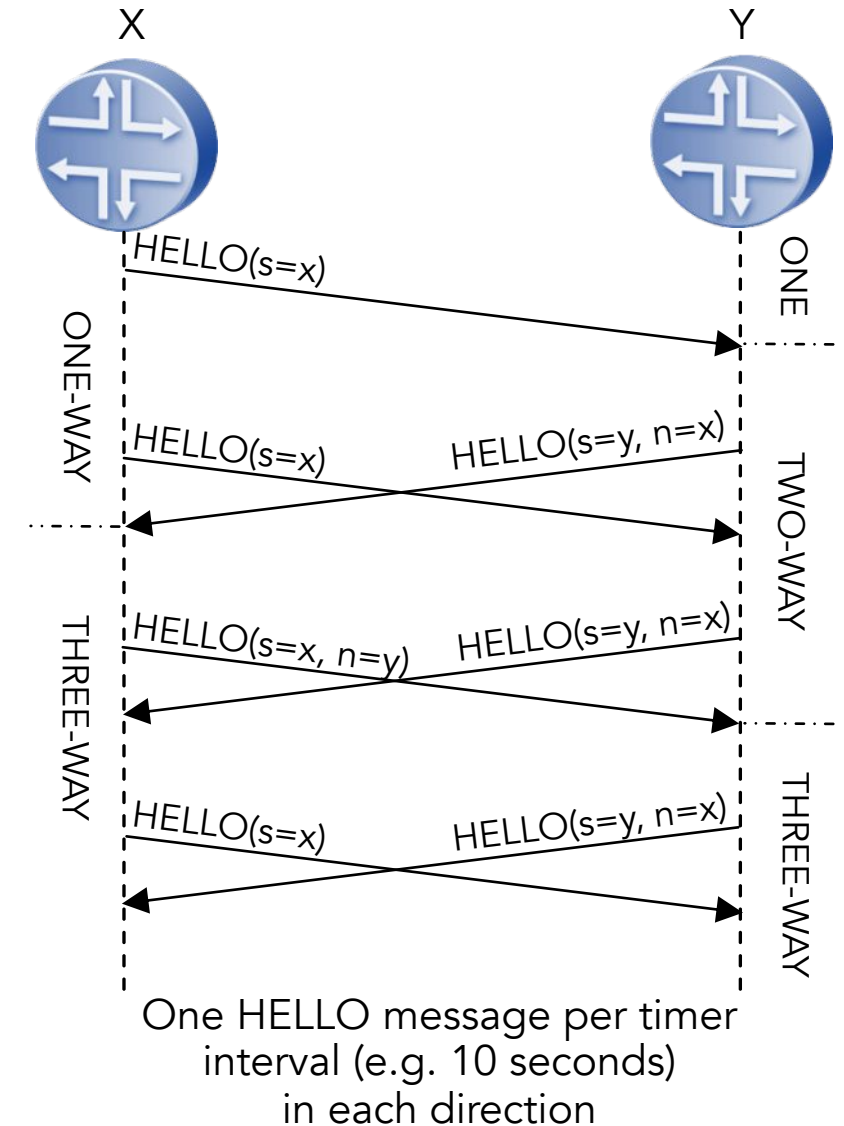
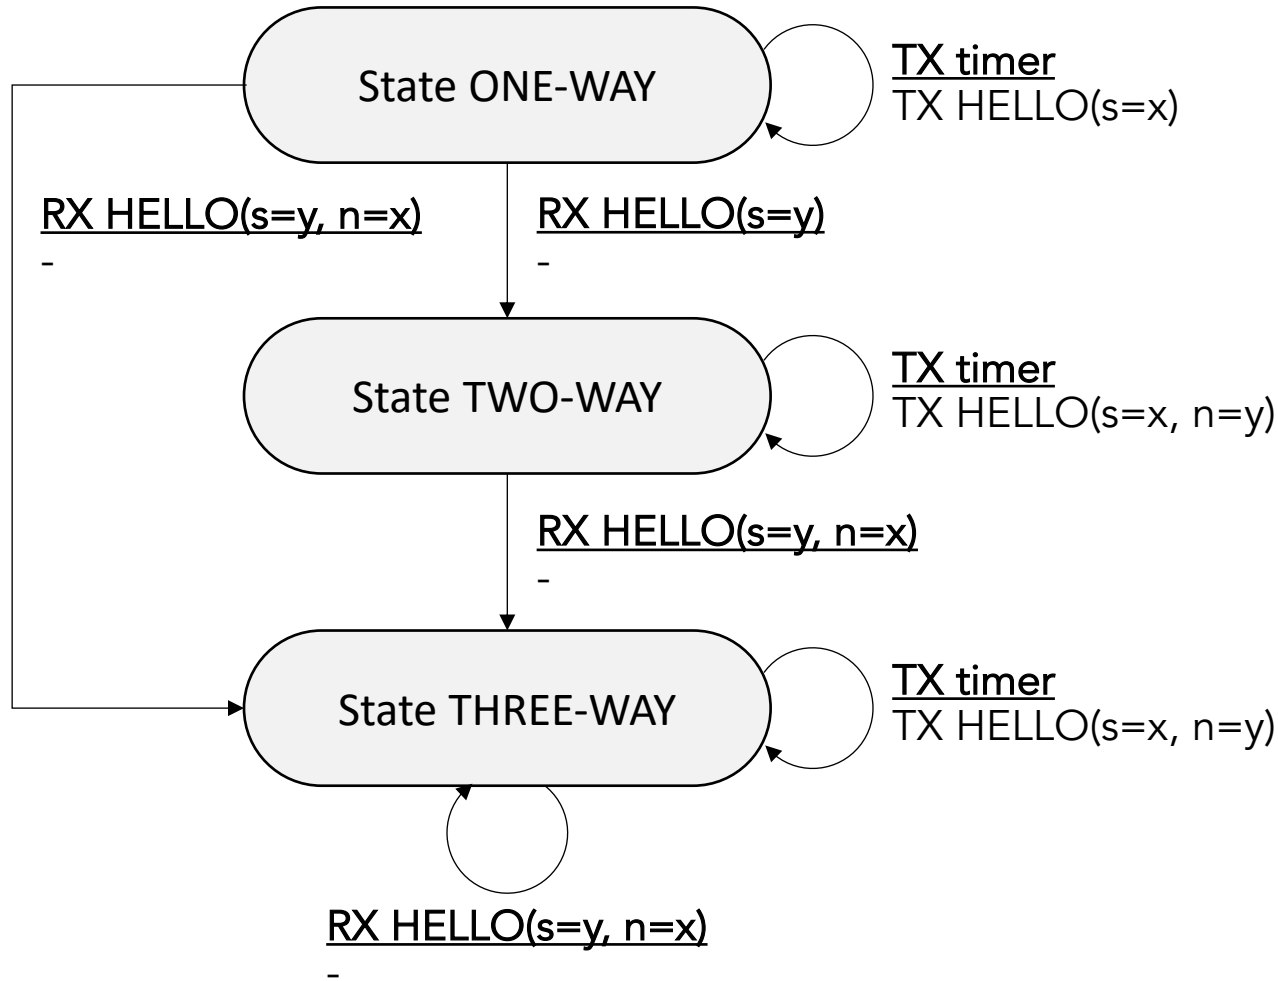


Wastes a lot of link capacity (no multipath)


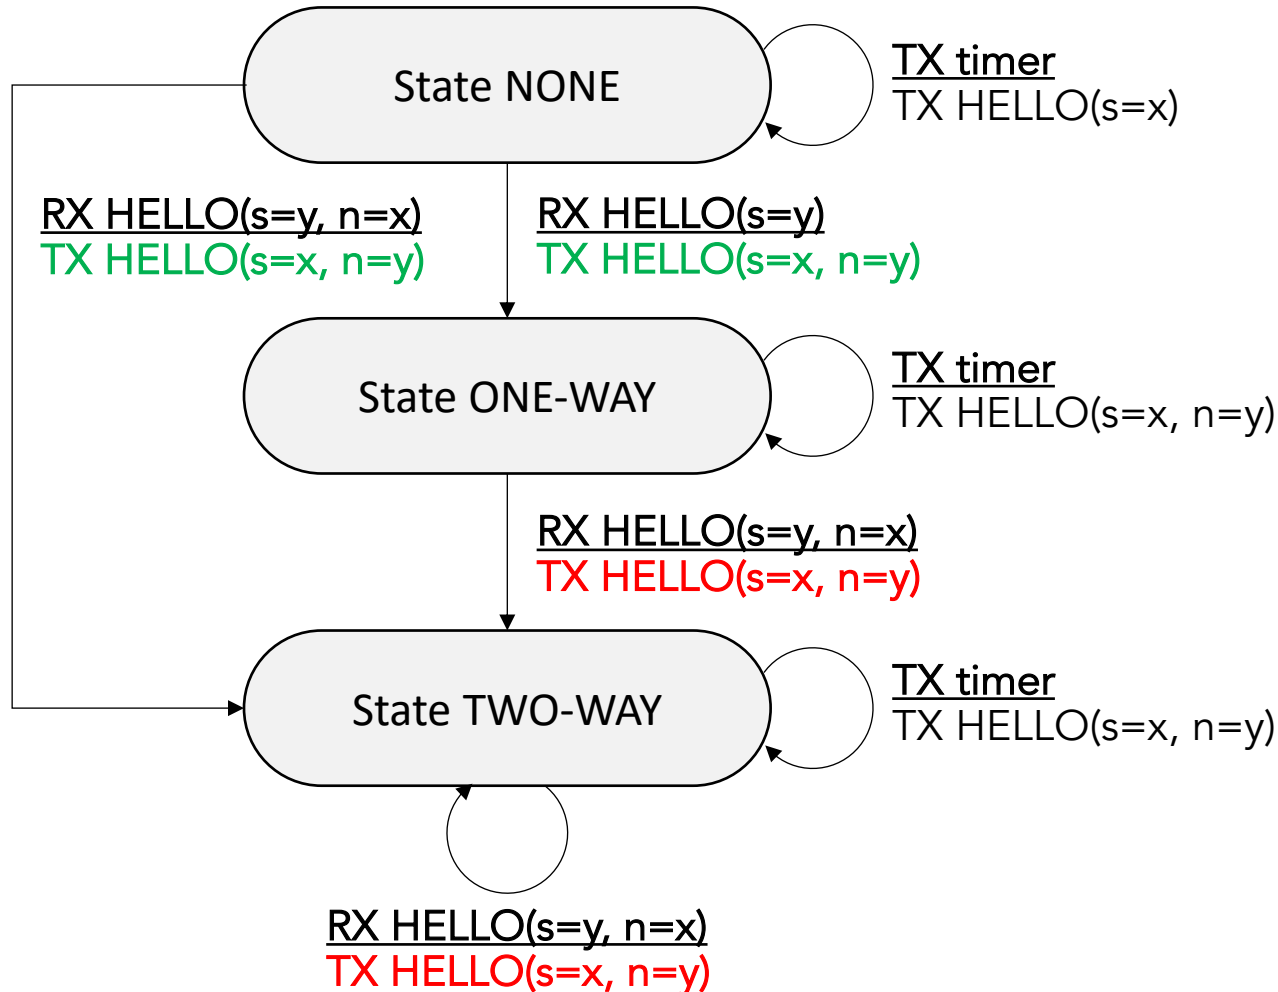
Network melts down if STP doesn't converge perfectly.

Network meltdown due to
Run-away state machine

Typical 3-way adjacency finite state machine



Run-away 3-way adjacency FSM




Real life example: early draft of RIFT


```
000143 2018-08-11 08:53:02,032 Transition NEIGHBOR_OFFER [UPDATING_CLIENTS] > update_or_remove_offer [None]
000144 2018-08-11 08:53:02,032
000145 2018-08-11 08:53:02,032 RX LIE ProtocolPacket(header=PacketHeader(level=0, minor_version=0, major_version=11, sender=1), content=PacketContent(tie=None, tire=None, lie=LIEPacket(name='node
000146 2018-08-11 08:53:02,033 Push LIE_RECEIVED
000147 2018-08-11 08:53:02,033 RX LIE ProtocolPacket(header=PacketHeader(level=0, minor_version=0, major_version=11, sender=1), content=PacketContent(tie=None, tire=None, lie=LIEPacke
000148 2018-08-11 08:53:02,033 Push LIE_RECEIVED
000149 2018-08-11 08:53:02,033 Push NEIGHBOR_OFFER
000150 2018-08-11 08:53:02,033
000151 2018-08-11 08:53:02,033 Transition LIE_RECEIVED [TWO_WAY] > process_lie,MULTIPLE_NEIGHBORS [None]
000152 2018-08-11 08:53:02,034 TX LIE ProtocolPacket(header=PacketHeader(level=1, minor_version=0, major_version=11, sender=2), content=PacketContent(tie=None, tire=None, lie=LIEPacket(name='node
000153 2018-08-11 08:53:02,034 Transition MULTIPLE_NEIGHBORS [TWO_WAY] > cleanup,send_lie [ONE_WAY]
000154 2018-08-11 08:53:02,034 Push NEIGHBOR_OFFER
000155 2018-08-11 08:53:02,034
000156 2018-08-11 08:53:02,034 Transition LIE_RECEIVED [TWO_WAY] > process_lie,MULTIPLE_NEIGHBORS [None]
000157 2018-08-11 08:53:02,034 TX LIE ProtocolPacket(header=PacketHeader(level=2, minor_version=0, major_version=11, sender=3), content=PacketContent(tie=None, tire=None, lie=LIEPacke
000158 2018-08-11 08:53:02,035 Transition MULTIPLE_NEIGHBORS [TWO_WAY] > cleanup,send_lie [ONE_WAY]
000159 2018-08-11 08:53:02,035 Transition NEIGHBOR_OFFER [UPDATING_CLIENTS] > update_or_remove_offer,LOST_HAL [None]
000160 2018-08-11 08:53:02,035 Transition LOST_HAL [UPDATING_CLIENTS] > start_timer_on_lost_hal,HOLD_DOWN_EXPIRED [HOLDING_DOWN]
000161 2018-08-11 08:53:02,035 Transition HOLD_DOWN_EXPIRED [HOLDING_DOWN] > purge_offers,stop_hold_down_timer,level_compute,COMPUTATION_DONE [COMPUTE_BEST_OFFER]
000162 2018-08-11 08:53:02,035 Transition COMPUTATION_DONE [COMPUTE_BEST_OFFER] > update_all_lie_fame [UPDATING_CLIENTS]
000163 2018-08-11 08:53:02,035 Transition NEIGHBOR_OFFER [UPDATING_CLIENTS] > update_or_remove_offer,LOST_HAL [None]
000164 2018-08-11 08:53:02,035 Transition LOST_HAL [UPDATING_CLIENTS] > start_timer_on_lost_hal,HOLD_DOWN_EXPIRED [HOLDING_DOWN]
000165 2018-08-11 08:53:02,036 Transition HOLD_DOWN_EXPIRED [HOLDING_DOWN] > purge_offers,stop_hold_down_timer,level_compute,COMPUTATION_DONE [COMPUTE_BEST_OFFER]
000166 2018-08-11 08:53:02,036 Transition COMPUTATION_DONE [COMPUTE_BEST_OFFER] > update_all_lie_fame [UPDATING_CLIENTS]
000167 2018-08-11 08:53:02,036 RX LIE ProtocolPacket(header=PacketHeader(level=2, minor_version=0, major_version=11, sender=3), content=PacketContent(tie=None, tire=None, lie=LIEPacket(name='node3-if1', nmc
000168 2018-08-11 08:53:02,036 Push LIE_RECEIVED
000169 2018-08-11 08:53:02,037 RX LIE ProtocolPacket(header=PacketHeader(level=2, minor_version=0, major_version=11, sender=3), content=PacketContent(tie=None, tire=None, lie=LIEPacket(name='node
000170 2018-08-11 08:53:02,037 Push LIE_RECEIVED
000171 2018-08-11 08:53:02,037
000172 2018-08-11 08:53:02,037 Push NEIGHBOR_OFFER
000173 2018-08-11 08:53:02,037
000174 2018-08-11 08:53:02,037 Transition LIE_RECEIVED [ONE_WAY] > process_lie,NEW_NEIGHBOR [None]
000175 2018-08-11 08:53:02,037 Transition NEW_NEIGHBOR [ONE_WAY] > SEND_LIE [TWO_WAY]
000176 2018-08-11 08:53:02,038 TX LIE ProtocolPacket(header=PacketHeader(level=0, minor_version=0, major_version=11, sender=1), content=PacketContent(tie=None, tire=None, lie=LIEPacket(name='node1-if1', nmc
000177 2018-08-11 08:53:02,038 Transition SEND_LIE [TWO_WAY] > send_lie [None]
000178 2018-08-11 08:53:02,038 Push NEIGHBOR_OFFER
000179 2018-08-11 08:53:02,038
000180 2018-08-11 08:53:02,038 Transition LIE_RECEIVED [ONE_WAY] > process_lie,NEW_NEIGHBOR [None]
000181 2018-08-11 08:53:02,038 Transition NEW_NEIGHBOR [ONE_WAY] > SEND_LIE [TWO_WAY]
000182 2018-08-11 08:53:02,038 TX LIE ProtocolPacket(header=PacketHeader(level=1, minor_version=0, major_version=11, sender=2), content=PacketContent(tie=None, tire=None, lie=LIEPacket(name='node
000183 2018-08-11 08:53:02,039 Transition SEND_LIE [TWO_WAY] > send_lie [None]
000184 2018-08-11 08:53:02,039 Transition NEIGHBOR_OFFER [UPDATING_CLIENTS] > update_or_remove_offer [None]
000185 2018-08-11 08:53:02,039 Transition NEIGHBOR_OFFER [UPDATING_CLIENTS] > update_or_remove_offer,BETTER_HAL [None]
000186 2018-08-11 08:53:02,039 Transition BETTER_HAL [UPDATING_CLIENTS] > stop_hold_down_timer,level_compute,COMPUTATION_DONE [COMPUTE_BEST_OFFER]
000187 2018-08-11 08:53:02,040 Transition COMPUTATION_DONE [COMPUTE_BEST_OFFER] > update_all_lie_fame [UPDATING_CLIENTS]
000188 2018-08-11 08:53:02,040
000189 2018-08-11 08:53:02,040 RX LIE ProtocolPacket(header=PacketHeader(level=0, minor_version=0, major_version=11, sender=1), content=PacketContent(tie=None, tire=None, lie=LIEPacket(name='node
000190 2018-08-11 08:53:02,040 Push LIE_RECEIVED
000191 2018-08-11 08:53:02,041 RX LIE ProtocolPacket(header=PacketHeader(level=0, minor_version=0, major_version=11, sender=1), content=PacketContent(tie=None, tire=None, lie=LIEPacke
000192 2018-08-11 08:53:02,041 Push LIE_RECEIVED
000193 2018-08-11 08:53:02,041 Push NEIGHBOR_OFFER
000194 2018-08-11 08:53:02,041
000195 2018-08-11 08:53:02,041 Transition LIE_RECEIVED [TWO_WAY] > process_lie,MULTIPLE_NEIGHBORS [None]
```

Network oscillation due to
Run-away feedback loop

Shortest path routing, bandwidth metrics

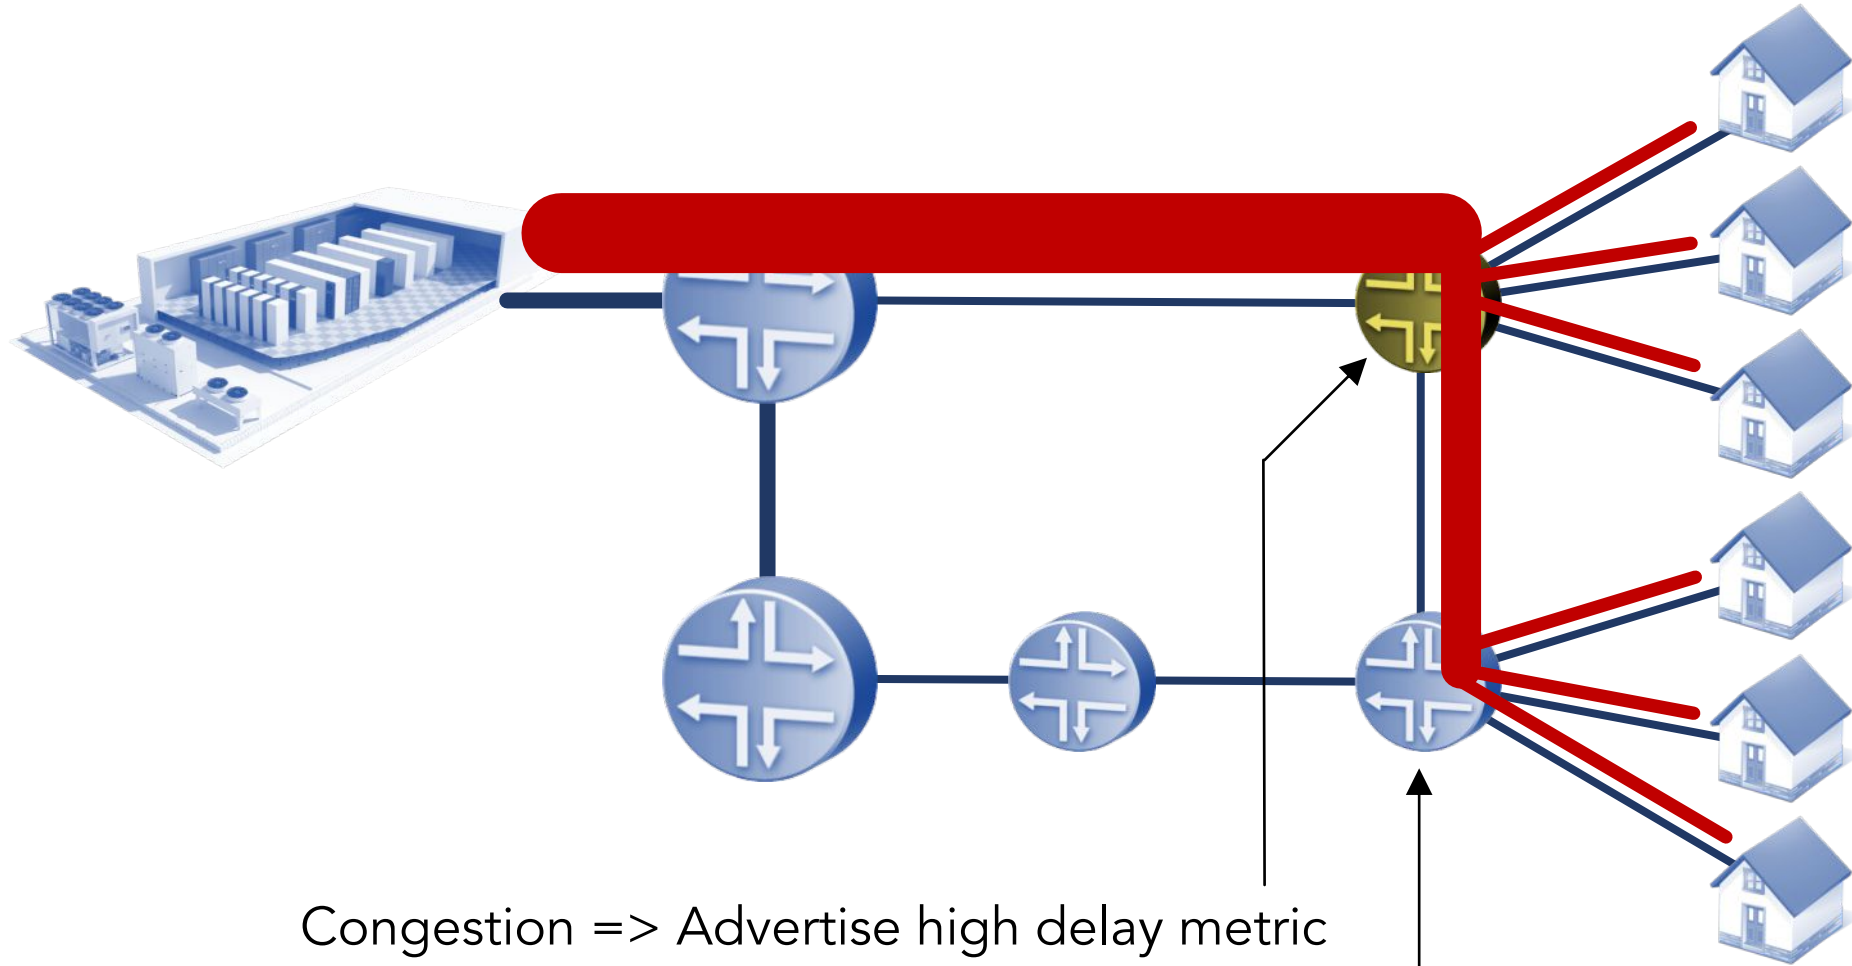


Desired situation: spread the load




This is what traffic engineering would do.

Failed attempts at congestion-sensitive routing




Failed attempts at congestion-sensitive routing



No congestion => Advertise low delay metric

Congestion => Advertise high delay metric

Oscillation due to unstable feedback loop



Delay-based metrics have been attempted many times, and have failed as many times

Take-aways for quantum control protocols

- Make sure your protocols don't have run-away scenarios
- Beware of run-away replication
- Beware of run-away state machines
- Beware of run-away feedback loops

Thank you.